

Lu Sun, and many more.

A Notebook on Machine Learning Updates



*To my family, friends and communities members who
have been dedicating to the presentation of this
notebook, and to all students, researchers and faculty
members who might find this notebook helpful.*



Contents

Foreword	ix
Preface	xi
List of Figures	xiii
List of Tables	xv
I Technologies	1
1 Transformer	3
1.1 It Seems That Attention is All You Need	3
1.2 CNN, RNN, and Their Limitations	4
1.2.1 Brief Review of CNN	5
1.2.2 Brief Review of RNN	6
1.3 Attention Mechanism	8
1.4 Transformer	8
1.5 Research Trends	9
2 Large Language Model	11
2.1 Brief Introduction to LLM	11
2.2 LLM Training from Scratch	11
2.2.1 Base Model	11
2.2.2 Assistant Model	11
2.3 LLM Limitations	12
2.4 LLM Technology Trends	12
2.5 Good Practices to Use LLM	12
2.6 LLM Environment	12
3 Model Fine Tuning and Knowledge Distillation	13
3.1 Transfer Learning VS Model Fine Tuning	14
3.2 Model Fine Tuning	14
3.2.1 Full Fine Tuning	14
3.2.2 Partial Fine Tuning	15
3.3 Low-Rank Adaptation of LLMs (LoRA)	16
3.3.1 Original LoRA	16
3.3.2 LoRA Extensions	16

3.3.3	LoRA Practice	17
3.4	Knowledge Distillation	17
II	Products	19
4	OpenAI	23
4.1	General Introduction	23
4.1.1	Background	23
4.1.2	Business Model	23
4.2	ChatGPT	23
4.3	OpenAI API	24
4.4	DELL-E-2	24
5	Google AI	25
5.1	General Introduction	25
5.1.1	Background	25
5.1.2	Business Model	25
5.2	TensorFlow	25
5.2.1	Summary of Capability	26
5.2.2	Examples	26
5.3	Bard	26
6	Meta AI	27
6.1	General Introduction	27
6.1.1	Background	27
6.1.2	Business Model	27
6.2	PyTorch	28
6.2.1	Summary of Capability	28
6.2.2	Examples	28
6.2.3	PyTorch Lightning	28
6.3	LLaMA	28
6.3.1	LLaMA: Open and Efficient Foundation Language Models	29
6.3.2	LLaMA Services	32
6.3.3	LLaMA-Relevant Use Cases, Tools, and Projects	32
III	Platforms	35
7	Microsoft Azure	37
7.1	Azure Support for OpenAI	38
7.2	Windows Copilot Solution: Microsoft 365 Copilot	39
7.2.1	Microsoft 365 Copilot	39
7.2.2	AI-Assisted Code Generation	39
7.3	User-Defined APPs Copilot Solution: Copilot Stack	40
7.3.1	VS and VS Code Copilot Chat	40
7.3.2	Create Plug-in for AI Models	40

<i>Contents</i>	vii
7.4 Azure AI Studio	41
7.5 Microsoft Fabric	41
Bibliography	43



Foreword

If software or e-books can be made completely open-source, why not a notebook?

This brings me back to the summer of 2009 when I started my third year as a high school student in Harbin No. 3 High School. In around the end of August when the results of Gaokao (National College Entrance Examination of China, annually held in July) are released, people from photocopy shops would start selling notebooks photocopies that they claim to be from the top scorers of the exam. Much curious as I was about what these notebooks look like, never have I expected myself to actually learn anything from them, mainly for the following three reasons.

First of all, some (in fact many) of these notebooks were more difficult to understand than the textbooks. I guess we cannot blame the top scorers for being so smart that they sometimes make things extremely brief or overwhelmingly complicated.

Secondly, why would I want to adapt to notebooks of others when I had my own notebooks which in my opinion should be just as good as theirs.

And lastly, as a student in the top-tier high school myself, I knew that the top scorers of the coming year would probably be a schoolmate or a classmate. Why would I want to pay that much money to a complete stranger in a photocopy shop for my friend's notebook, rather than requesting a copy from him or her directly?

However, things had changed after my becoming an undergraduate student in 2010. There were so many modules and materials to learn in a university, and as an unfortunate result, students were often distracted from digging deeply into a module (For those who were still able to do so, you have my highest respect). The situation became even worse as I started pursuing my Ph.D. in 2014. As I had to focus on specific research areas entirely, I could hardly split much time on other irrelevant but still important and interesting contents.

This motivated me to start reading and taking notebooks for selected books and articles, just to force myself to spent time learning new subjects out of my comfort zone. I used to take hand-written notebooks. My very first notebook was on *Numerical Analysis*, an entrance level module for engineering background graduate students. Till today I still have on my hand dozens of these notebooks. Eventually, one day it suddenly came to me: why not digitalize them, and make them accessible online and open-source, and let everyone read and edit it?

As most of the open-source software, this notebook (and it applies to the other notebooks in this series as well) does not come with any “warranty” of any kind, meaning that there is no guarantee for the statement and knowledge in this notebook to be absolutely correct as it is not peer reviewed. **Do NOT cite this notebook in your academic research paper or book!** Of course, if you find anything helpful with your research, please trace back to the origin of the citation and double confirm it yourself, then on top of that determine whether or not to use it in your research.

This notebook is suitable as:

- a quick reference guide;
- a brief introduction for beginners of the module;
- a “cheat sheet” for students to prepare for the exam (Don’t bring it to the exam unless it is allowed by your lecturer!) or for lecturers to prepare the teaching materials.

This notebook is NOT suitable as:

- a direct research reference;
- a replacement to the textbook;

because as explained the notebook is NOT peer reviewed and it is meant to be simple and easy to read. It is not necessary brief, but all the tedious explanation and derivation, if any, shall be “fold into appendix” and a reader can easily skip those things without any interruption to the reading experience.

Although this notebook is open-source, the reference materials of this notebook, including textbooks, journal papers, conference proceedings, etc., may not be open-source. Very likely many of these reference materials are licensed or copyrighted. Please legitimately access these materials and properly use them.

Some of the figures in this notebook is drawn using Excalidraw, a very interesting tool for machine to emulate hand-writing. The Excalidraw project can be found in GitHub, [excalidraw/excalidraw](https://github.com/excalidraw/excalidraw).

Preface

Artificial intelligence (AI) and specifically machine learning were originally a research topic under the scope of control systems, and it is primarily used for system identification (“machine learning was used to be called system identification”, as some professors say). Its artificial neuron network (ANN) structure is a promising approach for building highly nonlinear self-tuning functions.

Due to the limitations of computational power and data storage in the past years, training large-size ANN with massive data points was hardly possible. With the advancement of computer and material science in the beginning of the 21st century, in particular the development of GPU and TPU, nowadays we can manage networks with dozens of layers, each containing hundreds of neurons. This is referred as the deep learning (DL) network structure. With DL as the building block, effective convolutional neural networks (CNN) and recurrent neural networks (RNN) have been proposed which are effective at finding trends in spacial and sequential data. The proposition of transformers, yet another attention-based deep learning structure, has taken natural language processing to the next level.

AI has grown so significantly in the past decades that it is now considered a separate area apart from classic control systems, and it is drawing more and more attentions. To put it in perspective, while a well-known textbook on control systems may have 10,000+ citations records after years since its publication, a famous conference paper on modern AI can easily have 50,000+ citations after a few months.

This notebook does not focus on the introduction of AI mechanisms, as most of which are already covered in control systems related notebooks. The purpose of this notebook is mainly to keep up-to-date with the latest technology in machine learning, starting from transformer and its applications in chatbot, the most famous of which being ChatGPT developed by OpenAI. The notebook cares not only the techniques, but also the business models used in these technologies and companies.

Some part of this notebook is written by ChatGPT-4.



List of Figures

1.1	A demonstration of CNN kernel. The input is given by the white box (3D), and the kernel by the red box.	6
1.2	A example of RNN.	7



List of Tables

6.1	Data sources for LLaMA.	30
6.2	LLaMA models by size.	31





Part I

Technologies



1

Transformer

CONTENTS

1.1	It Seems That Attention is All You Need	3
1.2	CNN, RNN, and Their Limitations	4
1.2.1	Brief Review of CNN	5
1.2.2	Brief Review of RNN	6
1.3	Attention Mechanism	8
1.4	Transformer	8
1.5	Research Trends	8

This chapter introduces the development of “transformer” in its main branch. Transformer is an AI framework initially proposed in 2017 in a paper “Attention is All You Need”, and it has magically revolutionized how we train AI for natural language processing and interpreting.

1.1 It Seems That Attention is All You Need

It started in 2017 when “transformer” is proposed in a paper “Attention is all you need”, by Vaswani, et al., a bunch of guys mainly from Google.

“Attention” was not a new concept previous to this paper. It was proposed a few years earlier than the transformer, and had been integrated into several CNN and RNN models to boost their performances to capture long-range dependencies. In these researches, attention mechanism is more often a supporting role than the main character. It is quite surprising indeed, that when getting rid of CNN and RNN and purely relying on attention mechanism, everything actually gets better: it is faster to train the network, and its performance gets improved.

Does this mean the RNN, which we had been relying on so much, is in fact the drawback to natural language processing? The following is the answer of ChatGPT-4 of this question.

Q: By getting rid of RNN and using attention mechanism alone, natural language processing performance gets better. Is this correct? Does this mean that RNN is a drawback to natural language processing?

A: Yes, by getting rid of RNNs and using the attention mechanism alone, the Transformer model demonstrated improved performance in natural language processing tasks. However, it's important to note that RNNs are not inherently a drawback to NLP; they were successfully used for many tasks before the advent of Transformers.

RNNs, particularly LSTM (Long Short-Term Memory) and GRU (Gated Recurrent Unit) variants, were designed to handle sequences and capture dependencies in the data. They have been effective in many NLP tasks, such as language modeling, machine translation, and sentiment analysis. However, RNNs do have certain limitations, such as difficulty in capturing long-range dependencies and sequential processing, which can lead to slower training and inference times.

The attention mechanism, as employed in the Transformer model, addresses some of these limitations. It enables parallel processing of input sequences, leading to faster training and inference, and effectively captures long-range dependencies without the vanishing gradient problem that can affect RNNs.

While Transformers have shown superior performance in many NLP tasks, RNNs still have their place in certain scenarios, especially when working with resource-constrained environments or smaller datasets. They can be more lightweight and computationally efficient compared to large-scale Transformer models.

In summary, RNNs are not a drawback to NLP per se, but their limitations have been mitigated by the introduction of the attention mechanism in Transformer models, which has led to improved performance in a wide range of NLP tasks.

(April 14, 2023, ChatGPT-4)

In the remaining part of the chapter, a quick review of CNN and RNN is given and their limitations introduced. Attention mechanism is illustrated. The original transformer model is explained. Finally, research trends of transformer as of today are briefly summarized.

1.2 CNN, RNN, and Their Limitations

This section does not focus on the detailed introduction of CNN and RNN mechanisms. Instead, only their main features are reviewed, and their limita-

tions discussed. This section does not reflect on the state-of-art of researches in CNN and RNN.

The conventional feedforward network (network without cycles) with dense ANN is widely used. It can be trained systematically using back-propagation methods such as stochastic gradient descent. However, dense ANN is not very good at interpreting correlations among the inputs. At least, it is not good at doing it efficiently.

For example, consider the input to be pixels of an image of size $3 \times 128 \times 128$, where 3 corresponds to the RGB of the image, and 128×128 its pixel size. It is obvious to a human that the pixel at (1, 1) should definitely has a closer relationship with the one at (1, 2) (they are probably forming an object in the image together) than (128, 128). The same applies to sequential inputs. Consider a signal sampled continuously. The first and second samples are very likely to be more strongly correlated than the first and last samples.

In a conventional dense ANN, the correlation cannot be captured efficiently. In contrast, a dense ANN would treat all the inputs “equally” in a symmetric manner. To enforce the ANN to “memorize” the correlation, it would require a lot more layers and nodes (which is often considered low-efficient), and requires more data points during the training. Other problems of conventional dense ANN include, for example, the lack of ability in handling data with arbitrary length.

CNN and RNN try to tackle the above problems by implementing a “pre-processing” stage, where the correlation of the spacial and sequential data is first abstracted using some mechanism, and the correlation information is sent as (additional) inputs to the followed dense ANN.


The details of CNN and RNN mechanisms are not discussed in this notebook. Brief reviews are given as follows.

1.2.1 Brief Review of CNN

CNN is a type of ANN structure designed to handle grid-like data. It is effective when dealing when data spatially correlated, thus becomes very popular in computer vision. It defines “kernel” that aggregates nearby pixels information before sending it to a dense network.

A CNN kernel, also known as a filter, is a small matrix of weights that slides across the input image or feature map to perform a mathematical operation called convolution. An demonstration of CNN kernel is given in Fig. 1.1. Multiple kernels can be defined on the same layer to handle the same feature map, each kernel associated with an output channel. In practice, each kernel or channel is designed to detect a specific features in the feature map. For example, there might be a kernel detecting edges, while a second kernel detect color codes.

Notice that CNN differs quite largely from transformer in the problems they are expected to address. CNN is more for spatial data processing such

**FIGURE 1.1**

A demonstration of CNN kernel. The input is given by the white box (3D), and the kernel by the red box.

as image processing, while transformer targets more on sequential data processing such as natural language processing and machine translation.

1.2.2 Brief Review of RNN

RNN is a connectionist model with the ability to selectively pass information across sequence steps [3]. It is good at handling sequence of data such as voice message, text contents, or a flow of images (videos). It is worth mentioning that the “sequence” does not necessarily mean a time sequence. Nevertheless, without losing generality, we will consider time sequence in the review for simplicity and convenience.

Denote inputs $x(1), x(2), \dots, x(k), \dots$ where $x(k)$ is a vector sampled at time instant k . The length of the sequence may be finite or infinite. In the case of finite sequence, its maximum sample index is denoted by T . For example, in the context of natural language processing, each input might be a word in a dictionary. For example, $x(1) = \text{“Pandas”}$, $x(2) = \text{“are”}$, $x(3) = \text{“so”}$, $x(4) = \text{“cute”}$, $x(5) = \text{“!”}$. The corresponding target output sequence is given by $y(1), y(2), \dots, y(k), \dots$, respectively.

RNN differs from the conventional dense ANN by introducing “recurrent edges”, which allows the output of hidden layers at $k - 1$ be used as additional inputs to the system at k . This means, at any time k , the input of the system includes both $x(k)$ and also selected $h(k - 1)$, where $h(\cdot)$ is the outputs of hidden layers. We can think of the “weights” of a trained RNN the “long-term memory” that does not change with specific sequence of inputs, while the information passing through recurrent edges the “short-term memory” that links previous inputs with future inputs.

A demonstration is given in Fig. 1.2. Notice that each hidden layer is a

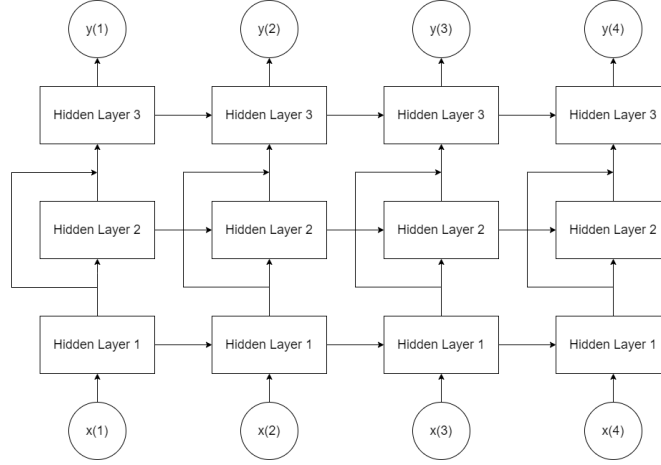


FIGURE 1.2
A example of RNN.

multi-input-multi-output subsystem containing multiple nodes. Different from a conventional dense ANN, each hidden layer takes additional inputs from its corresponding hidden layer in the previous instant. It is also common to see “bypass” (this is widely used in different ANN structures, not unique to RNN) for better performance of the system.

RNN has some limitations. One of the major problem is that it is difficult to train an RNN even for the basic standard feedforward networks. The optimization of RNN is NP-complete. It is especially difficult for RNN to learn long-range dependencies due to the vanishing and exploding gradients problem that could occur when backpropagating errors across many time-steps (long sequence) [3]. This is one of the main challenges why RNN has difficulties building on long-range dependencies. The vanishing and exploding gradient problems are caused by the structure of the system as well as the backpropagation-based training methods.

Different approaches have been proposed to prevent vanishing and exploding gradient problems. Famous ones among these approaches include strategic weight initialization, long short term memory (LSTM), gated recurrent units (GRUs), skip connections, and more. Many of these approaches try to reduce the effect of vanishing and exploding gradient problems by carefully design the ANN structures. For example, both LSTM and GRUs introduce “memory cells” with built-in “gates” that balance and control the flow of information from previous cells versus current inputs. The cells are used to replace the traditional perceptron nodes. These approaches have made the training of RNN a feasible problem. LSTM and GRUs are almost certainly used in modern RNNs.

Bidirectional RNN (BRNN) is proposed at about the same time with

LSTM. It allows information to travel not only from previous hidden layers to future layers, but also from future hidden layers to previous layers. LSTM and BRNN can be used together to boost the RNN performance. Notice that BRNN cannot run continuously as it requires fixed endpoints in both the future and the past. It is useful for prediction over a sequence of fixed length, such as part-of-speech tagging in natural language processing.

Another problem that people have found during the training of RNN is local optima. However, recent studies have shown that local optima is not as serious issue as we might thought when the network is large, since many critical points are actually saddle points rather than local minima.

Successful implementations of the above RNN structures include natural language translation such as [5] where an encoder-decoder structure is used, each is a LSTM. Another example is image captioning, where the AI tries to explain what is in an image using texts. A solution to this is to use CNN to encode the image, and use LSTM to decode to generate texts. Following similar ideas is hand-writing reorganization.

1.3 Attention Mechanism

“nobreak

1.4 Transformer

While CNN is mainly used for spatial data processing such as computer vision, both RNN and transformer are mainly used for sequential data processing. It is worth mention in the very beginning that transformer does not guarantee superior performance in all scenarios comparing with traditional RNN-based natural language processing models such as LSTM.

For one thing, transformer consumes larger computational capabilities. Transformer processes data in batch, while RNN does them in sequence, meaning that RNN can response faster in some real-time applications. It is difficult for the transformer to process super-long text due to the computational burden (quadratic computational complexity with respect to sequence length), while RNN can process arbitrarily long text, although the later suffers from capturing long-range dependencies.

With the above been said, the transformer does demonstrated superior performance than RNN in many tasks. The mechanism and the reasons why it performs better in these occasions are introduced as follows.

1.5 Research Trends



2

Large Language Model

CONTENTS

2.1	Brief Introduction to LLM	11
2.2	LLM Training from Scratch	11
2.2.1	Base Model	11
2.2.2	Assistant Model	11
2.3	LLM Limitations	11
2.4	LLM Technology Trends	12
2.5	Good Practices to Use LLM	12
2.6	LLM Environment	12

“nobreak

2.1 Brief Introduction to LLM

“nobreak

2.2 LLM Training from Scratch

“nobreak

2.2.1 Base Model

“nobreak

2.2.2 Assistant Model

“nobreak

2.3 LLM Limitations

“nobreak

2.4 LLM Technology Trends

“nobreak

2.5 Good Practices to Use LLM

“nobreak

2.6 LLM Environment

3

Model Fine Tuning and Knowledge Distillation

CONTENTS

3.1	Transfer Learning VS Model Fine Tuning	13
3.2	Model Fine Tuning	14
3.2.1	Full Fine Tuning	14
3.2.2	Partial Fine Tuning	15
3.3	Low-Rank Adaptation of LLMs (LoRA)	16
3.3.1	Original LoRA	16
3.3.2	LoRA Extensions	16
3.3.3	LoRA Practice	17
3.4	Knowledge Distillation	17

Training an LLM from scratch is usually time and computational power consuming. For cost efficiency, it is usually preferable to train a new LLM from an existing one.

The first approach of such kind is to re-use the weights of the existing model in the new model. For example, one may train a new model with its weights initialized as the existing model weights. Alternatively, one may modify the weights of the existing model, and save them as the new model. This approach is often known as model fine tuning.

Alternatively, consider building a new and smaller model. Take the new model as the “student”, and let the existing LLM model be the “teacher” that teaches the new model. The new model is trained to mimic the outputs of the existing model in a specific domain. The new model is more efficient and runs faster with less cost compared with the existing model due to the smaller size. This approach is often known as knowledge distillation.

Model fine tuning and knowledge distillation are introduced in this chapter.

3.1 Transfer Learning VS Model Fine Tuning

Transfer learning and model fine tuning are sometimes used interchangeably. They share similarities and sometimes may indeed refer to the same operation. However, it is worth clarifying the differences of the two concepts here.

Transfer learning refers to any attempt that tries to build a model for a new task using an existing model originally designed for another task. For example, consider an existing CNN for edge and shape detection. It is possible to add a new dense network on top of the CNN for cat detection. The weights of the original CNN is not changed (hence, it is not model fine tuning). It is just that its output is used as the input to a newly designed downstream network.

Model fine tuning is a widely used technique for transfer learning. It refers to the attempt of building a new network by modifying fully or partially the weights of an existing network. In a full fine tuning, the new network uses the weights of the existing network as a starting point, and further update them using new training data. In partial fine tuning, some of the weights of the existing network are fixed (usually bottom layers, i.e. the layers close to the input) while others are updated (usually top layers, i.e., the layers close to the output) using new training data. Partial fine-tuning often has the advantage of saving computational and memory burden, as less parameters are updated.

3.2 Model Fine Tuning

Both full and partial model fine tuning methods are introduced in this section.

3.2.1 Full Fine Tuning

There are many studies on model fine tuning. The most intuitive fine tuning method is full fine tuning, which basically continue training an existing model using new training data to boost its performance. There are many tools for full fine tuning an LLM (or other AI models). Technically speaking, the methods used to train the model from scratch can be used here to continue training it on new data. An example of such methods is ADAM optimizer, which is widely used in LLM training.

In the context of LLM, a full fine tuning algorithm often tries to find such weights Φ that maximizes

$$\sum_{(x,y) \in Z} \sum_{t=1}^{|y|} \log(P_{\Phi}(y_t|x))$$

where x, y are tokenized sequence, $P_{\Phi}(y|x)$ a pre-trained autoregressive language model parameterized by Φ , and $Z = \{(x_i, y_i), i = 1, \dots, N\}$ the context-target pairs. The same objective functions applies for training an LLM from scratch, in which case Φ is initialized with random values instead of a pre-trained model.

The disadvantage of full fine tuning is obviously the computational and memory cost, as there are many values to be parameters to be updated. This is true especially in LLM, and there are often billions of parameters in a model (for example, GPT-3 has 175 billions of parameters).

3.2.2 Partial Fine Tuning

Partial fine tuning of an AI model refers to fixing some of the parameters in an existing model and updating only selected parameters when training it on new datasets. It is widely used in ANN, not only to LLM. A general way of partial fine tuning is to fix the bottom layers of an AI model and only update the top dense layers, as it is widely accepted that top layers contain more comprehensive features of a task.

An advantage of partial fine tuning over full fine tuning is saving computational cost, as there are less parameters to be tuned. In addition, it also saves storage space when there are multiple fine tuned models, as part of the weights are shared and do not need to be stored in duplicated pieces.

When comes to LLM partial fine tuning, many algorithms have been proposed with details implemented carefully to enhance computational efficiency. Some examples are given below.

Prompt Engineering

Prompt engineering refers to the fine tuning of an LLM base model to assistant model, i.e., to train it using “request(prompt)-and-response” datasets on top of the LLM base model so that it can answer human questions. Notice that both full and partial fine tuning can be applied in prompt engineering.

Parameter-Efficient Adaptation

Parameter-efficient adaptation refers to the technique where “adapter layers” are inserted into the existing LLM. Only the adapter layers weights are updated during the fine tuning. This changes the topology and schematic of the LLM, which potentially adds complexity and computational load to the network.

It is argued that since the adapter layers are often smaller in size ($< 1\%$) compared with the original LLM, thus the small bottleneck may not largely affect the general performance of the system. However, in LLM practice, since it relies very heavily on hardware parallelism to keep the latency low, and adapter layers can be processed only sequentially, the introduced inference latency is quite noticeable.

Low-Rank Structure Based Adaption

Instead of fine tuning the weights directly, consider finding the optimal “delta of weights”. The ΔW is obtained using the new training data, and it is added to the original weights $W = W_0 + \Delta W$ to form the new model. In practice, $\Delta W = BA$ is formed by the product of two low-rank matrices $B \in \mathbb{R}^{d \times r}$, $A \in \mathbb{R}^{r \times k}$ obtained during the fine tuning, hence the name low-rank structure based adaption. An example of such techniques is Low-Rank Adaption of LLMs (LoRA), which will be introduced in more details in later sections.

LoRA has some advantages. It has a low computational cost when a small r is used. The value of r serves as a balancing factor, where the larger r , the more close it is to the full fine tuning. Unlike parameter-efficient adaptation, LoRA does not change the structure of the model, thus not adding computational complexity. Finally, the original model W_0 can be recovered easily from W by subtracting ΔW . These advantages make LoRA a popular fine tuning method.

3.3 Low-Rank Adaptation of LLMs (LoRA)

LoRA, its extensions, and its practice usage are introduced in this section.

3.3.1 Original LoRA

As introduced earlier, LoRA tries to find the optimal $B \in \mathbb{R}^{d \times r}$, $A \in \mathbb{R}^{r \times k}$ that form the “delta of weights” $\Delta W = BA$. The fine tuned weights $W = W_0 + \Delta W$ is the sum of the original weights and the delta of weights.

Hyper parameter r decides the complexity of the fine tuning. When r is chosen small, the fine tuning is computationally lite; when chosen large, the fine tuning is computationally heavy and the performance of the model would converge to a full fine tuning model.

The model stores W_0 , B and A separately. When executing the model, $W = W_0 + \Delta W = W_0 + BA$ is calculated and used, which integrates the fine tuning knowledge into the original model and reduces inference latency. When fine tuning for a different task, only B , A needs to be re-calculated and there is no need to change W_0 .

Studies have been carried out on how ΔW relates to W , and how to choose r in practice, etc. Details are not given here. See [2] for more details.

3.3.2 LoRA Extensions

There are some extensions for LoRA. For example, QLoRA modifies LoRA and claims that it can train a model with much smaller GPU (780GB memory

GPU to single core 48GB memory GPU) and less time consumption (< 1 day) than the state-of-the-art, while not sacrificing too much performance. See [1] for more details. QLoRA makes 33B and 65B LLMs more accessible to regular PCs.

3.3.3 LoRA Practice

“nobreak

3.4 Knowledge Distillation





Part II

Products



Many data companies, including Microsoft, Google, Meta (formerly known as Facebook) and Amazon, provide their own data models and AI engines of different kinds. Famous and AI engines and models include TensorFlow by Google, PyTorch by Meta, and ChatGPT by OpenAI, which is a company supported by Microsoft.

These data companies usually open (part of) the source code of these models and engines in the community and encourage individual developers to learn and use them. At the same time, they provide enterprise level services such as model fine-tuning to make profit from the tools.

Some state-of-art models, engines, as well as the business models of these companies are reviewed in this part of the notebook.



4

OpenAI

CONTENTS

4.1	General Introduction	23
4.1.1	Background	23
4.1.2	Business Model	23
4.2	ChatGPT	23
4.3	OpenAI API	23
4.4	DELL-E-2	24

“nobreak

4.1 General Introduction

“nobreak

4.1.1 Background

“nobreak

4.1.2 Business Model

“nobreak

4.2 ChatGPT

“nobreak

4.3 OpenAI API

“nobreak

4.4 DELL-E-2

5

Google AI

CONTENTS

5.1	General Introduction	25
5.1.1	Background	25
5.1.2	Business Model	25
5.2	TensorFlow	25
5.2.1	Summary of Capability	25
5.2.2	Examples	26
5.3	Bard	26

“nobreak

5.1 General Introduction

“nobreak

5.1.1 Background

“nobreak

5.1.2 Business Model

“nobreak

5.2 TensorFlow

“nobreak

5.2.1 Summary of Capability

“nobreak

5.2.2 Examples

“nobreak

5.3 Bard

6

Meta AI

CONTENTS

6.1	General Introduction	27
6.1.1	Background	27
6.1.2	Business Model	27
6.2	PyTorch	27
6.2.1	Summary of Capability	28
6.2.2	Examples	28
6.2.3	PyTorch Lightning	28
6.3	LLaMA	28
6.3.1	LLaMA: Open and Efficient Foundation Language Models	29
6.3.2	LLaMA Services	32
6.3.3	LLaMA-Relevant Use Cases, Tools, and Projects	32

Meta, formerly known as Facebook, provides variety of solutions for AI, including computer vision, stream data processing, natural language processing, and many more.

6.1 General Introduction

“nobreak

6.1.1 Background

“nobreak

6.1.2 Business Model

“nobreak

6.2 PyTorch

PyTorch, similar with TensorFlow, is a Python-based generic AI engine developed by Meta. It is easy to deploy and has been widely used in variety of applications including computer vision and sequential signal processing. PyTorch is freely accessible can be installed on customer machines.

6.2.1 Summary of Capability

“nobreak

6.2.2 Examples

“nobreak

6.2.3 PyTorch Lightning

AI engines have tried hard to make themselves easy to use. Ultimately, the AI solution architect should only concern about the network architecture and not the programming. PyTorch Lightning, a PyTorch add-on provided by Lightning.AI, is one of the many attempts to make that happen. PyTorch Lightning comes in the form of an additional Python package.

Notice that PyTorch packages and libraries are still required, as PyTorch Lightning is an add-on, but not a replacement, to PyTorch. PyTorch changes only the way of programming, but not the capability or performance of the models.

Examples are given in *pytorchlightning.ai* where PyTorch programs with and without PyTorch Lightning are compared. The examples show that PyTorch Lightning is helpful with simplifying the pipeline deployment and tidying the programming logic.

6.3 LLaMA

LLaMA is the large language model published by Meta AI. The associated paper that introduced LLaMA for the first time was published in February 2023 on ArXiv as [6]. The initial LLaMA is trained on publicly available datasets exclusively.

Different levels of complexity of the models are provided. Like many other LLMs, Meta AI provides charged service for fine-tuning the model. There is also a re-write of LLaMA, namely Lit-LLaMA, that is completely open source

under Apache 2.0 license. Both LLaMA and Lit-LLAMA can be fine-tuned via tools such as LoRA and Stanford Alpaca.

Comparing with other models such as ChatGPT-3, LLaMA is claimed to use smaller models (with less number of parameters) to achieve the same performance.

Notice that LLaMA is released as the base model, not the assistant model. This is different from GPT-4, in which case up till this moment only assistant model is released.

6.3.1 LLaMA: Open and Efficient Foundation Language Models

LLaMA is another implication of (modified) transformer proposed in [7]. A GitHub repository has been created and made public about LLaMA. The repository is given here: github.com/facebookresearch/llama.

Budget-Oriented LLM Design

There is a trad-off between using large model versus small model, given a fixed amount of budget to create and run the model. A large model often implies better performance potentially. But this is true only if it is trained on large amount of data. Otherwise, over-fitting may happen, which increases the variance of the system. A large model plus a large amount of data is sometimes economically inefficient.

Given a fixed budget, comparing using large model with small amount of data which gives only mediocre performance, it might be wise to use a relatively smaller model, but and spend the remaining budget training it with a large amount of data. It is possible to match the performance of the large but under-trained model using the small but well-trained model. The later approach would often mean longer training time, but less resources at the same time and lower execution cost in its daily run post training.

LLaMA implements the later approach. It assumes different levels of fixed budgets, and discusses what might be the optimal choice of LLM under each budget.

Data Sources

Only publicly available data is used in the training. The data sources for LLaMA is summarized Table 6.1. The total amount of training data is about 4.75TB, or 1.4T tokens after tokenization. To put it into perspective, all the work that Shakespeare has done, if converted to plain text, is about 5MB (1/950000 of 4.75TB).

Model

Different sizes of models are tested. The size of the model, i.e., number of parameters, are 6.7, 13.0, 32.5, and 65.2 billion respectively, which is sum-

TABLE 6.1

Data sources for LLaMA.

Data Source	Percentage	Description
Common Crawl	67%	Common Crawl English data is used. Common Crawl builds and maintains an open repository of web crawl data that can be accessed and analyzed by anyone. The data can be used for research, analysis, and education.
C4 Dataset	15%	C4 is a colossal, cleaned version of Common Crawl's web crawl corpus. It was based on Common Crawl dataset.
Github	4.5%	Github repositories under Apache, BSD and MIT licenses are used.
Wikipedia	4.5%	Wikipedia data with different Latin and Cyrillic scripts are used.
Public domain books	4.5%	Free ebooks from Gutenberg project, etc., are used.
ArXiv	2.5%	Latex files from ArXiv are used.
Stack Exchange	2%	Stack Exchange is a website of high-quality questions and answers of variety of domains.

marized in Table 6.2. Notice that the number of parameters listed above is significantly smaller than many of the available models.

TABLE 6.2

LLaMA models by size.

Model Name	Number of Parameters	Data Size for Training
LLaMA 7B	6.7 billion	1T token
LLaMA 13B	13.0 billion	1T token
LLaMA 33B	32.5 billion	1.4T token
LLaMA 65B	65.2 billion	1.4T token

Notice that 1.4T token of training data is approximately 4.75TB of plain text data.

The model architecture is based on the original transformer proposed in [7], with following modifications:

- Add pre-normalization to the input of each transformer sub-layer, instead of normalizing the output of the layers.
- Use SwiGLU activation function instead of ReLU.
- Use rotary positional embedding instead of absolute positional embedding. Rotary Position Embedding, or RoPE, is a type of position embedding which encodes absolute positional information with rotation matrix and naturally incorporates explicit relative position dependency in self-attention formulation. See [4] for details.

AdamW optimizer with dynamic learning rate is used.

Mini-batch training of 4M tokens are used for all models. Models with size 6.7B 13.0B are trained with 1.0T tokens, and the rest are trained with all 1.4T tokens, as shown in Table 6.2. Variety of methods are implemented to speed up the training.

Performance Evaluation

The benchmarks to evaluate and compare the performance of LLaMA with other models are mainly from the following categories:

- Common sense reasoning. E.g., “If someone is holding an umbrella, what might be the reason?”
- Closed-book question answering. E.g., “What is the capital city of France?”
- Reading comprehension. E.g., give a paragraph that introduces generic search, then ask “what is the major advantage of generic search?”

- Mathematical reasoning. E.g., “what is the sum of the first 5 prime numbers?”
- Code generation. E.g., “Generate a code in Python that calculates the first 100 values in Fibonacci series.”
- Multitask language understanding. E.g., “Translate English into French, then give a summary.”
- Bias, toxicity and misinformation.

LLaMA has shown improved performance compared with the model with similar size.

Training Effort

Take the largest LLaMA-65B as an example. An “A100-80GB” GPU (cost approx 15k USD) is used and trained using 1022k GPU-hours. The total power consumption is 449MWh, which is worth approximately 90k USD household price.

6.3.2 LLaMA Services

LLaMA is not yet a fully commercialized AI model, and its code as given in the LLaMA Github repository is under GPL license. Research team behind LLaMA suggests that the primary intended uses of LLaMA should be research on LLMs. The model (at least version 1 of the model) is a base or foundational model, and should not be used on downstream applications without any tuning and mitigation of risks.

The instruction to download and use LLaMA can be found in its Github repository. In short, the user needs ask approval to use the trained LLaMA model. The user needs to setup prerequisite Python environment in order to run the model. With the approval, the user can download the trained LLaMA model and the tokenizer, and run the model on its own machine. An example in Python is provided, which helps the user to verify the downloaded items, test the environment, and get a quick start.

Variety of tools and projects have been developed, either as a demonstration of LLaMA use case, or as a add-on to make the use of LLaMA more convenient. Examples include LoRA, a tool by Microsoft that can be used to fine-tuning models, and Lit-LLaMA, an open-source recreation of LLaMA, and a few more. Details are introduced in later sections.

6.3.3 LLaMA-Relevant Use Cases, Tools, and Projects

LoRA

LoRA (Low-Rank Adaptation of Large Language Models) is an open-source tool developed by Microsoft that enables transfer learning /knowledge distillation/model fine-tuning of LLMs. This is quite a important tool, as training LLMs from scratch can be very time and energy consuming.

LoRA is first proposed in [2]. A Github repository github.com/microsoft/LoRA contains the source code of LoRA implementation in Python package `loralib`, which currently support PyTorch. Indeed, LoRA can be used to fine-tune ChatGPT (as under Microsoft), LLaMA and Lit-LLaMA.

The latest version of LoRA is supported by github.com/huggingface/peft, which is a collection of state-of-the-art parameter-efficient fine-tuning methods.

Lit-LLaMA

Lit-LLaMA is an open-source recreation of LLaMA by Lightning.AI. Think of Lightning.AI as something similar with GNU project (Linux project). GNU wants to create a open-source version of UNIX that can be freely distributed. Companies like Red Hat can make a profit from it by providing enterprise level support to the subscribers who uses RHEL. Lightning.AI is essentially trying to do the same: it creates user-friendly interfaces for existing AI models such as PyTorch, and sometimes recreates open-source version of existing (commercialized) models, and make them open-source. It encourages users to use their interfaces and models. It makes profits by providing add-on services such as training models using their servers, etc.

Lit-LLaMA is a recreation of LLaMA model based on nanoGPT, another open-source transformer model. Different with LLaMA which is under GPL license, nanoGPT is under MIT license and Lit-LLaMA Apache license, both of which more friendly for private projects. Lit-LLaMA Github repository is given here github.com/Lightning-AI/lit-llama.

Unlike LLaMA where approval is required to download the trained model, under the scope of Lit-LLaMA, it is fine to download the trained 7B model freely. The customer can run the trained model on its server. A decently powerful GPU is required. Such a power GPU often causes a few thousands (e.g., NVIDIA GeForce RTX 3090 24GB, which can be used to run and fine-tune the model, approx 3k) to tens of thousands (e.g., NVIDIA A100 80GB, which can be used to train LLaMA and Lit-LLaMA from scratch, approx 15k) USD.

Both LLaMA and Lit-LLaMA provides interfaces and example scripts for the users to use the model on their own server. In addition, they also provide interface for LoRA, which can be used to fine-tune the model. Python script examples are provided to use LoRA on Lit-LLaMA in the repository.

Stanford Alpaca

It is quite unfortunate that, in the area of AI, academic researchers and universities have fallen behind. The most powerful models are usually trained using very expensive resources including high-cost GPU clusters, months or

years of human labors, and tens or hundreds of terabytes of data. It is often done by data driven companies who would not make them completely open-source to the public even for research purposes. LLaMA is an outlier in the good sense, and it gives Stanford a chance to make a break through.

Stanford Alpaca 7B is a fine-tuned model from LLaMA 7B model. The cost of the fine-tuning is impressively small (hundreds of USD), and the performance of the system after fine-tuning is comparable with ChatGPT-3.5 (text-davinci-003) which has a significantly larger model of 335B than 7B.

It is common that a fine-tuned model gets its performance boosted. However, it is quite surprising to see a huge boost of performance happening on LLaMA 7B, as this model is famous for being small and already being trained by a huge amount of data. How much potential does LLaMA 7B have that can be further improved? Not to mention, everything is done using only a few hundred USD, given that a GPU alone that can train the model would usually cause thousands of dollars.

It is worth mentioning that text-davinci-003 outputs are used to fine-tune LLaMA 7B to create Stanford Alpaca 7B. This may imply:

- Training Stanford Alpaca 7B is essentially fine-tuning a less-power model (LLaMA 7B) using the outputs from a more powerful model (text-davinci-003), instead of using plain text. This may explain partially why it is so inexpensive, as data collection and preprocessing is usually one of the most expensive parts in AI training and it is now taken care of by text-davinci-003.
- It is suspicious whether the performance of LLaMA 7B can surpass text-davinci-003 in general, if it is trained by text-davinci-003 in the first place. This is because the bias of information can be enlarged during training.
- The above suggests that very likely other techniques than training with text-davinci-003 been used to further improve its performance.

An introduction of Stanford Alpaca is given here crfm.stanford.edu. An introduction to using the model is given in its Github repository github.com/tatsu-lab/stanford_alpaca. As of April 2023, Stanford Alpaca has released the fine-tuning receipt, and it claims that the trained model will be released in the future.



Part III

Platforms



7

Microsoft Azure

CONTENTS

7.1	Azure Support for OpenAI	38
7.2	Windows Copilot Solution: Microsoft 365 Copilot	38
7.2.1	Microsoft 365 Copilot	39
7.2.2	AI-Assisted Code Generation	39
7.3	User-Defined APPs Copilot Solution: Copilot Stack	40
7.3.1	VS and VS Code Copilot Chat	40
7.3.2	Create Plug-in for AI Models	40
7.4	Azure AI Studio	41
7.5	Microsoft Fabric	41

Microsoft Build 2023 was hold on May 24. In the event the hosts claimed that “everyone is a developer” and “what we would have not believed to happen very soon six months ago is happening today”. The above servers as a good summary of the AI spirit as of today: with its help, anything can happen, and anything can happen soon, and everyone can be part of it.

Azure as a platform provides integrated hardware and software supports to the developers in the community. In the hardware portion, Azure enables training and executing the AI models on cut-edge high performance computers. In the software portion, variety of ideas and tools for AI development related tasks have been proposed, including bringing Bing (Microsoft Search Engine) to ChatGPT as a plug-in for up-to-date information retrieving, Windows Copilot, Copilot stack (user-defined copilots for their APPs), Azure AI Studio (user-defined AI pipeline that involves smart plug-ins), Microsoft Fabric (unified data analysis and visualization tool, with many data models supported), and many more. Some of them are introduced in this chapter.

What is the next big thing? There are a few items in the list.

- Domain knowledge AI model.
- Advanced copilot systems and AI platforms, with elevated user experience and data security.
- Smart plug-ins that enhance and complete AI model capabilities.
- Smaller and faster AI model while achieving similar performance tiers, meantime more edge-computing friendly.

7.1 Azure Support for OpenAI

Notice that Azure support for OpenAI and Azure AI Studio indeed share a lot of resources and interfaces. Many operations introduced in this section are done under Azure AI Studio framework in practice. OpenAI is indeed a big chunk in Azure AI Studio, as many of its features are backed up by OpenAI. Azure AI also provides other tools than OpenAI support, and those will be introduced in later sections.

The most important Azure OpenAI feature is that it allows transfer learning of AI model on customer's domain data fairly easily. Azure guarantees the data security, to make sure that customer data is isolated from one another. The transfer learning is done jointly with Azure AI Studio. When creating OpenAI resource on Azure, the customer has the choice to fine-tuning the model using customer's data via Azure AI Studio.

The customer's data may come from different URIs, including Azure data warehouse or other third-party data sources such as AWS. Everything including data used inside Azure AI Studio is kept closed only to the customer. Azure AI Studio has made it very easy to import data into the workspace for training. It is only a few button clicks from Azure UI. When the model answers a question using the customer data, it will also displays the original text from where the information is retrieved.

Plug-ins, both those recommended or even already integrated into GPT-4 and those open-source tools available on GitHub, etc., can be easily enabled in the AI model that the customer fine-tuned using its own data. Some popular plug-ins may enable the following functions: multi-language translation, advanced searching (for example, Azure Cognitive Search service, which can be enabled in the AI model as a plug-in), internet connectivity for up-to-date information, etc.

Azure Cognitive Search

To enhance data retrieving capability, Azure provides Azure Cognitive Search service, which is essential a NoSQL vector-based database what is good at clustering by nature. This demonstrates Azure's multi-model approach for searching (AI model fine-tuned with customer's data, together with SQL/NoSQL database).

Some promising areas of Azure support for OpenAI include smart data retrieving, text summarizing, code generation, and copilot (current focus is mainly on text/contents generation). Some promising domains include smart health care, robot-aided data analysis, and many more.

7.2 Windows Copilot Solution: Microsoft 365 Copilot

Microsoft 365 has been made much smarter than before with the help of advanced AI model and graph database techniques. It is also more flexible, as it supports plug-ins to be integrated into the copilot. AI-assisted code generation is becoming increasingly popular. Microsoft provides multiple tools to help developers with code generation. These tools are introduced as follows.

7.2.1 Microsoft 365 Copilot

Microsoft 365 copilot is the assistant tool built-in to Microsoft 365 that helps the user to quickly retrieve data globally. The copilot has access to Teams, Outlooks, OneDrive, and other Microsoft 365 software and even Power APPs. It can respond to user's requests such as check up-coming events from a calendar, find a past email, and summarizes the email contents. The copilot has an intuitive user interface: it appears as a chatbot in Teams.

Copilot supports user-defined plug-ins. A plug-in can be created from the third-party software APIs systematically following Microsoft's instructions. The plug-in support shows the extendability of copilot.

Copilot supports semantic index, vector database, etc. Think of semantic index and vector database as Microsoft's NoSQL database that can be used to organize information efficiently. In specific, semantic index is Microsoft's graph database (similar with semantic web) to manage data in Microsoft 365 and Power APPs. Copilot is able to create, insert and retrieve data from these database.

Copilot can be used as a human-in-the-loop pipeline when multiple people are working on the same project using Microsoft 365.

Microsoft Teams Share and Microsoft Mesh are proposed. Microsoft Teams Share allows multiple people working on the same project while having a Teams meeting. Microsoft Mesh allows multiple people working in a virtual environment (think of VR games).

7.2.2 AI-Assisted Code Generation

Microsoft provides at least 2 different tools for AI-assisted code generation: GPT-assisted code generation, and GitHub copilot (also known as codex). These two models are trained from different training set, therefore, may perform differently solving the same problem.

Both tools take in the "prompt" (instructions and descriptions) from the developer, and generate codes accordingly. GPT-assisted code is more natural language friendly, while GitHub copilot is usually more specific on the construction of the prompt.

Regarding the copyright of the codes and images generated by the AI,

there is an undergoing debate. Given that the AI model is trained mostly by public data, very likely the generated codes would be regarded as public codes and can be used freely. When comes to images, things can get a bit more complicated. The current walk around to this is to state clearly that the images are generated by AI model when such images are used.

When using GitHub copilot to generate codes, it is recommended to also check the source of the codes that is also provided by copilot along with the codes.

7.3 User-Defined APPs Copilot Solution: Copilot Stack

Here copilot refers to the “assistance” that comes with the software. It can be an interactive chatbot that provides data retrieving or document editing functions. In this regard, ChatGPT by itself is a general-purpose copilot. Every automated pipeline has a copilot working in the background to sort out all the tasks need to be done one by one. Here is another example, Podcast Copilot github.com/microsoft/PodcastCopilot, which is used in Microsoft Build 2023, is an example of a copilot generating contents that can be posted to social media.

Microsoft Azure provides platform and tools, namely Copilot Stack, for developers to develop user-defined copilots for their APPs. Many features are enabled in Copilot Stack, including fine-tuning models using existing models (such as ChatGPT 3.5 and 4), using plug-ins, and calling open-source models found on GitHub.

7.3.1 VS and VS Code Copilot Chat

Visual Studio and VS Code Copilot Chat is essentially natural-language-interface GitHub Copilot built-in to VS Code that helps to make programming easier. It can create code using the comments left in the programming script, and interpret raised error, and fix the code.

7.3.2 Create Plug-in for AI Models

It is easy to create a plug-in on GitHub using its built-in tools, and call that API from AI models such as ChatGPT. Just pass the URL of the plug-in to ChatGPT, and ChatGPT will try to retrieve information using that plug-in automatically. The similar applies to the AI mode created in users own copilots built inside Copilot Stack.

The plug-in can be made either private or open to public, up to the developer’s choice. The plug-in development can be put into CI/CD using GitHub actions, and deployed in containers on Azure managed by k8s.

Try this function using the demo given by github.com/Azure-Samples/openai-plugin-fastapi.

7.4 Azure AI Studio

Azure tries its best to help the customer develop his own AI model as easily, safely and less toxic as possible. Many tools such as Promptflow has been developed.

7.5 Microsoft Fabric



Bibliography

- [1] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms. *arXiv preprint arXiv:2305.14314*, 2023.
- [2] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models, 2021.
- [3] Zachary C. Lipton, John Berkowitz, and Charles Elkan. A critical review of recurrent neural networks for sequence learning, 2015.
- [4] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yurfeng Liu. Roformer: Enhanced transformer with rotary position embedding, 2022.
- [5] Ilya Sutskever, Oriol Vinyals, and Quoc V Le. Sequence to sequence learning with neural networks. *Advances in neural information processing systems*, 27, 2014.
- [6] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. Llama: Open and efficient foundation language models, 2023.
- [7] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2017.