Lu Sun, and many more.

A Notebook on Probability, Statistics, and Data Science

To my family, friends and communities members who have been dedicating to the presentation of this notebook, and to all students, researchers and faculty members who might find this notebook helpful.

Contents

Foreword

If software or e-books can be made completely open-source, why not a notebook?

This brings me back to the summer of 2009 when I started my third year as a high school student in Harbin No. 3 High School. In around the end of August when the results of Gaokao (National College Entrance Examination of China, annually held in July) are released, people from photocopy shops would start selling notebooks photocopies that they claim to be from the top scorers of the exam. Much curious as I was about what these notebooks look like, never have I expected myself to actually learn anything from them, mainly for the following three reasons.

First of all, some (in fact many) of these notebooks were more difficult to understand than the textbooks. I guess we cannot blame the top scorers for being so smart that they sometimes make things extremely brief or overwhelmingly complicated.

Secondly, why would I want to adapt to notebooks of others when I had my own notebooks which in my opinion should be just as good as theirs.

And lastly, as a student in the top-tier high school myself, I knew that the top scorers of the coming year would probably be a schoolmate or a classmate. Why would I want to pay that much money to a complete stranger in a photocopy shop for my friend's notebook, rather than requesting a copy from him or her directly?

However, things had changed after my becoming an undergraduate student in 2010. There were so many modules and materials to learn in a university, and as an unfortunate result, students were often distracted from digging deeply into a module (For those who were still able to do so, you have my highest respect). The situation became even worse as I started pursuing my Ph.D. in 2014. As I had to focus on specific research areas entirely, I could hardly split much time on other irrelevant but still important and interesting contents.

This motivated me to start reading and taking notebooks for selected books and articles, just to force myself to spent time learning new subjects out of my comfort zone. I used to take hand-written notebooks. My very first notebook was on *Numerical Analysis*, an entrance level module for engineering background graduate students. Till today I still have on my hand dozens of these notebooks. Eventually, one day it suddenly came to me: why not digitalize them, and make them accessible online and open-source, and let everyone read and edit it?

—

As most of the open-source software, this notebook (and it applies to the other notebooks in this series as well) does not come with any "warranty" of any kind, meaning that there is no guarantee for the statement and knowledge in this notebook to be absolutely correct as it is not peer reviewed. **Do NOT cite this notebook in your academic research paper or book!** Of course, if you find anything helpful with your research, please trace back to the origin of the citation and double confirm it yourself, then on top of that determine whether or not to use it in your research.

This notebook is suitable as:

- a quick reference guide;

- a brief introduction for beginners of the module;

- a "cheat sheet" for students to prepare for the exam (Don't bring it to the exam unless it is allowed by your lecturer!) or for lecturers to prepare the teaching materials.

This notebook is NOT suitable as:

- a direct research reference;

- a replacement to the textbook;

because as explained the notebook is NOT peer reviewed and it is meant to be simple and easy to read. It is not necessary brief, but all the tedious explanation and derivation, if any, shall be "fold into appendix" and a reader can easily skip those things without any interruption to the reading experience.

—

Although this notebook is open-source, the reference materials of this notebook, including textbooks, journal papers, conference proceedings, etc., may not be open-source. Very likely many of these reference materials are licensed or copyrighted. Please legitimately access these materials and properly use them.

Some of the figures in this notebook is drawn using Excalidraw, a very interesting tool for machine to emulate hand-writing. The Excalidraw project can be found in GitHub, *excalidraw/excalidraw.*

Preface

This notebook introduces probability, statistics, data science and engineering. They are the "must-have" ability in most, if not all, academic and industrial projects.

In Part I of the notebook, probability theory is introduced. Probability theory studies how likely an event is to occur. It offers rich models and tools to describe random values and stochastic events.

In Part II of the notebook, statistics is introduced. Statistics is a collection of methods to analyze and observe insights from data, verify statistics hypothesis and draw conclusions and predictions.

In Part III of the notebook, commonly used software and toolkits for statistics analysis and data science are introduced. Different from Parts I and II which focus more on theory, Part III focuses more on the tools to solve practical problems. Notice that Artificial Intelligence (AI) has become increasingly popular in recent years for data analysis. There is a separate notebook on introducing AI. In this notebook, the tools to process data using AI is only briefly introduced, again focusing only on the usage of software, not theory.

As a bonus, in Part IV, semantic web, the database framework defined under Web 3.0, is introduced. Semantic web does not necessarily contribute to solving a specific problem using a specific statistics theory. It is rather a concept or organizing and exchanging data efficiently.

Key references of this notebook are summarized as follows. Notice that these materials are so very widely used here and there in the entire notebook that it becomes improbable to address them each time they are used.

For probability and statistics parts:

- Spiegel, Murray, John Schiller, and Alu Srinivasan. *Probability and statistics.* 2020.

- Dekking, Frederik Michel, et al., *A Modern Introduction to Probability and Statistics: Understanding why and how.* Vol. 488. London: Springer, 2005.

For data science part:

- Kirill Eremenko, *R Programming A-Z: R For Data Science With Real Exercises*, Udemy Course.

- Lakshmanan, Valliappa, Martin Görner, and Ryan Gillard. Practical Machine Learning for Computer Vision. " O'Reilly Media, Inc.", 2021.

- Jose Portilla, *Complete TensorFlow 2 and Keras Deep Learning Bootcamp*, Udemy

Online materials such as tutorials from YouTube, Bilibili, etc., are also used in forming this notebook. ChatGPT-4 is used as a consultant in forming this notebook.

List of Figures

List of Tables

Part I

Probability

1

Basics

CONTENTS

This chapter introduces the basic concepts, axioms, theorems, and fundamental calculations in probability theory.

1.1 Randomness and Stochasticity

People use *random* and *stochastic* to describe a situation or a model whose outcome is not precisely predictable.

By saying *randomness*, we often refer to the case where the value of a variable (as in "random variable") is not entirely predictable. The value is not known precisely or its outcome not predictable until it is measured. The frequency of its outcome may follow some pattern which can be described as the statistic property of the variable. An example of a random variable would be the result of tossing a coin which can be either head (described by $X = 1$) and tail (described by $X = 0$).

By saying *stochasticity*, we often refer to the case where the outcome of a process (as in "stochastic process") cannot be modeled or predicted precisely. This might be a result of incomplete modeling, or maybe there are some random variable generation built in to the process. As a result, the process becomes non-deterministic, i.e., using the model configuration and the inputs to the model alone cannot give its precise output. A stochastic process can still

be described by a parametric model, just with some unknowns or disturbance coming into the model.

In this chapter, we will focus mostly on randomness. There are researches that try to handle stochastic process such as optimal stochastic control and robust control, etc.

1.1.1 Random Experiments

"Experiment" is one of the most important concepts in science and engineering. In many cases, experiments are used to verify a theory. In these cases, experiments are carefully designed so that its outcome is deterministic and predictable by the theory. By observing the results of the experiments matching the theory repeatedly, we build confidence in the theory.

However, there are other experiments where we do not have full control over the result due to the lack of information or incomplete modeling. Such experiments include tossing a coin, predicting the GDP of a country next year, etc. These experiments are known as *random experiments*.

1.1.2 Sample Space

A set S that consists of all possible outcomes of a random experiment is called a *sample space*.

If a sample space has a finite number of elements, it is called a *finite sample space*. If it has infinite elements as many as there are natural numbers (i.e., its elements can be mapped to $1, 2, 3, \ldots$), it is called a *countably infinite sample space*. If it has infinite elements as many as there are in some intervals on an axis, it is called a *noncountably infinite sample space*.

Infinity VS Infinity

Though both infinity, the total number of natural numbers, i.e., the cardinality of $\mathbb{N}$, is less than the total number of real numbers, i.e., the cardinality of $\mathbb{R}$.

The cardinality of $\mathbb{N}$ is known as $\aleph_0$. It is also the cardinality of all rational numbers. In computability theory, it is also the cardinality of all computable numbers (i.e., numbers that can be computed to within any desired precision by a finite terminating algorithm) and computable functions, i.e., algorithms. The total number of real numbers, on the other hand, is known as $\aleph_1$. It is also the cardinality of all irrational numbers and complex numbers, the number of points on an axis or axis interval, or a high dimensional space of $\mathbb{R}^n$ where n is a finite integer. Even larger infinity quantities, such as $\aleph_2, \aleph_3, \ldots, \aleph_\omega, \ldots$ are also defined, though they may not have an intuitive explanation as $\aleph_0$ and $\aleph_1$.

Deeper discussion of this topic requires advanced set theory, hence is not given here.

Finite sample space and countably infinite sample space are also known as *discrete sample space*, whereas noncountably infinite sample space, *nondiscrete sample space*.

1.1.3   Events

An event is a subset A of the sample space S, i.e., it is a set of possible outcomes. An event may or may not occur depending on the outcome of an experiment. In the special case where an event covers only a single point in S, it is also called an *elementary event*, or just a *sample*. The sample space S can also be taken as a event, which is known as the certain event.

1.2   Probability

Given an experiment setup and an event, there is always uncertainty as to whether the event will occur or not, and *probability* is a measurement of chance or likelihood that the event is going to occur. For example, by saying a probability of 100% or 1, we mean that we are certain that the event would occur.

1.2.1   Classical and Empirical Probability

There are different ways to calculate or estimate the probability of an event. In the *classical approach* where we know the total number of outcomes n, all of which have a equal chance of happening, and there are h ways for the event to occur, then the probability of the event is h/p. In the *frequency approach*, we can repeat the experiment n times where n is a large number, and record the number of instants where the event happened as h. The *empirical probability* of the event, which should be close to the actual probability of the event, is h/p.

1.2.2   Axioms

Both the classical approach and the frequency approach have some drawbacks. It is often difficult to define "equal chance" and "large number" in the two approaches respectively. Therefore, *axiomatic approach* is introduced as follows.

Let the sample space denoted by S, and events by A_i. For simplicity, assume that S is discrete, or at least part of S associated by the events is discrete. Define $P(\cdot)$ as the probability function, and $P(A)$ the probability of an event A, if the following axioms are satisfied:

1. For every event A_i, $P(A_i) \geq 0$.

2. For the certain event S, $P(S) = 1$.

3. For mutually exclusive events A_1 and A_2, $P(A_1 \cup A_2) = P(A_1) + P(A_2)$.

Using the above axioms, a bunch of well-known theorems can be derived, such as

- If $A_1 \subseteq A_2$, $P(A_2 - A_1) = P(A_2) - P(A_1)$.

- For every event A_i, $0 \leq P(A_i) \leq 1$.

- For the impossible event $\varnothing$, $P(\varnothing) = 0$.

- For the complement of an event A', $P(A') = 1 - P(A)$.

- For mutually exclusive events $A_i, i = 1, ..., n$, if $A = \bigcup_{i=1}^{n} A_i$, $P(A) = \sum_{i=1}^{n} P(A_i)$.

- For two events A_1 and A_2, $P(A_1 \cup A_2) = P(A_1) + P(A_2) - P(A_1 \cap A_2)$.

With the above axioms, we can revisit classical probability as follows.

Assume that a discrete and finite sample space S consists of the following elementary events (elementary events are always mutually exclusive) $A_i, i = 1, ..., n$, i.e.,

$$S = \bigcup_{i=1}^{n} A_i$$

and assume equal probabilities for all the elementary events, i.e., the probability of each event is given by

$$P(A_i) = \frac{1}{n}, i = 1, ..., n$$

Define an event A that is made up of h such elementary events out of A_i. The probability of A can then be calculated by

$$P(A) = \frac{h}{n}$$

where h, n are nothing but the cardinality of A and S with respect to the elementary events.

1.2.3 Conditional Probability

Assume two events A and B. The conditional probability $P(B|A)$ describes the probability of B given that A has occurred. The definition is given below. Think of this definition as S replaced by A since A has been confirmed occurred.

$$P(B|A) \equiv \frac{P(A \cap B)}{P(A)}$$

Or equivalently,

$$P(A \cap B) = P(A)P(B|A)$$

In the special case where $P(B|A) = P(B)$, or equivalently $P(A \cap B) = P(A)P(B)$, the two events A and B are said to be *independent events*.

From the above,

$$P(B|A) = \frac{P(B)P(A|B)}{P(A)}$$

which is known as the Bayes' rule.

1.3 Permutation and Combination

In classical probability, calculating the cardinalities of event A and sample space S is the key to solving the problem. Permutations and combinations are widely used for such calculation.

Suppose that there are n distinct objects, and we would like to select r objects from them, and arrange them in a particular order. The permutation

$$nPr = n(n-1)(n-2)\ldots(n-r+1) \tag{1.1}$$

gives the number of possible outcomes. In the special case $r = n$,

$$nPn = n(n-1)(n-2)\ldots 1$$

where $n(n-1)(n-2)\ldots 1$ is often denoted as $n!$. Therefore, (1.1) becomes

$$nPr = \frac{n!}{(n-r)!}$$

In permutation, the arrangement of the selected r objects matters. In the case where the order of the r objects in the outcome does not matter, combination

$$nCr = \frac{nPr}{r!}$$
$$= \frac{n!}{r!\,(n-r)!}$$

is used to calculate the total number of outcomes. Notice that nCr is also denoted by $\binom{n}{r}$.

2

Random Variables and Probability Distributions

CONTENTS

The outcome of an experiment can often be characterized by variables. For example, as introduced earlier, we can define X to be the outcome of tossing a coin, where $X = 1$ denotes head and $X = 0$ tail. Such variable is called a *random variable* A random variable is essentially a function whose output is an element in the sampling space. Therefore, a random variable is also called a *random function*. Since a random variable is often generated by a process, model or experiment, people also use *stochastic variable* or *stochastic function* to refer to a random variable.

2.1   Discrete and Continuous Random Variables

A random variable may be discrete or continuous depending on the sample space of the variable.

2.1.1   Discrete Random Variables

A random variable X may only take discrete values $x_1, x_2, ...,$ in which case X is called a discrete random variable.

    The probability of X taking a particular value x is denoted by $P(X =$

x). Sometimes for simplicity, it is simply denoted as $f(x)$ when there is no ambiguities. In this case, $P(X = x)$ or $f(x)$ is called the *probability function* of X.

Furthermore, define $P(X \leq x)$ or $F(x)$ as the *cumulative distribution function*. It is easy to prove that $F(x)$ is nondecreasing, and $\lim_{x\to 0} F(x) = 0$, $\lim_{x\to\infty} F(x) = 1$. Also, $F(x)$ "jumps" at each $P(X = x_k) > 0$ and it is continuous from the right.

2.1.2   Continuous Random Variables

A random variable X may also take continuous values in many applications. For example, let X denotes the total time consumption to finish the homework by a primary school student, which is likely to take any value between 0 minutes (no homework) to 120 minutes (very heavy homework).

In this case, for X to take a precise value, say for the student to finish his homework using precisely 25 minutes 13 seconds and 750 milliseconds, is very small (in fact, zero). The probability function $P(X = x)$ becomes useless. The cumulative distribution function $F(x) = P(X \leq x)$ still makes sense, as it essentially calculates the probability of X not given by a precise value, but within a range.

Inspired by this, define *probability density function (PDF)* or $f(x)$ for continuous random variable as follows. The probability density function $f(x)$ is such that

$$P(a < X < b) \quad = \quad \int_a^b f(x)dx$$

Or equivalently,

$$\begin{aligned} F(x) \quad &= \quad P(X \leq x) \\ &= \quad \int_{-\infty}^x f(\epsilon)d\epsilon \end{aligned}$$

Notice that $f(x)$ itself is not probability. It is $f(x)dx$ that accumulates in range $x \in (a, b)$ and forms the probability, hence the name "probability density".

In science and engineering, people are often more interested in continuous random variables than discrete random variables. This is because physical signals and measurements often come in the form of continuous signals. Therefore, the main focus of this notebook from this point on would be continuous signals. Discrete random variables can sometimes be taken as a special case of continuous random variables with impulse PDF.

2.2 Joint Distribution

Joint distribution is used to study the correlation of multiple random variables. It is widely used in system identification where insights of unknown variables are obtained by observing the known variables, assuming their correlation is known.

2.2.1 Joint Probability Density Function

Consider two random variables X and Y. In the case of discrete variables, define joint probability function and joint cumulative distribution function as follows.

$$\begin{aligned}
f(x,y) &= P\left(X = x, Y = y\right) \\
F(x,y) &= P\left(X \leq x, Y \leq y\right) \\
&= \sum_{u \leq x} \sum_{v \leq y} f(u,v)
\end{aligned}$$

In the case of continuous random variables, joint PDF of X and Y is such that

$$\int_{x=a}^{b} \int_{y=c}^{d} f(x,y) dx dy = P\left(a < X < b, c < y < d\right) \qquad (2.1)$$

Integrate (2.1) w.r.t an axis gives the PDF of the other variable. For example,

$$\int_{y=-\infty}^{\infty} f(x,y) = f_X(x) \qquad (2.2)$$

which becomes a function of a single variable, x. Here $f_X(x)$ is nothing but the PDF of X, completely ignoring the other variable Y.

We can define cumulative distribution function for joint distribution likewise as follows.

$$\begin{aligned}
F(x,y) &= P(X \leq x, Y \leq y) \\
&= \int_{u=-\infty}^{x} \int_{v=-\infty}^{y} f(u,v) du dv \\
F_X(x) &= P(X \leq x) \\
&= \int_{u=-\infty}^{x} \int_{v=-\infty}^{\infty} f(u,v) du dv
\end{aligned}$$

2.2.2 Conditional Distribution

One may get confused by (2.2) with the conditional distribution, in which case the values of some of the variables in the joint distribution is known and based on which the PDFs of other variables need to be calculated. For example, from (2.1), let Y be measured by a fixed value y, and we would like to calculate $f_{X|Y}(x|Y = y)$, i.e., the PDF of X given the known information $Y = y$. In many literatures, $f_{X|Y}(x|Y = y)$ is denoted by $f_{X|Y}(x|y)$ for simplicity.

The conditional PDF $f_{X|Y}(x|y)$ can be obtained as follows. It is essentially Bayes' rule applied on continuous variables.

$$f_{X|Y}(x|y) \quad = \quad \frac{f(x,y)}{f_Y(y)} \tag{2.3}$$

Equation (2.3) is given as an analytical equation of both x and y. Just substitute the measured value of y into (2.3) to get a function of x alone. Since (2.3) itself is a PDF of X, its integration w.r.t x is certainly 1. This can be easily verified as follows.

$$
\begin{aligned}
\int_{x=-\infty}^{\infty} f_{X|Y}(x|y)dx &= \int_{x=-\infty}^{\infty} \frac{f(x,y)}{f_Y(y)}dx \\
&= \frac{\int_{x=-\infty}^{\infty} f(x,y)dx}{f_Y(y)} \\
&= \frac{f_Y(y)}{f_Y(y)} \\
&= 1
\end{aligned}
$$

for any y. Notice that given a particular value of y, $f_Y(y)$ is a constant value independent of x, hence can be taken out from the integration in the intermediate derivation.

If X and Y are independent variables, $f(x,y) = f_X(x)f_Y(y)$. In this case, (2.3) becomes

$$f_{X|Y}(x|y) \quad = \quad f_X(x)$$

which implies that the information of $Y = y$ does not affect our understanding of X, just as if the information is completely ignored.

2.2.3 Parameter Estimation with Conditional Distribution

Equation (2.3) is widely used in parameter estimation. Let θ be the parameter(s) to be estimated, and x the measurement(s) that reflects θ via measurement model $f(\theta, x)$ which is given in the form of joint distribution of θ and X. The model is known as priori knowledge. Notice that parameters and measurements are correlated via joint PDF but not deterministic due to measurement noise and other uncertain factors.

Given measurement(s) $X = x$, the posteriori estimation of θ (in the form of PDF) can be obtained as follows. From (2.3),

$$f_{\theta|X}(\theta|x) \quad = \quad \frac{f_{X|\theta}(x|\theta)f_\theta(\theta)}{f_X(x)} \tag{2.4}$$

where on the right side $f_{X|\theta}(x)f_\theta(\theta)$ is known as the *likelihood function* that describes the likelihood of measuring $X = x$ if the actual parameter(s) is θ, $f_\theta(\theta)$ the *priori* knowledge of the distribution of θ, and finally $f_X(x)$ which is known as *evidence*, in this case a constant known value given $X = x$.

Equation (2.4) can be memorized as follows.

$$\text{posteriori} \quad = \quad \frac{\text{likelihood} \times \text{priori}}{\text{evidence}}$$

An estimation of θ can then be obtained by simply calculate its expectation as follows

$$\mathrm{E}(\theta) \quad = \quad \int_{-\infty}^{\infty} f_{\theta|X}(\theta|x)d\theta$$

Terms *priori* and *prior*, *posteriori* and *posterior* can be used interchangeably, respectively.

2.2.4 Geometric Probability

Geometric probability is an extension of classical probability to continuous random variables. Similar with classical probability where it is prerequisite that all elementary events share the same probability, in geometric probability it is assumed that all random variables (usually 2 or 3 variables in total) are uncorrelated and uniformly distributed. The probability of an event happening then can be obtained by measuring the area or volume of the event.

Theory wise, geometric probability is just one of the many joint PDF applications that is simple, intuitive and useful. An example is used to illustrate the use of geometric probability.

Geometric Probability Example

Consider two people trying to meet up at the park, but they forgot to tell each other the time to meet. Both of them will arrive at the park at anytime between $8:00$ AM and $9:00$ AM with equal chance, and wait for the other for 15 minutes or until $9:00$ AM, whichever is earlier, and then will leave the park if the other does not show up.

Calculate the chance of the two people successfully meeting up.

Draw the sample space in a 2-D plot as shown in Fig. 2.1, where the x-axis and y-axis are the arriving time of the two people, respectively. The shaded
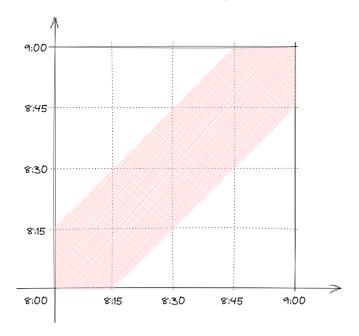
FIGURE 2.1
Sample space of two people arriving at part from 8:00 AM to 9:00 AM.

area represents the samples where the two would meet up successfully. Divide the shared area by the total sample space area to get $P = 7/16 = 43.75\%$ which is the probability that the two would meet up successfully.

2.3   Probability Density Distribution of Functions

Let X be a random variable with PDF $f_X(x)$. Let U be another random variable which is a function of X via $U = \phi(X)$. The PDF of U can be calculated as follows. For simplicity, assume that $U = \phi(X)$ is a injective function (one-to-one function), and $X = \phi^{-1}(U) = \psi(U)$. In that case, the PDF of U, $g(u)$, can be obtained as follows.

$$g(u) \quad = \quad |\psi'(u)| \, f\left(\psi(u)\right)$$

For example, let $U = aX$, $X = \frac{U}{a}$.

$$g(u) \quad = \quad \frac{1}{a} f\left(\frac{u}{a}\right)$$

Let X, Y be two random variables with joint distribution $f(x, y)$. Let

$U = X + Y$. The PDF of U can be obtained as follows.

$$g(u) \quad = \quad \int_{-\infty}^{\infty} f(x, u - x) dx \qquad (2.5)$$

In the special case where X and Y are independent, $f(x, y) = f_X(x) f_Y(y)$, and (2.5) becomes

$$\begin{aligned} g(u) \quad &= \quad \int_{-\infty}^{\infty} f_X(x) f_Y(u - x) dx \\ &= \quad f_X * f_Y \end{aligned}$$

where $*$ denotes the convolution operator.

3

Measures of Distribution

CONTENTS

Given the probability or probability density distribution of a random variable, a lot of information can be abstracted. For example, the expectation describes the average value of that random variable if the experiment is repeated many times.

Measures of distribution such as expectation and variance are introduced in this chapter.

3.1 Expectation

In probability and statistics sense, *expectation* (also known as *mean*) describes the average value of a random variable, if the variable is generated many times. For discrete random variable X, the expectation is given below.

$$\mathrm{E}(X) \quad = \quad \sum_{i=1}^{n} x_i P(x_i) \tag{3.1}$$

where $\mathrm{E}(\cdot)$ is used to denote the expectation, and n the cardinality of the sample space. For continuous random variable, it is

$$\mathrm{E}(X) \quad = \quad \int_{-\infty}^{\infty} x f(x) dx \tag{3.2}$$

Expectation is sometimes denoted as μ in literatures.

Some features of expectation calculation are given below.

$$\mathrm{E}(cX) \quad = \quad c\mathrm{E}(X)$$
$$\mathrm{E}(X+Y) \quad = \quad \mathrm{E}(X) + \mathrm{E}(Y)$$

where c is a constant and X, Y are two random variables. Furthermore, if X, Y are independent,

$$\mathrm{E}(XY) \quad = \quad \mathrm{E}(X)\mathrm{E}(Y)$$

These features can be easily derived from (3.1) and (3.2).

3.2   Variance and Standard Deviation

Variance and standard deviation describe how spread samples are from its expectation. It is defined as follows.

$$
\begin{aligned}
\mathrm{Var}(X) \quad &= \quad \mathrm{E}\left((X - \mathrm{E}(X))^2\right) \\
&= \quad \mathrm{E}\left(X^2 - 2X\mathrm{E}(X) + \mathrm{E}(X)^2\right) \\
&= \quad \mathrm{E}\left(X^2\right) - \mathrm{E}(X)^2 \\
\sigma_X \quad &= \quad \sqrt{\mathrm{Var}(X)}
\end{aligned}
\tag{3.3}
$$

Variance and standard deviation are sometimes denoted as σ^2 and σ respectively.

For zero-mean continuous random variables, from (3.3) the variance is

$$\mathrm{Var}(X) \quad = \quad \int_{-\infty}^{\infty} x^2 f(x) dx$$

This implies that for a non-bias estimation, the Mean Square Error (MSE) of the estimation is identical to the variance of the estimation.

Some features of variance calculation are given below.

$$\mathrm{Var}(cX) \quad = \quad c^2 \mathrm{Var}(X) \tag{3.4}$$
$$\tag{3.5}$$

For independent random variables X and Y,

$$\mathrm{Var}(X \pm Y) \quad = \quad \mathrm{Var}(X) + \mathrm{Var}(Y)$$

Mean and standard deviation can be used to standardize a ransom variable as follows.

$$X^* \quad = \quad \frac{X - \mu}{\sigma}$$

where μ, σ are the mean and standard deviation of random variable X respectively. The standardized random variable, X^*, has a mean of 0 and standard deviation of 1.

3.3   Moments

In mathematics, the r-th moment of a continuous function $f(x)$ about c is defined as follows.

$$\mu_n = \int_{-\infty}^{\infty} (x-c)^n f(x) dx$$

By simply saying *moment* without further explanation of c, it is by default that $c = 0$. Let $f(x)$ be a PDF. In this sense, the 0-th order and the 1-st order moment of a probability density function can be calculated as follows.

$$\mu_0 = \int_{-\infty}^{\infty} f(x) dx = 1$$

$$\mu_1 = \int_{-\infty}^{\infty} x f(x) dx = \mathrm{E}(X)$$

where it can be seen that the 0-th and 1-st moment of a PDF is 1 and its mean, respectively.

Further more, let $c = \mu_1$ be the mean of the random variable to calculate the second *central moment* μ_2 as follows.

$$\mu_2 = \int_{-\infty}^{\infty} (x - \mu_1)^2 f(x) dx = \mathrm{Var}(X)$$

which is the variance of the random variable.

Using mean and variance to further define *standardized moments* as follows. For $k \geq 3$,

$$\bar{\mu}_k = \frac{\mu_k}{\sigma^k}$$

$$\mu_k = \mathrm{E}\left((X - \mu)\right) = \int_{-\infty}^{\infty} (x - \mu)^k f(x) dx$$

$$\sigma^k = (\mu_2)^{\frac{k}{2}} = \left(\int_{-\infty}^{\infty} (x - \mu)^2 f(x) dx\right)^{\frac{k}{2}}$$

where μ_1 and μ_2 are the aforementioned mean (1-st order moment) and variance (2-nd order central moment) of the random variable, respectively.

The 3-rd and 4-th order standardized moments are known as the *skewness* and *kurtosis* of the PDF, respectively, as shown below.

$$\bar{\mu}_3 \equiv \gamma_1 = \frac{\mathrm{E}(x - \mu)^3}{\sigma^3}$$

$$\bar{\mu}_4 \equiv \gamma_2 = \frac{\mathrm{E}(x - \mu)^4}{\sigma^4}$$
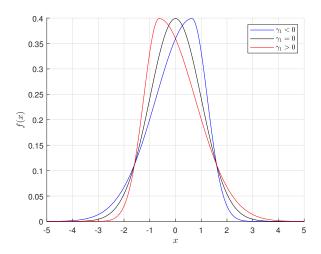
FIGURE 3.1
Demonstration of PDF with different skewness.

where $\mu = \mu_1$ and $\sigma = \sqrt{\mu_2}$ are the mean and standard deviation of the random variable, respectively.

The skewness γ_1 is a measure of asymmetry of the PDF. When $\gamma_1 > 0$ or positive skew, the distribution has a long tail on the right side of the PDF. When $\gamma_1 < 0$ or negative skew, the distribution has a long tail on the left side. When $\gamma_1 = 0$, the PDF might be symmetric (but not necessarily so). Examples are given in Fig. 3.1.

The kurtosis γ_2 measures the "tailedness" of a probability distribution, i.e., whether the PDF has heavy tail or thin tail. However, the kurtosis defined by the fourth standardized moment is often adjusted by subtracting 3 (which is the kurtosis of a standard normal distribution) to make it zero for a normal distribution.

3.4   Correlation

"nobreak

3.5   Important Rules

Part II

Statistics

4

Sampling

CONTENTS

The data used for statistics analysis is usually a small portion collected from a huge group. The collected data is called a sample, and the huge group from which the sample is collected, the population. For example, the incomes of a random 1000 citizens as a sample may reflect the incomes of the entire city. Another example is that the working efficiency of a machine for the past months as a sample may reflect its working efficiency in its entire lifespan.

An important underlying assumption behind statistics is that the insights observed from the sample also apply to the population. With that regard, questions naturally come up. Why does this assumption hold? When does this assumption hold? In what extend does this assumption hold? What can be done to make the sample more effective and efficient when reflecting the population?

This chapter tries to answer the above questions. As it is introduced later, the inference from sample to population is not certain, as there is a chance (however small it might be) that the samples are completely biased from the population. Therefore, we must use probability in any statement drawn from statistics analysis.

4.1 Sampling Methods

When selecting elements from the population, make sure that all elements have a equal probability of being selected, hence, random sampling. Depending on how many times a member can be sampled, we have

- Sampling with replacement: a member can be chosen more than once.

- Sampling without replacement: a member can be chosen no more than once.

Sometimes it is interesting to compare the differences of the two methods, especially then the population is finite. An obvious difference is that by using sampling with replacement all the samples can be considered as "independent event", while by using sampling without replacement, previous samples may change the distributions in the remaining population, thus making the samples relevant. In this case, using sampling with replacement can theoretically be considered as sampling from an infinite population (by thinking that the population is duplicated as many times as necessary).

In practice, the population is usually so large, that sampling from a finite population can be considered as sampling from an infinite population, and the two methods would make no differences as far as it is concerned.

Consider the following examples. A set of N random variables are generated from a Gaussian distribution as the population. Sample the population M times using sampling with replacement and sampling without replacement, respectively. Calculate the sampled mean and variance after each sampling instance, and see how it converges to the mean and variance of the population.

In the first example, let $N = 100$ and $M = 500$. Figures 4.1 and 4.2 gives the cumulative mean and variance of sampling with and without replacement, respectively. The mean and variance are given by red and blue curves, respectively. The statistics obtained from the cumulative samples and from the population are given by the solid and dashed curves, respectively. Notice that in Fig. 4.2, after number of samples exceeding 100, the entire population has been sampled, and thus the sampling stops. This explains why its mean and variance stop fluctuating and converge to the population mean and variance, respectively.

In practice, however, the population size is often orders of magnitudes larger than the number of samples. In the second example, let $N = 10000$ and $M = 500$. The corresponding figures are given in Figs. 4.3 and 4.4. There is no obvious differences of the two figures from statistics perspective.

4.2   Model of Population

The features of the population is often not known, or at least not known entirely. It is possible to make some preliminary assumptions to the distribution of population, with parameters to be further confirmed using the samples.

For example, let X be a variable of the population. It could be, for example, the heights of all teenagers in a city. We can make an assumption that X follows some distribution $f(x)$. A widely used assumption, in this scenario, is that $f(x)$ is a Gaussian distribution with mean μ and standard deviation σ, and each element in the population, X_i, can be taken as a random variable generated from $f(x)$. In the case of Gaussian distribution, since it is uniquely
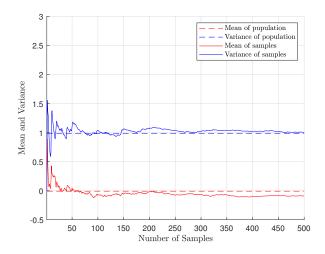
FIGURE 4.1
Sample with replacement, $N = 100$, $M = 500$.
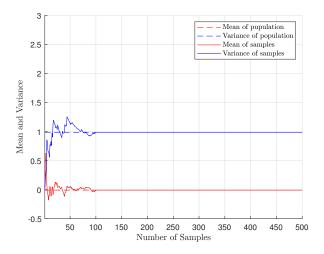
FIGURE 4.2
Sample without replacement, $N = 100$, $M = 500$.
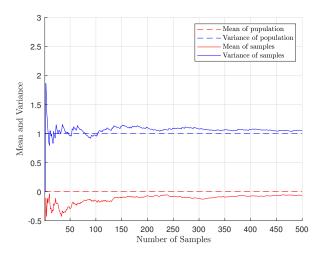
FIGURE 4.3
Sample with replacement, $N = 10000$, $M = 500$.
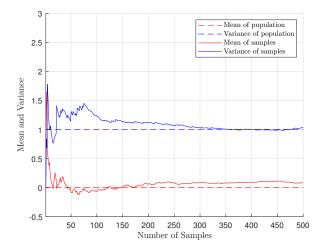
FIGURE 4.4
Sample without replacement, $N = 10000$, $M = 500$.

characterized by μ and σ, other quantities such as the median, moments, skewness, etc., can be derived once μ and σ is calibrated.

The questions rise sequentially are:

- What are the parameters in the assumed distribution?

- Does it indeed follow the assumed distribution?

The answers to the above questions need to be found out via the samples, or more precisely, from the sample statistics.

4.3   Sample Statistics

5

Estimation

CONTENTS

6

Regression

CONTENTS

7

Hypothesis Testing

CONTENTS

8

Bayesian Methods

CONTENTS

Part III

Tools

9

R Basics

CONTENTS

This chapter introduces R language basics, including the import (from files or databases), tidying, transformation, modeling, verification, and visualization of data. The use of R language in machine learning is also briefly covered.

Given that R is less popular than the more well-known and widely appreciated Python and MATLAB among engineering background researchers, this chapter serves as a basic introduction to R, before digging into its use in data science.

The first part of the chapter from Section 9.1 to Section 9.6 introduces

the basic grammar and visualization tools, and the second part of the chapter from Section 9.7 onwards more advanced practice of R.

9.1 R and RStudio Installation

R is a programming language for statistical computing and visualization. It is widely used among statisticians and data miners for developing statistical software and carrying out data analysis. R is free and can be downloaded from [2], where more details about R can also be found.

RStudio, also known as *Posit*, is an IDE widely used for R programming and testing. RStudio IDE is open-source and free of charge for personal use. It can be downloaded from [3].

Download R and RStudio from the aforementioned web sites, and install them sequentially.

9.2 R Packages Management

Before introducing the syntax, libraries, data frames and tools of R, it is worth introducing package management methods in R.

R packages, both built-in and third-party, provide power functions, data, and compiled codes for data analysis and visualization in a well-defined format. The packages can be published and shared online. CRAN is by far the most popular platform to store and share R packages.

To install or remove a package, use

```
install.packages("<package>")
remove.packages("<package>")
```

respectively. For example,

```
install.packages("pacman")
```

To load a package, use

```
library(<package>)
```

for example

```
library(pacman)
```

After loading a package, the data frames and functions defined in that package can be used normally. Otherwise, to refer to a data frame or a function, the package name has to be used as a prefix as `<package>::<resource>`, which is inconvenient if the resource is used frequently.

To unload a package, use

```
detach("package:<package>", unload = TRUE)
```

for example

```
detach("package:pacman", unload = TRUE)
```

The above methods work for both built-in packages (which often does not require installation) and third-party packages.

There are third-party packages that provides package management functions. The package `pacman` is an example of such package. With `pacman` installed and loaded, use the following commands to install, load and unload packages respectively.

```
p_install(<package>, ...) # install
p_load(<package>, ...) # install and load
p_unload(<package>, ...) # unload
p_unload(all) # unload all
```

An example of using `pacman` to load packages are given as follows.

```
pacman::p_load(
pacman, # package management
dplyr, # data manipulation
GGally, # data visualization
ggplot2, # data visualization
ggthemes, # data visualization
ggvis, # data visualization
httr, # url and http
lubridate, # date and time manipulation
plotly, # data visualization
rio, # io
rmarkdown, # documentation
shiny, # web apps development
stringr, # string operation
tidyr # data tidying
)
```

where notice that the above command can be executed before the loading of `pacman` itself, which is the reason `pacman::p_load()` prefix is used. These packages are commonly used in R projects. A brief explanation to them are given as comments following `#`. Notice that the first time installation of all the above packages may take a few minutes.

RStudio provides a graphical interface to manage packages as shown in Fig. 9.1.

FIGURE 9.1
Graphical interface to manage packages provided by RStudio.

9.3   R Programming Basics

This section introduces the basics of R programming, including the data types and syntax for basic R commands.

As the fundamentals, it is worth introducing here that R is case sensitive. Use # to lead a comment in R. Use print() to print a variable on the console. Typing the name of a variable often also prints it out, with a few exceptions such as when in a loop. Finally, use ? followed by a function or a data frame name to check the help document for that function or data frame.

9.3.1   Data Types

R provides many data types. Commonly used data types are summarized in Table 9.1, where notice that <- is used to assign a value to a variable. Use typeof() to check the type of a variable. Alternatively, use is.numeric(), is.integer(), is.double(), is.character(), etc., to check whether a variable belongs to a particular data type.

Examples of assigning variables and checking their types are given as follows.

```
> n <- 2L
> typeof(n)
```

TABLE 9.1

Commonly used data types.

Data Type	Syntax (Example)	Description
integer	n <- 2L	An integer. Define an integer by a value followed by L.
double	x <- 2	An double float value.
complex	z <- 3+2i	A complex value.
character	a <- "a"	A character or a string.
logical	q <- T	A boolean value. Use T, TRUE and F, FALSE to represent true and false repectively.

```
[1] "integer"

> x <- 2
> typeof(x)
[1] "double"

> z <- 3+2i
> typeof(z)
[1] "complex"

> a <- "h"
> typeof(a)
[1] "character"

> q <- T
> typeof(q)
[1] "logical"
```

To transform data from one type to another, use `as.<data-type>()`. Examples of transforming data types are given as follows.

```
> n1 <- as.integer(2)
> typeof(n1)
[1] "integer"

> n2 <- as.integer("2")
> typeof(n2)
[1] "integer"

> x1 <- as.double(2L)
> typeof(x1)
[1] "double"

> x2 <- as.double("2")
> typeof(x2)
```

```
[1] "double"

> z1 <- as.complex("3+2i")
> typeof(z1)
[1] "complex"

> a1 <- as.character(2L)
> typeof(a1)
[1] "character"

> a2 <- as.character(2)
> typeof(a2)
[1] "character"
```

R supports arithmetic calculations of variables, including +, -, *, /, %/% (integer division), %% (modulus) and ^ exponential. Examples of arithmetic calculations are given as follows.

```
> a <- 16
> b <- 3
> add <- a + b
> sub <- a - b
> multi <- a * b
> division <- a / b
> int_division <- a %/% b
> modulus <- a %% b
> exponent <- a ^ b
> add
[1] 19
> sub
[1] 13
> multi
[1] 48
> division
[1] 5.333333
> int_division
[1] 5
> modulus
[1] 1
> exponent
[1] 4096
```

R supports built-in and third-party functions which extend the capability of data manipulation. There is a rich set of functions for numerical calculations, string operations, probability density calculations and statistics analysis. Some of them are summarized in Tables 9.2, 9.3, 9.4, 9.5 and 9.6.

9.3.2   Conditionals and Loops

The if statement syntax is given as follows.

TABLE 9.2

Numerical calculations.

Syntax (Example)	Description
abs(x)	Absolute value.
sqrt(x)	Square root.
ceiling(x)	Smallest larger/equal integer.
floor(x)	Largest smaller/equal integer.
trunc(x)	Integer part of a variable.
round(x, n=0)	Round to n digit after decimal.
sin(x)	Trigonometric sin function.
cos(x)	Trigonometric cos function.
tan(x)	Trigonometric tan function.
log(x)	Natural logarithm.
log10(x)	Common logarithm.
exp(x)	Exponent.

TABLE 9.3

Logical comparisons.

Syntax (Example)	Description
x == y	Equal.
x != y	Not equal.
x > y, x < y	Greater than; less than.
x >= y, x <= y	Greater than or equal to; less than or equal to.
!x	Not.
x & y	And.
x | y	Or.
isTRUE(x)	Is true.

TABLE 9.4

String operations.

Syntax (Example)	Description
substr(s, n1, n2)	Segment of a string, from the n_1-th character to n_2-th character, both characters included.
grep(p, s)	Searching of a pattern in a string.
sub(s1, s2, s)	Find and replace patterns in a string.
paste(s1, s2, ..., p="")	Concatenate strings with selected pattern to separate them.
strsplit(s, p)	Split string into multiple strings at selected split points.
tolower(s)	Convert to lower case.
toupper(s)	Convert to upper case.

TABLE 9.5

Probability density related operations.

Syntax (Example)	Description
dnorm(x, m=0, std=1)	Calculate the PDF of Gaussian distribution.
pnorm(x, m=0, std=1)	Calculate the CDF of Gaussian distribution.
qnorm(x, m=0, std=1)	Inverse function of pnorm().
rnorm(n, m=0, std=1)	Generate Gaussian distribution samples.
dbinom(n, size, p)	Calculate the probability of a binominal distribution.
pbinom(n, size, p)	Calculate the comulative probability of a binominal distribution.
qbinom(x, size, p)	Inverse function of pbinom().
rbinom(n, size, p)	Generate binominal distribution samples.
dpois(x, lambda)	Calculate the probability of a Poisson distribution.
ppois(n, lambda)	Calculate the comulative probability of a Poisson distribution.
qpois(x, lambda)	Inverse function of ppois().
rpois(n, lambda)	Generate Poisson distribution samples.
dunif(x, min=0, max=1)	Calculate the PDF of uniform distribution.
punif(x, min=0, max=1)	Calculate the CDF of uniform distribution.
qunif(x, min=0, max=1)	Inverse function of punif().
runif(n, min=0, max=1)	Generate uniform distribution samples.

TABLE 9.6

Aggregate and statistics functions.

Syntax (Example)	Description
mean(l)	Mean.
sd(l)	Standard deviation.
median(l)	Median.
range(l)	Minimum and maximum.
min(l)	Minimum.
max(l)	Maximum.
sum(l)	Sum

```
if(<condition>){
        <command>
} else if(<condition>){
        <command>
} else{
        <command>
}
```

An example of using `if` statement is givne below.

```
> x <- rnorm(1)
> if(x > 0){
        +   y <- x
        + } else if(x < 0){
        +   y <- -x
        + } else{
        +   y <- 0
        + }
> print(x)
[1] -1.981445
> print(y)
[1] 1.981445
```

The `for` loop syntax is given as follows.

```
for(<variable> in <vector>){
        <command>
}
```

where the `<vector>` can be a list of not only numbers but also characters. Examples of using `for` loop are given below.

```
> for(i in 1:5){
        +     print(i)
        + }
[1] 1
[1] 2
[1] 3
[1] 4
[1] 5

> for(i in c("a", "b", "c")){
        +     print(i)
        + }
[1] "a"
[1] "b"
[1] "c"
```

The `while` loop syntax is given as follows.

```
while(<condition>){
    <command>
}
```

An example of using `while` loop is given below.

```
> counter <- 0
> while(counter < 5){
+     print(counter)
+     counter <- counter + 1
+ }
[1] 0
[1] 1
[1] 2
[1] 3
[1] 4
```

An example of using the above to verify the law of the large number is given below.

```
> N <- 1000000
> counter <- 0
> for(i in rnorm(N)){
      +   if(-1 < i & i < 1){
              +     counter <- counter + 1
              +   }
      + }
> answer <- counter / N # expected: 0.682689492
> print(answer)
[1] 0.682407
```

9.3.3   User-Defined Functions

Define a simple function as follows. Note to run the codes where the function is described before calling the function.

```
<function-name> <- function(<input>[=<default>], ...){
      <command>
      [return(<variable>)]
}
```

9.3.4   Vectors

There are different types of vectors in R. The commonly used vector types include numeric vector (including both double and integer vector) and character vector. All elements in a vector must have the same data type. When different data type values are stored in a vector, they will be transferred to the most general data type. A single number or character is stored as a vector of length 1.

Notice that the index of a vector in R starts from 1 instead of 0. This is different with many other computer languages.

Use the following syntax to create a vector.

```
<vector> <- c(<value>, ...)
```

where `<value>` can be single element or a vector. For example,

```
> l <- c(1,2,3,4,5)
> print(l)
[1] 1 2 3 4 5
> typeof(l)
[1] "double"
```

Alternative ways to create a vector are given as follows. Use sequence to create a vector as follows.

```
<vector> <- seq(<from>, <to>, <by=1>)
<vector> <- <from>:<to> # equivalent to seq() with by=1
```

Use replica to create a vector as follows.

```
<vector> <- rep(<value>, <repeate>)
```

where `<value>` can be a numeric number, a character, or a vector. For example,

```
> l <- rep(c("a", "b", "cde"), 2)
> print(l)
[1] "a"   "b"   "cde" "a"   "b"   "cde"
```

Replica can also be used to create empty vector vy `rep(NA, n)`.

A character vector can also be created by splitting strings using `strsplit()`. For example,

```
> a <- "Hello World!"
> b <- strsplit(a, "")
> print(b)
[[1]]
[1] "H" "e" "l" "l" "o" " " "W" "o" "r" "l" "d" "!"
```

To access the element in a vector, use `<vector>[<index>]`. Again, notice that the first element in a vector has the index of 1 instead of 0. The index can be an integer, or a list of integer such as a sequence. Examples are given below.

```
> s <- c("a", "b", "c", "d", "e", "f", "g")
> s[1]
[1] "a"
> s[7]
[1] "g"
> s[2:5]
[1] "b" "c" "d" "e"
> s[c(1L, 3L, 5L)] # s[c(1.1, 3.5, 5.9)] gives the same result; data
    type auto transferred
[1] "a" "c" "e"
```

Notice that accessing a single element in a vector is rarely used in practice, because most operations in R are done by the vector basis. Vectorization operation, also known as single-instruction-multiple-data operation, significantly speeds up the calculation in R, which is quite commonly seen in high-layer languages such as R and Python. This is due to the intepreting and wrapping techniques a high-layer language uses to communicate with the underlying low-layer languages, and also the support many processors have for parallel computing.

Most, if not all, of the numerical calculations, including +, -, *, /, %/%, %%, ^. Examples are given below.

```
> a <- c(1,2,3,4,5)
> b <- c(5,4,3,2,1)
> a + b
[1] 6 6 6 6 6
> a - b
[1] -4 -2 0  2  4
> a * b
[1] 5 8 9 8 5
> a / b
[1] 0.2 0.5 1.0 2.0 5.0
> a %/% b
[1] 0 0 1 2 5
> a %% b
[1] 1 2 0 0 0
> a ^ b
[1]  1 16 27 16 5
```

It is also possible to apply logic operations using vectors. Examples are given below.

```
> a <- c(1,2,3,4,5)
> b <- c(5,4,3,2,1)
> a < b
[1]  TRUE  TRUE FALSE FALSE FALSE
> a > b
[1] FALSE FALSE FALSE TRUE TRUE
> a == b
[1] FALSE FALSE TRUE FALSE FALSE
```

When the sizes of the vectors are not consistent, the shorter vector will repeat and populate to align with the longer vector. Examples are given below.

```
> a <- c(1,10)
> b <- c(1,2,3,4)
> a + b
[1]  2 12  4 14
> a - b
[1]  0  8 -2 6
> a * b
```

```
[1]  1 20  3 40
> a / b
[1] 1.0000000 5.0000000 0.3333333 2.5000000
> a %/% b
[1] 1 5 0 2
> a %% b
[1] 0 0 1 2
> a ^ b
[1]     1   100     1 10000
```

The vector can also play as the input argument or output return of a function, which can be difficult for some computer languages by nature. Again, a scalar is treated as a vector with length 1 in R.

An example of using the above to analyze the profit of a company is given below.

```
> revenue <- round(rnorm(12, 10000, 500), 2)
> expenses <- round(rnorm(12, 9500, 500), 2)
>
> # profit for each month
> profit.month <- revenue - expenses
> print(profit.month)
[1] -861.33 974.12 665.84 275.72 2374.99 1231.07 953.87 396.83 -536.62
    1202.72 529.28 522.62
> # profit after tax for each month (tax rate 30%)
> profit.month.aftertax <- 0.3*profit.month
> print(profit.month.aftertax)
[1] -258.399 292.236 199.752  82.716 712.497 369.321 286.161 119.049
    -160.986 360.816 158.784 156.786
> # profit margin for each month
> profit.month.margin <- round(100 * profit.month.aftertax / revenue,
    2)
> print(profit.month.margin)
[1] -2.85 2.69 1.97  0.84  6.88  3.40  2.99  1.16 -1.71 3.25  1.60  1.56
> # is good month
> profit.month.aftertax > mean(profit.month.aftertax)
[1] FALSE TRUE  TRUE FALSE TRUE  TRUE  TRUE FALSE FALSE TRUE FALSE FALSE
> # is bad month
> profit.month.aftertax < mean(profit.month.aftertax)
[1]  TRUE FALSE FALSE TRUE FALSE FALSE FALSE TRUE TRUE FALSE TRUE TRUE
> # is the best month
> profit.month.aftertax == max(profit.month.aftertax)
[1] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
    FALSE
> # is the worst month
> profit.month.aftertax == min(profit.month.aftertax)
[1]  TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
    FALSE
```

FIGURE 9.2
A demonstration of a matrix in R.

9.3.5 Matrices

A matrix in R is a recording of a table of data. Matrices are important because they are how data is often naturally organized, and they are also the build blocks of data frame in R.

A demonstration of a matrix in R is given by Fig. 9.2. Let the matrix be named `A`. The elements in the matrix can be accessed by the name of the table followed by the index coordinates. For example, in the figure, `A[1,1]` refers the first element and `A[4,6]` the last element.

It is possible to refer to an entire row or column. For example, use `A[1,]` to represent the first row of the matrix, by not specifying the column index. The same applies to the column.

Notice that all elements in a matrix must have the same data type.

A matrix can be created from scratch by stacking rows as follows. First, consider creating rows in the matrix. Then, use `rbind()` to bind rows. Finally, give names to each column and row.

```
# build rows
<row1> <- c(<value11>, ..., <value1n>)
...
<rowm> <- c(<valuem1>, ..., <valuemn>)
# build matrix
<matrix> <- rbind(<row1>, ..., <rowm>)
# (optional) clean rows
rm(<row1>, ..., <rowm>)
# give names
colnames(<matrix>) <- c("<column-name1>", ..., "<column-namen>")
rownames(<matrix>) <- c("<row-name1>", ..., "<row-namem>")
```

There are alternative ways, other than `rbind()`, to create a matrix. For example, `matrix()` convert a vector into a matrix. Similar with `rbind()`, `cbind()` binds the columns to form a matrix. Examples to create matrices using different methods are given below.

```
> A <- matrix(1:9, 3, 3)
> print(A)
[,1] [,2] [,3]
[1,]    1    4    7
[2,]    2    5    8
[3,]    3    6    9
> B <- rbind(c(1, 4, 7), c(2, 5, 8), c(3, 6, 9))
> print(B)
[,1] [,2] [,3]
[1,]    1    4    7
[2,]    2    5    8
[3,]    3    6    9
> C <- cbind(c(1, 2, 3), c(4, 5, 6), c(7, 8, 9))
> print(C)
[,1] [,2] [,3]
[1,]    1    4    7
[2,]    2    5    8
[3,]    3    6    9
```

The name of the columns and rows can also be used to access an element, just by replacing the index with the name (with quotation mark) of the associated column or row. The same applies to vectors, as they can be treated as a one dimensional matrix. More details about naming a vector and columns and rows of a matrix are illustrated as follows.

To check the names relevant to a matrix, use `names(<vector>)`, `rownames(<matrix>)` and `colnames(<matrix>)`, depending on dealing with either a vector or a matrix. These commands can also be used to assign names. Examples are given below.

```
> v <- c(1, 2, 3, 4, 5)
> names(v) <- c("e1", "e2", "e3", "e4", "e5")
> print(v)
e1 e2 e3 e4 e5
 1  2  3  4  5
> print(v[3])
e3
3
> print(v["e3"])
e3
3
> A <- matrix(1:9, 3, 3)
> colnames(A) <- c("col1", "col2", "col3")
> rownames(A) <- c("row1", "row2", "row3")
> print(A)
col1 col2 col3
```

```
row1   1   4   7
row2   2   5   8
row3   3   6   9
> print(A[2,2])
[1] 5
> print(A["row2", "col2"])
[1] 5
> print(A[2,])
col1 col2 col3
2    5    8
> print(A["row2",])
col1 col2 col3
2    5    8
> print(A[,2])
row1 row2 row3
4    5    6
> print(A[,"col2"])
row1 row2 row3
4    5    6
```

To remove the names, simply assign `NULL` to the name.

Like the vector, operators are defined in matrix level as well. For example, for two matrices with the same shape, numerical operations such as `+`, `-`, `*`, `/`, `%/%`, `%%` and `^` can be implemented.

R provides flexible and powerful data visualization tools, many of which more advanced than what is to be introduced in this section. This section introduces a simple matrix visualization function called `matplot()`, which plots the columns of a matrix against each other.

To demonstrate `matplot()`, consider the following example.

```
professor <- c(1130, 1026, 893, 922, 776)
student <- c(2, 14, 24, 49, 46)
citation <- rbind(professor, student)
colnames(citation) <- c("2018", "2019", "2020", "2021", "2022")
rownames(citation) <- c("Professor", "Student")
print(citation)
citation.ratio <- citation
citation.ratio["Professor",] <- round(citation["Professor",] / mean(
    citation["Professor",]) * 100, 1)
citation.ratio["Student",] <- round(citation["Student",] / mean(
    citation["Student",]) * 100, 1)
print(citation.ratio)
matplot(
        2018:2022, # x axis
        t(citation.ratio), # y axis
        type="b", # line and point selection
        pch = 15:16, # point shape
        col = 1:2, # color
        xlab = "Year",
```
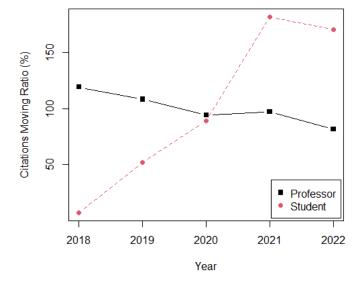
FIGURE 9.3
A demonstration of using `matplot` to plot trends.

```
        ylab = "Citations Moving Ratio (%)"
)
legend("bottomright", inset = 0.01, legend = rownames(citation.ratio),
    pch = 15:16, col = 1:2, horiz = F)
```

where `t()` used inside `matplot()` calculates the transpose of a matrix. Save
the above in a script and execute the code, to get the following Fig. 9.3.

Notice that `matplot()` is not widely used in particular in R.

As introduced earlier, a matrix or a vector can be split and segmented to
form a smaller matrix or vector. It is worth mentioning that when a single
column or row is selected, R will automatically treated the return as a vector
instead of a matrix. An example is given below. When a matrix downgrades
to a vector, the row name (if it has only one row), or the column name (if it
has only one column) will be removed.

```
> A <- matrix(1:9, 3, 3)
> is.matrix(A)
[1] TRUE
> is.vector(A)
[1] FALSE
> is.matrix(A[1,])
[1] FALSE
> is.vector(A[1,])
[1] TRUE
```

To get consistent results, when segmenting matrix to get a single row or

column vector, deliberately ask R to not drop the matrix dimensions. This can be done as follows. By doing this, the names assigned to columns and rows preserve.

```
> A <- matrix(1:9, 3, 3)
> is.matrix(A[1,,drop=F]) # select a row/column
[1] TRUE
> is.matrix(A[2,3,drop=F]) # select an element
[1] TRUE
```

An example of using the above to analyze the performance of players through a series of basketball games are given below.

```
# generate table
player_name <- c("player1", "player2", "player3")
match_name <- c("match1", "match2", "match3", "match4", "match5", "
    match6", "match7", "match8", "match9", "match10")
penalty_attempt <- abs(matrix(round(rnorm(3*10, 5, 2)), 3, 10))
penalty_point <- abs(penalty_attempt - matrix(abs(round(rnorm(3*10, 1,
    1))), 3, 10))
throw_attempt <- abs(matrix(round(rnorm(3*10, 15, 3)), 3, 10))
total_point <- abs(3*throw_attempt - abs(matrix(round(rnorm(3*10, 5, 1)
    ), 3, 10))) + penalty_point
rownames(penalty_attempt) <- player_name
colnames(penalty_attempt) <- match_name
rownames(penalty_point) <- player_name
colnames(penalty_point) <- match_name
rownames(throw_attempt) <- player_name
colnames(throw_attempt) <- match_name
rownames(total_point) <- player_name
colnames(total_point) <- match_name

# claim function
myplot <- function(table, xlab, ylab){
    row_name = rownames(table)
    column_name = colnames(table)
    matplot(
        1:length(column_name), # x axis
        t(table), # y axis
        type="b", # line and point selection
        pch = 1:length(row_name), # point shape
        col = 1:length(row_name), # color
        xlab = xlab,
        ylab = ylab
    )
    legend("bottomleft", inset = 0.01, legend = row_name, pch = 1:
        length(row_name), col = 1:length(row_name), horiz = F)
}

# plot
```
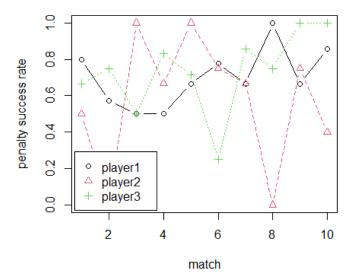
FIGURE 9.4
Plot of penalty success rate of the 3 players in 10 matches.

```
myplot(penalty_point / penalty_attempt, "match", "penalty success rate
    ") # penalty successful rate
myplot((total_point - penalty_point) / throw_attempt, "match", "average
     gained point per throw") # average point gained per throw
```

The results of the above codes are given in Figs. 9.4, 9.5.

9.4 Data Frames

Data frame, just like vector and matrix, is another data structure defined in R.

Both matrix and data frame use a table structure to store data, but data frame does not require all data to be with the same data type. Therefore, data frame is by nature the most closest format to represent a data structure from the real life. The aforementioned flexibility makes it maybe the most important and commonly used data structure in R.

In many applications and sample examples, data are stored and processed in data frame structure. When importing data from the real world, such as from a CSV file, the data is often read into a data frame before further processing.
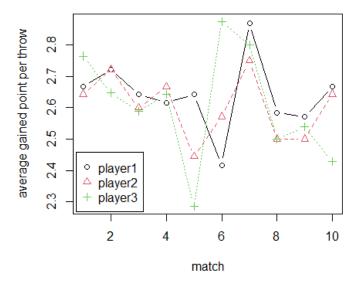
FIGURE 9.5
Plot of average point gained per throw attempt for the 3 players in 10 matches.

9.4.1 Data Import

One of the most common sources of data is CSV files. R provides convenient functions to read data from CSV files into data frames. Use the following commands to import data from a CSV file into a data frame.

The following command pops up a separate window that allows the user to choose a CSV file manually.

```
<data-frame> <- read.csv(file.choose()) # manual selection
```

The following commands import a specified CSV file.

```
setwd("<directory>") # navigate to the directory of the csv file
<data-frame> <- read.csv("<csv-file>.csv")
```

where notice that `getwd()` and `setwd()` are used to get and set current working directory, respectively.

9.4.2 Basic Operations

There are a few ways to access an element in a data frame. The methods used for accessing matrix element, including

```
<df>[<row-index>, <column-index>]
<df>[<row-index>, "<column-name>"]
```

still work fine. Do notice that different from a matrix, the rows in a data frame

TABLE 9.7

Commonly used commands for data frame exploration.

Syntax (Example)	Description
`nrow(df)`	Number of rows.
`ncol(df)`	Number of columns.
`head(df, n=6L)`	Display the first few rows.
`tail(df, n=6L)`	Display the last few columns.
`str(df)`	A summary of the data frame, including the structure of each column.
`summary(df)`	A summary of the data frame, including some of its statistics features.
`levels(df$<column>)`	The level of the column.

have only indices but not names, while columns have both indices and names. In the case of data frame, `$` can be used to access a column as follows.

```
<df>$<column-name> # equivalent to <df>[, <column-name>]
```

which returns all elements in the column as a vector. The row index `[<row-index>]` can follow up to further specify an element if necessary.

Table 9.7 summarizes the commonly used commands for data frame exploration, such as checking its shape and data types.

An example of applying the above functions to `iris` data frame from the built-in `datasets` package is given below.

```
> library(datasets)
> nrow(iris)
[1] 150
> ncol(iris)
[1] 5
> head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4         0.2 setosa
2          4.9         3.0          1.4         0.2 setosa
3          4.7         3.2          1.3         0.2 setosa
4          4.6         3.1          1.5         0.2 setosa
5          5.0         3.6          1.4         0.2 setosa
6          5.4         3.9          1.7         0.4 setosa
> tail(iris)
    Sepal.Length Sepal.Width Petal.Length Petal.Width Species
145          6.7         3.3          5.7         2.5 virginica
146          6.7         3.0          5.2         2.3 virginica
147          6.3         2.5          5.0         1.9 virginica
148          6.5         3.0          5.2         2.0 virginica
149          6.2         3.4          5.4         2.3 virginica
150          5.9         3.0          5.1         1.8 virginica
> str(iris)
```

```
'data.frame':  150 obs. of 5 variables:
 $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species     : Factor w/ 3 levels "setosa","versicolor",..: 1 1 1 1 1
     1 1 1 1 1 ...
> summary(iris)
  Sepal.Length   Sepal.Width    Petal.Length   Petal.Width        Species
 Min.   :4.300  Min.   :2.000  Min.   :1.000  Min.   :0.100  setosa    :50
 1st Qu.:5.100  1st Qu.:2.800  1st Qu.:1.600  1st Qu.:0.300  versicolor
     :50
 Median :5.800  Median :3.000  Median :4.350  Median :1.300  virginica
     :50
 Mean   :5.843  Mean   :3.057  Mean   :3.758  Mean   :1.199
 3rd Qu.:6.400  3rd Qu.:3.300  3rd Qu.:5.100  3rd Qu.:1.800
 Max.   :7.900  Max.   :4.400  Max.   :6.900  Max.   :2.500
> levels(iris$Species) # only works discrete-value columns
[1] "setosa"    "versicolor" "virginica"
```

Data frame sub-setting works similarly with matrix. A sub-setting of multiple rows and columns of a data frame is a data frame. Notice that unlike the matrix case where if only one row is segmented the return is treated as a vector by default, in the case of a data frame the structure preserves. When a single column is segmented, however, in both matrix and data frame scenarios, the result will be treated as a vector by default, and `drop=F` can be used to preserve data frame structure.

An example is given below.

```
> library(datasets)
> print(iris[1:5,])
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4         0.2 setosa
2          4.9         3.0          1.4         0.2 setosa
3          4.7         3.2          1.3         0.2 setosa
4          4.6         3.1          1.5         0.2 setosa
5          5.0         3.6          1.4         0.2 setosa
> print(iris[1,])
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4         0.2 setosa
> is.data.frame(iris[1,])
[1] TRUE
> print(iris[,1]) # equivalent to print(iris@Sepal.Length)
  [1] 5.1 4.9 4.7 4.6 5.0 5.4 4.6 5.0 4.4 4.9 5.4 4.8 4.8 4.3 5.8 5.7
      5.4 5.1 5.7 5.1 5.4 5.1 4.6 5.1
 [25] 4.8 5.0 5.0 5.2 5.2 4.7 4.8 5.4 5.2 5.5 4.9 5.0 5.5 4.9 4.4 5.1
      5.0 4.5 4.4 5.0 5.1 4.8 5.1 4.6
 [49] 5.3 5.0 7.0 6.4 6.9 5.5 6.5 5.7 6.3 4.9 6.6 5.2 5.0 5.9 6.0 6.1
      5.6 6.7 5.6 5.8 6.2 5.6 5.9 6.1
```

```
 [73] 6.3 6.1 6.4 6.6 6.8 6.7 6.0 5.7 5.5 5.5 5.8 6.0 5.4 6.0 6.7 6.3
      5.6 5.5 5.5 6.1 5.8 5.0 5.6 5.7
 [97] 5.7 6.2 5.1 5.7 6.3 5.8 7.1 6.3 6.5 7.6 4.9 7.3 6.7 7.2 6.5 6.4
      6.8 5.7 5.8 6.4 6.5 7.7 7.7 6.0
[121] 6.9 5.6 7.7 6.3 6.7 7.2 6.2 6.1 6.4 7.2 7.4 7.9 6.4 6.3 6.1 7.7
      6.3 6.4 6.0 6.9 6.7 6.9 5.8 6.8
[145] 6.7 6.7 6.3 6.5 6.2 5.9
> is.data.frame(iris[,1])
[1] FALSE
> print(iris[,1,drop=F]) # preserve data frame
    Sepal.Length
1           5.1
2           4.9
3           4.7
... # WRAPPED # ...
148         6.5
149         6.2
150         5.9
> is.data.frame(iris[,1,drop=F])
[1] TRUE
```

To add a new column to an existing data frame, just assign values to a new column name as follows.

```
<df>$<new-column> <- <vector>
```

if there is a mismatch in size, `<vector>` will be cycled.

Ro remove a column, assign `NULL` to all the elements in that column as follows.

```
<df>$<column> <- NULL
```

9.4.3   Filtering

Filtering is about selecting specific rows from a data frame that meet specific criteria. A true-false vector can be used as a filter as follows.

```
<filter-name> <- <true-false-vector> # use true-false vector as filter
<df>[filter,] # implement filter on data frame
```

An example is given below.

```
> library(datasets)
> filter <- iris$Sepal.Length >= 7
> print(iris[filter,])
    Sepal.Length Sepal.Width Petal.Length Petal.Width Species
51           7.0         3.2          4.7         1.4 versicolor
103          7.1         3.0          5.9         2.1 virginica
106          7.6         3.0          6.6         2.1 virginica
108          7.3         2.9          6.3         1.8 virginica
110          7.2         3.6          6.1         2.5 virginica
```

```
118          7.7          3.8          6.7          2.2 virginica
119          7.7          2.6          6.9          2.3 virginica
123          7.7          2.8          6.7          2.0 virginica
126          7.2          3.2          6.0          1.8 virginica
130          7.2          3.0          5.8          1.6 virginica
131          7.4          2.8          6.1          1.9 virginica
132          7.9          3.8          6.4          2.0 virginica
136          7.7          3.0          6.1          2.3 virginica
> filter <- iris$Sepal.Length >= 7 & iris$Sepal.Width >= 3.5
> print(iris[filter,])
    Sepal.Length Sepal.Width Petal.Length Petal.Width Species
110          7.2          3.6          6.1          2.5 virginica
118          7.7          3.8          6.7          2.2 virginica
132          7.9          3.8          6.4          2.0 virginica
```

As shown above, it is possible to use &, | to form a more complex filter. The commands can be merged together as follows.

```
> print(iris[iris$Sepal.Length >= 7,])
    Sepal.Length Sepal.Width Petal.Length Petal.Width Species
51           7.0          3.2          4.7          1.4 versicolor
103          7.1          3.0          5.9          2.1 virginica
106          7.6          3.0          6.6          2.1 virginica
108          7.3          2.9          6.3          1.8 virginica
110          7.2          3.6          6.1          2.5 virginica
118          7.7          3.8          6.7          2.2 virginica
119          7.7          2.6          6.9          2.3 virginica
123          7.7          2.8          6.7          2.0 virginica
126          7.2          3.2          6.0          1.8 virginica
130          7.2          3.0          5.8          1.6 virginica
131          7.4          2.8          6.1          1.9 virginica
132          7.9          3.8          6.4          2.0 virginica
136          7.7          3.0          6.1          2.3 virginica
> nrow(iris[iris$Sepal.Length >= 7,]) # count the result number
[1] 13
```

9.4.4 Building Data Frames

To create a data frame from scratch, use function `data.frame()` as follows.

```
<df> <- data.frame(<vector>, ...) # add a column
colnames(<df>) <- c("<column-name>", ...)
```

or

```
<df> <- data.frame(<column-name> = <vector>, ...)
```

An example is given below, where a data frame of mortgage price at 3 types of areas, namely "CBD", "city" and "suburbs", are is created. Arbitrary data is used.

```
# create data frame
vec_region <- c(rep("CBD", 100), rep("City", 100), rep("Suburbs", 100))
vec_size_cbd <- rnorm(100, 75, 10)
vec_size_city <- rnorm(100, 100, 15)
vec_size_suburbs <- rnorm(100, 150, 25)
vec_size = c(vec_size_cbd, vec_size_city, vec_size_suburbs)
vec_price_cbd <- vec_size_cbd*rnorm(100, 12500, 2500)
vec_price_city <- vec_size_city*rnorm(100, 7500, 1000)
vec_price_suburbs <- vec_size_suburbs*rnorm(100, 5000, 1000)
vec_price <- c(vec_price_cbd, vec_price_city, vec_price_suburbs)
mortgage_price <- data.frame(Region = vec_region, Size = vec_size,
    Price = vec_price)
rm(vec_region, vec_size_cbd, vec_size_city, vec_size_suburbs, vec_size,
     vec_price_cbd, vec_price_city, vec_price_suburbs, vec_price)
```

which gives the following result

```
> head(mortgage_price)
  Region     Size     Price
1    CBD 81.84889 1154873.0
2    CBD 77.78946 831468.7
3    CBD 84.60477 735265.2
4    CBD 62.42625 829977.5
5    CBD 65.42723 933851.3
6    CBD 82.43867 1208589.0
```

A data frame can also be created from two existing data frames by merging them together. It works like the "JOIN" function in SQL, and in this sense it supports all "INNER JOIN", "LEFT JOIN", "RIGHT JOIN" and "OUTER JOIN". The syntax is given below.

```
<df> <- merge(
    <df1>,
    <df2>,
    by.x = <column-in-df1>,
    by.y = <column-in-df2>,
    all=FALSE,
    all.x = all, # left join, by default FALSE
    all,y = all, # right join, by default FALSE
    sort = TRUE # sort by the join column
)
```

In case the two data frames have duplicated columns other than the joining columns pair, use `<df>$<column> <- NULL` to unnecessary columns.

9.5 Basic Data Visualizations Using `qplot()`

The package `ggplot2` provides useful tools for visualization of a data frame. For example, both `qplot()` and `ggplot()` in `ggplot2` provide plot function. Notice that in the late versions of `ggplot2`, `qplot()` is deprecated to encourage using of the more powerful `ggplot()`. With that been said, both functions are smart and flexible enough to produce many different types of plots.

An example of `qplot()` is given below, just to show some of its capability. Run the following codes, and Fig. 9.6 is displayed. It can be seen that `qplot()` is smart enough to automatically choose plot dype, background color, etc., to simplify the plot function.

```
library(datasets)
library(ggplot2)
qplot(
    data=iris,
    x=Sepal.Length*Sepal.Width,
    y=Petal.Length*Petal.Width,
    color=Species,
    size=I(3),
    xlab = "Sepal Area",
    ylab = "Petal Area"
)
```

As a recap, the `mortgage_price` data frame created previously can be visualized as follows. Figures 9.7 and 9.8 can be obtained.

```
library(ggplot2)
rm(list=ls())
# create data frame
vec_region <- rep(c("CBD","City","Suburbs"),each = 100)
vec_size_cbd <- rnorm(100, 75, 10)
vec_size_city <- rnorm(100, 100, 15)
vec_size_suburbs <- rnorm(100, 150, 25)
vec_size = c(vec_size_cbd, vec_size_city, vec_size_suburbs)
vec_price_cbd <- vec_size_cbd*rnorm(100, 12500, 2500)
vec_price_city <- vec_size_city*rnorm(100, 7500, 1000)
vec_price_suburbs <- vec_size_suburbs*rnorm(100, 5000, 1000)
vec_price <- c(vec_price_cbd, vec_price_city, vec_price_suburbs)
mortgage_price <- data.frame(Region = vec_region, Size = vec_size,
    Price = vec_price)
rm(vec_region, vec_size_cbd, vec_size_city, vec_size_suburbs, vec_size,
     vec_price_cbd, vec_price_city, vec_price_suburbs, vec_price)
# processing
mortgage_price$Price.Unit <- mortgage_price$Price / mortgage_price$Size
# plot
qplot(data=mortgage_price, x=Size, y=Price, color=Region, geom=c("point
    ", "smooth"))
```
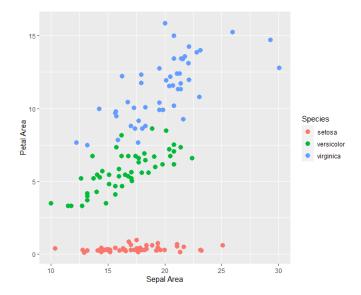
FIGURE 9.6
A demonstration of `qplot`.
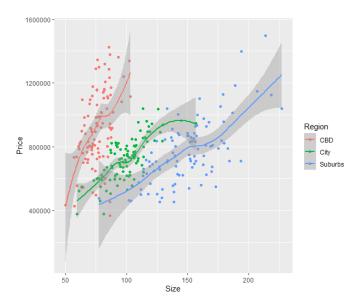
```
qplot(data=mortgage_price, x=Region, y=Price.Unit, geom="boxplot")
```

9.6 Advanced Data Visualizations Using `ggplot()`
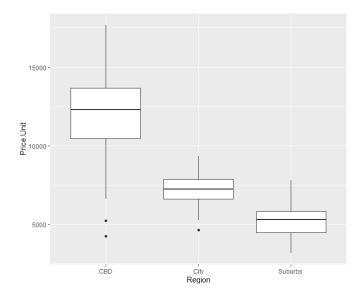
Function `ggplot()` is the main data visualization tool in `ggolot2` package. It provides very flexible approaches for data plotting.

9.6.1 Grammar of Graphics

As proposed by Leland Wilkinson's Grammar of Graphics, a chart shall contain multiple independent and reusable layers including "data" (as data in data frames), "aesthetics" (how data maps to the chart, i.e., the logic of the plot; for example sample dots, curve, color block, or length of lines/bars), "geometries" (the actual color and shape of each element on the chart), "statistics" (information derived from the data being represented in the chart), "facets" (subplots of the same style align together for comparison), "coordinates" (the meaning and range of axis) and "theme" (overall design, such as title, label, etc.). A demonstrative Fig. 9.9 is given to illustrate the different layers in a chart.

FIGURE 9.7
A demonstration of `qplot` on mortgage price data frame.

FIGURE 9.8
A second demonstration of `qplot` on mortgage price data frame.
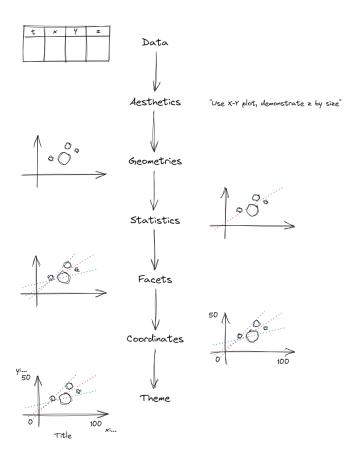
FIGURE 9.9
Multiple layers in chart design.

TABLE 9.8
Commonly used commands for data frame exploration.

Geom	Description
`geom_point()`	Scatter plots and dot plots.
`geom_line()`	Line plots.
`geom_bar()`	Bar plots.
`geom_histogram()`	Histograms.
`geom_boxplot()`	Box plots.
`geom_violin()`	Violin plots.
`geom_density()`	Density plots.
`geom_density2d()`	2-dimensional Density plots.
`geom_text()`	Text Annotation.
`geom_label()`	Label on the observations.

9.6.2 Data, Aesthetics and Geometries Layers

Function `ggplot()` is a very good practice of implementing the above chart design and plotting philosophy. A simple example for `ggplot()`, just for quick demonstration purpose, is given below.

```
p <- ggplot(data=<df>, aes(x=<column>, y=<column>, color=<column>, size
    =<column>))
p <- p + geom_point()
p
```

where `aes()` is used to build mappings in the aesthetics.

An interesting fact when using `ggplot()` is that, when adding a layer to the chat, the layer is literally added to `ggplot()`. In the program, this step by step build up an object, where `ggplot()` provides the most basic layers. Therefore, the above simple example is equivalent to

```
p <- ggplot(data=<df>, aes(x=<column>, y=<column>, color=<column>, size
    =<column>))
p + geom_point()
p + geom_line() # another selection of geometries choice
p + geom_line() + geom_point() # overlying both geometries choices
p + geom_point(aes(size=<another-column>)) # overwrite aesthetics
    settings.
```

and the added layers are able to inherit the aesthetics settings, if it is not overwritten. And speaking of overwriting, even the x and y axis can be overwritten. The displaying name of the labels can be overwritten by stack `xlab("")` and `ylab("")` into the chart.

Function `ggplot()` provides many choices for geometries. The most commonly used ones are summarized in Table 9.8.

9.6.3 Statistics Layers

Similar to the case of geometries layers, statistics layers can also be stacked to `ggplot()`. As introduced earlier, statistics layers are often "add-on" layers that derives statistical features from the data and provide additional insights to the users.

Many functions in Table 9.8 are statistics layer built-in, such as `geom_boxplot()` which by nature is a statistics result presentation in the first place. Regression functions such as `geom_smooth()` also reveals statistical insights of the data. More details of these functions are as follows.

Consider using `geom_boxplot()` to visualize the `mortgage_price` data frame that was used in the previous section. Examples are given below.

```
library(ggplot2)
# create data frame
vec_region <- rep(c("CBD","City","Suburbs"), each = 100)
vec_size_cbd <- rnorm(100, 75, 10)
vec_size_city <- rnorm(100, 100, 15)
vec_size_suburbs <- rnorm(100, 150, 25)
vec_size = c(vec_size_cbd, vec_size_city, vec_size_suburbs)
vec_price_cbd <- vec_size_cbd*rnorm(100, 12500, 2500)
vec_price_city <- vec_size_city*rnorm(100, 7500, 1000)
vec_price_suburbs <- vec_size_suburbs*rnorm(100, 5000, 1000)
vec_price <- c(vec_price_cbd, vec_price_city, vec_price_suburbs)
mortgage_price <- data.frame(Region = vec_region, Size = vec_size,
    Price = vec_price)
rm(vec_region, vec_size_cbd, vec_size_city, vec_size_suburbs, vec_size,
     vec_price_cbd, vec_price_city, vec_price_suburbs, vec_price)
# processing
mortgage_price$Price.Unit <- mortgage_price$Price / mortgage_price$Size
# plot
p <- ggplot(data=mortgage_price, aes(x=Region, y=Price.Unit, color=
    Region)) + ggtitle("Mortgage uprice per area VS region") + xlab("
    Region") + ylab("Price per area")
p + geom_boxplot() + geom_jitter(aes(size=Size, color=Region), alpha
    =0.25)
```

and the result is shown in Fig. 9.10. Notice that `ggtitle()`, `xlab()`, `ylab()`, `alpha` are used in the plot. They are self-explanatory. A new geometry `geom_jitter()` is used, which works similarly with `geom_point()` except the additional vibration in the horizontal axis which makes the points clearer to see.

Function `geom_smooth()` is widely used for curve fitting. An example is given below.

```
library(ggplot2)
# generate data
t <- 1:500
var1 <- 1.5*t + rnorm(500, 0, 100)
var2 <- 0.5*t + rnorm(500, 200, 10) + t^1.3*rnorm(500, 0, 0.1)
```
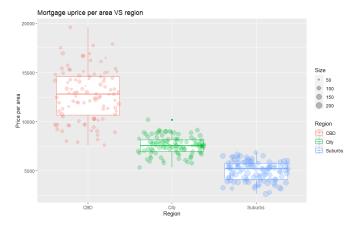
FIGURE 9.10
An example of box plot of the mortgage price data frame using `ggplot()` and `geom_boxplot()`.

```
df <- data.frame(t=t, x=var1, y=var2)
# plot data
p <- ggplot(data=df) +
ggtitle("Plot of x and y VS t.") +
xlab("t") +
ylab("x and y") +
geom_point(aes(x=t, y=x), color="blue", shape=1, size=1.5) +
geom_smooth(aes(x=t, y=x), color="blue") +
geom_point(aes(x=t, y=y), color="red", shape=2, size=1.5) +
geom_smooth(aes(x=t, y=y), color="red")
p
```

Do note that aesthetics needs to be given to `geom_smooth()` in the above example. This is because aesthetics is not given in the base `ggolot()`. Notice that `geom_smooth()` can inherit aesthetics from the previous `ggplot()`, but not from the previous `geom_point()`. The plot is given by Fig. 9.11.

More functions similar to `geom_smooth()` are summarized in Table 9.9.

9.6.4 Facets Layers

The facets layer allows subplot of data. Consider the following example, where the distribution of mortgage price is studied using histogram. The following code can be used to plot the result in a single plot without the facets layer. The plot is given in Fig. 9.12.

```
library(ggplot2)
# create data frame
Region = rep(c("CBD","City","Suburbs"), each = 500)
```
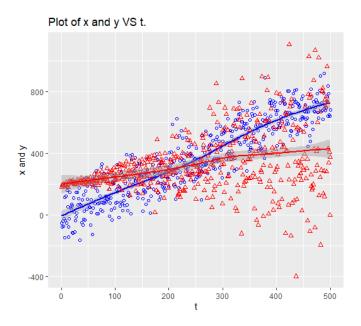
FIGURE 9.11
An example of using `geom_smooth()` for scatter point fitting.

TABLE 9.9
Functions that fit smooth lines to scatter points.

Function	Description
`loess()`	Non-parametric method for fitting a smooth line to a scatter plot using locally weighted regression algorithm.
`smooth.spline()`	Fits a smoothing spline to the data, which is a type of regression spline where the degree of smoothing is chosen automatically by cross-validation.
`lm()`	Linear Model, fits a linear relationship between independent and dependent variables by minimizing the residuals between the data points and the line.
`glm()`	Generalized Linear Model, similar to linear model, but it allows different distribution of error other than normal.
`gam()`	Generalized Additive Model, it is similar to GLM, but it allows non-parametric smooth functions to be added to the linear predictor.
`geom_smooth()`	A function in ggplot2 that is used to add a smooth line to a scatter plot, it uses method = "loess" by default but also allow to use other smoothing method like lm, gam etc.
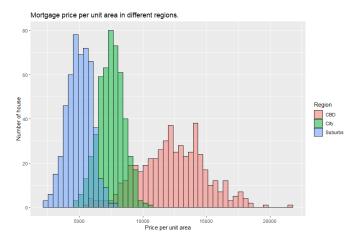
FIGURE 9.12
An example of histogram plot of house price per unit area in different regions
in a single plot.

```
vec_size = list(vec_size_cbd = rnorm(500, 75, 10), vec_size_city =
    rnorm(500, 100, 15), vec_size_suburbs = rnorm(500, 150, 25))
vec_price = list(vec_price_cbd = vec_size$vec_size_cbd*rnorm(500,
    12500, 2500),
                vec_price_city = vec_size$vec_size_city*rnorm(500,
                    7500, 1000),
                vec_price_suburbs = vec_size$vec_size_suburbs*rnorm
                    (500, 5000, 1000))
mortgage_price <- data.frame(Region = Region,
                            Size = unlist(vec_size),
                            Price = unlist(vec_price))
mortgage_price$Region <- as.factor(mortgage_price$Region)
mortgage_price$Price.Unit <- mortgage_price$Price / mortgage_price$Size
# plot data
p <- ggplot(data=mortgage_price, aes(x=Price.Unit))
p + geom_histogram(aes(fill=Region), bins=50, color="black", alpha=0.5,
    position="identity") +
  ggtitle("Mortgage price per unit area in different regions.") +
  xlab("Price per unit area") +
  ylab("Number of house")
```

To use facets layer, revise the code as follows. Notice that `facet_grid()`
is added to the plot, and its input `<column>~.` or `.~<column>` (it is okay to
use `<column1>~<column2>` as well) decide the design of the subplots (how to
arrange the rows and columns of the subplots).

```
library(ggplot2)
# create data frame
Region = rep(c("CBD","City","Suburbs"), each = 500)
```

```
vec_size = list(vec_size_cbd = rnorm(500, 75, 10),
                vec_size_city = rnorm(500, 100, 15),
                vec_size_suburbs = rnorm(500, 150, 25))
vec_price = list(vec_price_cbd = vec_size$vec_size_cbd*rnorm(500,
    12500, 2500),
                vec_price_city = vec_size$vec_size_city*rnorm(500,
                    7500, 1000),
                vec_price_suburbs = vec_size$vec_size_suburbs*rnorm
                    (500, 5000, 1000))
mortgage_price <- data.frame(Region = Region,
                        Size = unlist(vec_size),
                        Price = unlist(vec_price))
mortgage_price$Region <- as.factor(mortgage_price$Region)
mortgage_price$Price.Unit <- mortgage_price$Price / mortgage_price$Size
# plot data
p <- ggplot(data=mortgage_price, aes(x=Price.Unit))
p <- p + geom_histogram(aes(fill=Region), bins=50, color="black", alpha
    =0.5, position="identity") +
  ggtitle("Mortgage price per unit area in different regions.") +
  xlab("Price per unit area") +
  ylab("Number of house")
p
p + facet_grid(Region~.) # put subplots for different regions in rows
p + facet_grid(.~Region) # put subplots for different regions in
    columns
```

The results are given in Figs. 9.13 and 9.14, depending on the subplot designs.

9.6.5 Coordinates Layers

Coordinate control is important. The coordinate layer allows setting limits to the axis and zooming in to the chart. An example of adding coordinates layers to a plot is given as follows. The same mortgage price data frame is used for illustration.

```
library(ggplot2)
# create data frame
Region = rep(c("CBD","City","Suburbs"), each = 500)
vec_size = list(vec_size_cbd = rnorm(500, 75, 10),
                vec_size_city = rnorm(500, 100, 15),
                vec_size_suburbs = rnorm(500, 150, 25))
vec_price = list(vec_price_cbd = vec_size$vec_size_cbd*rnorm(500,
    12500, 2500),
                vec_price_city = vec_size$vec_size_city*rnorm(500,
                    7500, 1000),
                vec_price_suburbs = vec_size$vec_size_suburbs*rnorm
                    (500, 5000, 1000))
mortgage_price <- data.frame(Region = Region,
```

FIGURE 9.13
Use facets to plot the histogram of price per unit are of the house in different regions (subplots in rows).
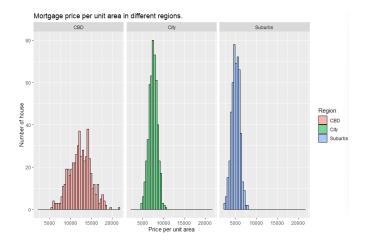
FIGURE 9.14
Use facets to plot the histogram of price per unit are of the house in different regions (subplots in columns).
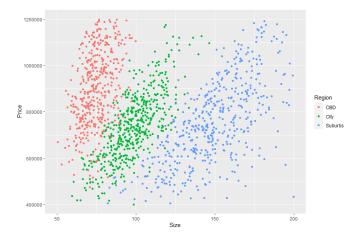
FIGURE 9.15
Add coordinates layer using `xlim()` and `ylim()`.

```
                            Size = unlist(vec_size),
                            Price = unlist(vec_price))
mortgage_price$Region <- as.factor(mortgage_price$Region)
mortgage_price$Price.Unit <- mortgage_price$Price / mortgage_price$Size
# plot data
p <- ggplot(data=mortgage_price, aes(x=Size, y=Price))
p <- p + geom_point(aes(color=Region))
p + xlim(50, 200) + ylim(400000, 1200000) # first chart
p + coord_cartesian(xlim=c(50, 200), ylim = c(400000, 1200000)) #
    second chart
```

where notice that two charts are generated. The first chart using `xlim()`, `ylim` removes all samples outside the boundary from the chart. While in the second chart using `coord_cartesian()`, all samples preserves and the chart zooms in towards the boundary. The results are given in Figs. 9.15 and 9.16, respectively. The difference can be observed near the boundary.

9.6.6 Themes Layers

Theme layers mainly refer to titles, labels, and other comments on the chart that help with understanding the content of the chart. As already demonstrated in previous examples, use `xlab()`, `ylab()` to add labels, `ggtitle()` to add title.

Use `theme()` to change the themes of the labels. An example is given below.

```
library(ggplot2)
# create data frame
Region = rep(c("CBD","City","Suburbs"), each = 500)
```
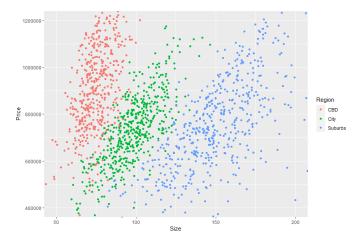
FIGURE 9.16
Add coordinates layer using `coord_cartesian()`.

```
vec_size = list(vec_size_cbd = rnorm(500, 75, 10), vec_size_city =
    rnorm(500, 100, 15), vec_size_suburbs = rnorm(500, 150, 25))
vec_price = list(vec_price_cbd = vec_size$vec_size_cbd*rnorm(500,
    12500, 2500),
                 vec_price_city = vec_size$vec_size_city*rnorm(500,
                     7500, 1000),
                 vec_price_suburbs = vec_size$vec_size_suburbs*rnorm
                     (500, 5000, 1000))
mortgage_price <- data.frame(Region = Region,
                             Size = unlist(vec_size),
                             Price = unlist(vec_price))
mortgage_price$Region <- as.factor(mortgage_price$Region)
mortgage_price$Price.Unit <- mortgage_price$Price / mortgage_price$Size
# plot data
p <- ggplot(data=mortgage_price, aes(x=Price.Unit))
p <- p + geom_histogram(aes(fill=Region), bins=50, color="black", alpha
    =0.5, position="identity")
p + ggtitle("Mortgage price per unit area in different regions.") +
  xlab("Price per unit area") +
  ylab("Number of house") +
  theme(axis.title.x = element_text(color = "DarkGreen", size=15),
        axis.title.y = element_text(color = "DarkRed", size=15),
        axis.text.x = element_text(size=10),
        axis.text.y = element_text(size=10),
        legend.title = element_text(size=10),
        legend.text = element_text(size=8),
        legend.position = c(1,1), # right top corner of chart
        legend.justification = c(1,1), # legend align point
        plot.title = element_text(color = "DarkBlue", size = 15)
```
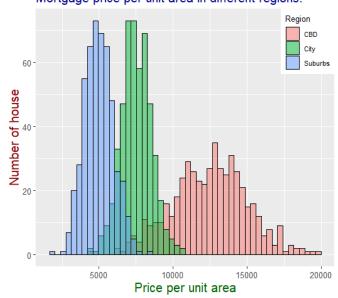
FIGURE 9.17
Mortgage price chart with theme.

```
)
```

The resulted chart is given in Fig. 9.17. Compare it with Fig. 9.12 to see the differences by applying `theme()` in the themes layer.

9.7   Data Preparation

The data downloaded from sensors usually needs to go through pre-processing procedures such as filtering, normalization, etc., before it can be used by a controller or an AI engine.

9.7.1   Data Type Conversion

Data preparation including data tidy is one of the most tedious and time consuming parts when using R for data analysis. The section introduces useful techniques helpful with data preparation.

It is importnt that the data types of all the columns meet expectation, especially for numeric and factor (categorical) data types. Use `str(<df>)` to check the column data types of a data frame, and convert data types as follows.

```
<df>$<column> <- factor(<df>$<column>) # character/numeric to factor
<df>$<column> <- as.numeric(<df>$<column>) # character to numeric
<df>$<column> <- as.numeric(as.character(<df>$<column>)) # factor to
    numeric
```

Notice that when converting factor type to other types, R may deal with the factor using the underlying "factorization integers" instead of the factor item names. An example is given below. It can be seen that the original 5.1, after being converted to factor then back to numeric, becomes 9. This is because the factorization integer for 5.1 is 9, as given by printing `my_factor` to the console.

```
> library(datasets)
> iris$Sepal.Length
  [1] 5.1 4.9 4.7 4.6 5.0 5.4 4.6 5.0 4.4 4.9 5.4 4.8 4.8 4.3 5.8 5.7
 [17] 5.4 5.1 5.7 5.1 5.4 5.1 4.6 5.1 4.8 5.0 5.0 5.2 5.2 4.7 4.8 5.4
 [33] 5.2 5.5 4.9 5.0 5.5 4.9 4.4 5.1 5.0 4.5 4.4 5.0 5.1 4.8 5.1 4.6
 [49] 5.3 5.0 7.0 6.4 6.9 5.5 6.5 5.7 6.3 4.9 6.6 5.2 5.0 5.9 6.0 6.1
 [65] 5.6 6.7 5.6 5.8 6.2 5.6 5.9 6.1 6.3 6.1 6.4 6.6 6.8 6.7 6.0 5.7
 [81] 5.5 5.5 5.8 6.0 5.4 6.0 6.7 6.3 5.6 5.5 5.5 6.1 5.8 5.0 5.6 5.7
 [97] 5.7 6.2 5.1 5.7 6.3 5.8 7.1 6.3 6.5 7.6 4.9 7.3 6.7 7.2 6.5 6.4
[113] 6.8 5.7 5.8 6.4 6.5 7.7 7.7 6.0 6.9 5.6 7.7 6.3 6.7 7.2 6.2 6.1
[129] 6.4 7.2 7.4 7.9 6.4 6.3 6.1 7.7 6.3 6.4 6.0 6.9 6.7 6.9 5.8 6.8
[145] 6.7 6.7 6.3 6.5 6.2 5.9
> my_factor <- factor(iris$Sepal.Length)
> my_numeric <- as.numeric(my_factor)
> my_numeric
  [1]  9  7  5  4  8 12  4  8  2  7 12  6  6  1 16 15 12 9 15 9 12  9
 [23]  4  9  6  8  8 10 10  5  6 12 10 13  7  8 13  7  2  9  8  3  2  8
 [45]  9  6  9  4 11  8 28 22 27 13 23 15 21  7 24 10  8 17 18 19 14 25
 [67] 14 16 20 14 17 19 21 19 22 24 26 25 18 15 13 13 16 18 12 18 25 21
 [89] 14 13 13 19 16  8 14 15 15 20  9 15 21 16 29 21 23 33  7 31 25 30
[111] 23 22 26 15 16 22 23 34 34 18 27 14 34 21 25 30 20 19 22 30 32 35
[133] 22 21 19 34 21 22 18 27 25 27 16 26 25 25 21 23 20 17
> typeof(my_factor)
[1] "integer"
> my_factor
  [1] 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 5.4 4.8 4.8 4.3 5.8 5.7
 [17] 5.4 5.1 5.7 5.1 5.4 5.1 4.6 5.1 4.8 5 5 5.2 5.2 4.7 4.8 5.4
 [33] 5.2 5.5 4.9 5 5.5 4.9 4.4 5.1 5 4.5 4.4 5  5.1 4.8 5.1 4.6
 [49] 5.3 5  7   6.4 6.9 5.5 6.5 5.7 6.3 4.9 6.6 5.2 5 5.9 6 6.1
 [65] 5.6 6.7 5.6 5.8 6.2 5.6 5.9 6.1 6.3 6.1 6.4 6.6 6.8 6.7 6 5.7
 [81] 5.5 5.5 5.8 6 5.4 6   6.7 6.3 5.6 5.5 5.5 6.1 5.8 5 5.6 5.7
 [97] 5.7 6.2 5.1 5.7 6.3 5.8 7.1 6.3 6.5 7.6 4.9 7.3 6.7 7.2 6.5 6.4
[113] 6.8 5.7 5.8 6.4 6.5 7.7 7.7 6 6.9 5.6 7.7 6.3 6.7 7.2 6.2 6.1
[129] 6.4 7.2 7.4 7.9 6.4 6.3 6.1 7.7 6.3 6.4 6 6.9 6.7 6.9 5.8 6.8
[145] 6.7 6.7 6.3 6.5 6.2 5.9
35 Levels: 4.3 4.4 4.5 4.6 4.7 4.8 4.9 5 5.1 5.2 5.3 5.4 5.5 ... 7.9
```

When converting factor to other types, special caution is required. To

convert a factor to other types such as numeric, consider converting it to character first as given in the following example.

```
> my_numeric <- as.numeric(as.character(my_factor))
> my_numeric
  [1] 5.1 4.9 4.7 4.6 5.0 5.4 4.6 5.0 4.4 4.9 5.4 4.8 4.8 4.3 5.8 5.7
 [17] 5.4 5.1 5.7 5.1 5.4 5.1 4.6 5.1 4.8 5.0 5.0 5.2 5.2 4.7 4.8 5.4
 [33] 5.2 5.5 4.9 5.0 5.5 4.9 4.4 5.1 5.0 4.5 4.4 5.0 5.1 4.8 5.1 4.6
 [49] 5.3 5.0 7.0 6.4 6.9 5.5 6.5 5.7 6.3 4.9 6.6 5.2 5.0 5.9 6.0 6.1
 [65] 5.6 6.7 5.6 5.8 6.2 5.6 5.9 6.1 6.3 6.1 6.4 6.6 6.8 6.7 6.0 5.7
 [81] 5.5 5.5 5.8 6.0 5.4 6.0 6.7 6.3 5.6 5.5 5.5 6.1 5.8 5.0 5.6 5.7
 [97] 5.7 6.2 5.1 5.7 6.3 5.8 7.1 6.3 6.5 7.6 4.9 7.3 6.7 7.2 6.5 6.4
[113] 6.8 5.7 5.8 6.4 6.5 7.7 7.7 6.0 6.9 5.6 7.7 6.3 6.7 7.2 6.2 6.1
[129] 6.4 7.2 7.4 7.9 6.4 6.3 6.1 7.7 6.3 6.4 6.0 6.9 6.7 6.9 5.8 6.8
[145] 6.7 6.7 6.3 6.5 6.2 5.9
```

where `my_factor` is generated previously.

It is possible for some columns in the data frame to look like a factor and a character string, but indeed should be handled as numeric values. For example, S$6,125.50 in many occasions should be treated just as 6125.0. These factor or character values cannot be converted to numeric values directly.

In such cases, consider using `sub()` or `gsub()` to replace patterns in a character, then convert it into numeric values. Notice that `sub()` replaces only the first encounter of the pattern, while `gsub()` replaces all the encounters. An example of using `gsub()` is given below.

```
> money_character <- c("S$6,273.15", "S$215.3", "S$8,987,756.00")
> typeof(money)
[1] "character"
> a <- gsub(",", "", money) # replace "," with ""
> a <- gsub("S\\$", "", a) # replace "S$" with ""
> money_numeric <- as.numeric(a)
> money_numeric
[1]    6273.15    215.30 8987756.00
> typeof(money_numeric)
[1] "double"
```

where notice that $ is a special character defined in R, and to escape from that \\$ is used. Notice that applying `sub()` and `gsub()` on a factor automatically converts it to character as a hidden step.

9.7.2 Handling Missing Data

There can be missing data in the data frame. There are a few ways to deal with missing data as follows.

- If the missing data can be derived from other columns, derive the missing data and fill in the blanks.

- If the missing data does not affect the rest analysis, leave it blank.

- Delete the row.

- Use interpolations to fill in the blank.

- Use correlations and similarities to fill in the blank.

- Argument a new column add a "data-missing" flag to that row.

If the missing data can be derived from other columns, derive the missing data and fill in the blanks.

In R, `NA` is a special variable used to indicate a missing value, and it is by itself of logical data type in addition to `TURE` and `FALSE`. Operations involving `NA` often return `NA`. Examples are given below.

```
> typeof(NA)
[1] "logical"
> TRUE == 1 # TRUE is equivalent with 1
[1] TRUE
> TRUE == 2
[1] FALSE
> FALSE == 0 # FALSE is equivalent with 0
[1] TRUE
> FALSE == -1
[1] FALSE
> TRUE == FALSE
[1] FALSE
> NA == NA
[1] NA
> NA == TRUE
[1] NA
> NA == FALSE
[1] NA
```

Use the following to filter for all rows with/without at least one `NA`.

```
<df>[complete.cases(<df>),] # all complete rows
<df>[!complete.cases(<df>),] # all incomplete rows
```

where `complete.cases(<df>)` returns a list made up of `TRUE` and `FALSE` indicating whether the associated row is complete or now.

Sometimes a blank string `""` that we would expected to be treated as `NA` is not treated as so. To fix that, while importing the data frame (say, from a CSV file), use the following

```
df <- read.csv("<csv-name>", na.string=c("<pattern>", ...))
```

where `"<pattern>"` are the patterns in the original file to be replaced by `NA`, for instance, `""`, `"ERROR"`, etc.

9.8   Connectivity with Data Sources

This section introduces the connectivity of R to the data sources, such as a file, or a database.

10

R for Data Science

CONTENTS

11

Python Basics

CONTENTS

Python has been increasingly popular for data science in the past few years. Many libraries and tools have been developed for Python to enhance its data analysis and visualization capabilities, just to name a few, `numpy`, `scipy`, `scikit-learn`, `pandas`, `matplotlib tensorflow` and `pytorch`.

This chapter together with a few consequent chapters introduces commonly used tools that data science adopts using Python. This part of the notebook is more application driven, and only the basic implementations are introduced. We are not digging into the theory supporting machine learning and artificial intelligence.

It has been increasingly popular today to use Python together with Conda and jupyter-lab/jupyter notebook. Conda is an open-source language-agnostic package and environment management system. Jupyter notebook is an interactive computing platform for Python and other computer programming languages. The detailed introduction to the installation and usage of Conda and Jupyter notebook is not covered here. They are used when demonstrating the examples in this chapter and the consequent chapters.

11.1 NumPy

When comes to any sort of numerical computation, one of the most popular Python packages is definitely NumPy. NumPy allows quickly deployment of

numerical vectors, matrices and tensors, as well as associated efficient numerical calculations. It is the "MATLAB" package in Python.

Details of NumPy can be found at *numpy.org*.

The following commands can be used to create NumPy arrays and matrices.

```python
import numpy as np

# create numpy array from python list
x = np.array([1, 2, 3, 4, 5]) # 1d
x = np.array([[1, 2], [3, 4]]) #2d

# create numpy array/matrix using built-in functions
x = np.arange(0, 5) # array([0, 1, 2, 3, 4])
x = np.linspace(0, 5, 3) # array([0, 2.5, 5])
x = np.zeros(5) # 1d zero vector
x = np.zeros([5, 5]) # 2d zero matrix
x = np.ones(5)
x = np.ones([5, 5])
x = np.eye(5)

# create random vector/matrix
np.random.seed(1) # set seed; optional
x = np.random.rand(5, 5) # uniform distribution [0, 1)
x = np.random.randn(5, 5) # standard normal distribution N(0, 1)

# reshape
x = np.random.randn(16)
y = x.reshape(4, 4)
y = x.reshape(1, 16) # return is always 2d matrix format, not 1d vector
    format
```

Aggregation functions are defined. These functions are used to calculate maximum, minimum, sum, etc., of a vector or a matrix. Examples are given below.

```python
import numpy as np

x = np.random.randn(10)
xmax = np.max(x)
xargmax = np.argmax(x)
xmin = np.min(x)
xargmin = np.argmin(x)
xsum = np.sum(x)
xprod = np.prod(x)
xmean = np.mean(x)
xstd = np.std(x)
xvar = np.var(x)
xmedian = np.median(x)
```

When some of these functions such as `sum()` are applied to matrix, it is important to specify the axis along which the calculation would be carried out.

```
import numpy as np

x = np.array([[1,2,3,4,5], [6,7,8,9,10]])
np.sum(x, axis=0) # array([ 7, 9, 11, 13, 15])
np.sum(x, axis=1) # array([15, 40])
```

Accessing (reading and modifying) values in numpy arrays or matrices using the index. Examples are given below. Notice that all examples are about vector accessing.

```
import numpy as np

x = np.random.randn(10)
x[1:5] # x[1], x[2], x[3] and x[4] are returned in an numpy array
x[:5] # same as x[0:5]
x[5:] # same as x[5:10]
```

Matrix accessing is a bit more tricky. A matrix in NumPy is stored like a nested array. Examples are given below to access a single item, a row, and a column or a matrix.

```
import numpy as np

x = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
x[1, 2] # 6 (2nd row, 3rd column)
x[0] # 1st row as a numpy array
x[:,1] # 2nd column as a numpy array
```

NumPy provides efficient and convenient vector and matrix level calculations, such as broad casting, matrix multiplication, etc. Broad casting essentially allows element-by-element basis adding, subtracting or assigning scalar value to a vector or a matrix. An example is given below.

```
import numpy as np

x = np.array([1, 2, 3, 4, 5])
x[:] = 10 # all elements become 10
```

NumPy array and matrix support boolean selection. An example is given below.

```
import numpy as np

x = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9])
x > 5 # array([False, False, False, False, False, True, True, True,
    True])
x[x>5] # array([6, 7, 8, 9])
```

NumPy supports both element-by-element or vector-level operations, including `+`, `-`, `*`, `/`, `**`, `numpy.sqrt()`, `numpy.log()`, etc. Most of these operators are executed element-by-element by default.

11.2   SciPy

SciPy is a library collections of "comprehensive" algorithms widely used in scientific and technical calculations. Notice that NumPy also has built in basic and commonly algorithms in the library, such as FFT. For those algorithms not included in NumPy, there is a chance that it is in SciPy, such as K-means clustering.

A detailed documentation of SciPy functions put into different categories are given here *docs.scipy.org/doc/scipy/reference/.*

11.3   Matplotlib and Seaborn

Data visualization is important through out the entire data analysis process. It is about not only demonstrating the results to the audiences, but also helping the developers to understand the data and improving the inefficient designs in the pipeline. Matplotlib and Seaborn are two important data visualization libraries. They are briefly introduced in this section.

11.3.1   Matplotlib

Matplotlib is the "basic" visualization library. Many data visualization features provided by other packages such as pandas are essentially realized using Matplotlib internally.

Simple line and scatter plots using Matplotlib can be drawn easily. Examples are given below. Pandas series is used as the axis to the plots, but in reality they can be Python arrays or NumPy arrays. The scatter plot used in the example is given in Fig. 11.1.

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

index_s = pd.Series([1, 2, 3, 4, 5, 6, 7, 8, 9, 10])
values_s = pd.Series([1, 1, 2, 3, 5, 8, 13, 21, 34, 55])
df_dict = {
        "F_Index": index_s,
        "F_Values": values_s
}
fibonacci = pd.DataFrame(df_dict)
plt.plot(index_s, values_s) # line
plt.scatter(fibonacci["F_Index"], fibonacci["F_Values"], color="red") #
    scatter
```

```
plt.title("Fibonacci␣Series")
plt.xlabel("n")
plt.ylabel("f(n)")
```
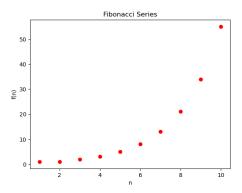

FIGURE 11.1
Plot Fibonacci series as scatter plot.

The presentation of the plot can be customized. For example, use `plt.xlim(x1, x2)`, `plt.ylim(y1, y2)` to change the axis limitations, etc. It is possible to change curve color, size, style, and marker size, style, etc.

11.3.2 Seaborn

Seaborn is another visualization library built on top of the Matplotlib library. It tries to standardize the plots with a simple function call. It scarifies some flexibility that Matplotlib offers, but makes plotting easier.

An example of using Seaborn to plot a histogram is given below. The result is given in Fig. 11.2.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

x = np.random.randn(1000)
plt.figure()
sns.histplot(x, color="red")
plt.title("Histogram␣Example")
plt.xlabel("x")
plt.ylabel("Count")
```

An example of using Seaborn to plot a count plot for discrete values (usually category values) is given below. Notice that the attribute whose values are put into categories is assigned to x of `seaborn.countplot()`. The result is given in Fig. 11.3.
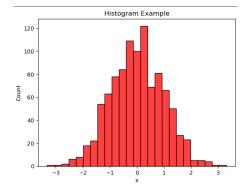
FIGURE 11.2
Histogram plot using Seaborn.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

response = ["yes", "no", "no", "yes", "yes", "yes", "yes", "yes", "yes"
    , "no", "yes", "yes", "NA", "NA"]
plt.figure()
sns.countplot(x=response, color="blue")
plt.title("Count␣Plot␣Example")
plt.xlabel("response")
plt.ylabel("Count")
```
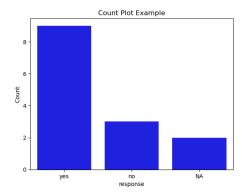

FIGURE 11.3
Count plot using Seaborn.

Bot plot gives the maximum, minimum, and interquartile range (IQR) of the data. In the box plot, the data is split into 4 parts, namely $x < Q_1$

$(0\% - 25\%)$, $Q_1 < x < Q_2$ $(25\% - 50\%)$, $Q_2 < x < Q_3$ $(50\% - 75\%)$, and $Q_3 < x$ $(75\% - 100\%)$. The IQR gives the range between Q_1 and Q_3, i.e., the half in the middle $25\% - 75\%$. An example is given below. The result is given in Fig. 11.4.

```python
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

x = np.random.randn(1000)
plt.figure()
sns.boxplot(x, color="blue")
plt.title("Box Plot Example")
plt.xlabel("Input")
plt.ylabel("Value")
```
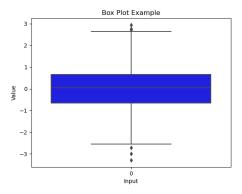

FIGURE 11.4
Box plot using Seaborn. The box gives IQR. The bars below and above the box give $Q_1 - 1.5 \times$ IQR and $Q_3 + 1.5 \times$ IQR, respectively. The dots are outliers.

There are many more plot functions in Seaborn. For example, scatter plot `seaborn.scatterplot()` works similar with Matplotlib as shown by Fig. 11.1. The marker size and color can be adjusted to reflect different parameters, just like in R language. Pairplot `seaborn.pairplot()` generates a matrix of plots to demonstrate associations of columns in a data frame.

11.4   Pandas

Many Python packages provide functions to handle structured data such as tables, series, and data frames. Among all these packages, `pandas` is the all-

time star that is very widely used by developers and data scientists. With `pandas`, Python gains the ability to easily, flexibly and efficiently deal with data frames. The `pandas` package is introduced in this section.

A large portion of this chapter, including codes and examples, are from online resources such as *Data Analysis by Pandas and Python* on Udemy. My special thanks go to them.

The data frames used in the examples of this chapter may come from different public data resources. It is worth mentioning *kaggle.com*, a place where tens of thousands of data sets, code examples, and notebooks are collected and shared. Some data frames used in the examples come from Kaggle.

Two Python packages, `numpy` and `pandas`, are almost certainly used in all the examples to be presented in this chapter. Import these packages as follows.

```
import numpy as np
import pandas as pd
```

Unless otherwise mentioned, these packages are always assumed imported in this chapter. To display the packages version, use something like

```
print("Numpy version: {}.".format(np.__version__))
print("Pandas version: {}".format(pd.__version__))
```

11.4.1 Data Importing

Pandas provide variety of ways to import data from different resources, including plain texts, CSV files, databases, etc. One of the most commonly seen data sources is CSV files. Importing data from CSV files is introduced here.

Pandas provide `.read_csv()` function to import data from CSV files. Its basic usage is introduced below.

```
best_selling_games_df = pd.read_csv("best-selling-video-games.
    csv")
best_selling_games_df
```

which reads all the information in the CSV table into a data frame as shown by Fig. 11.5. It is possible to import only selected columns as follows.

```
best_selling_games_df = pd.read_csv("best-selling-video-games.
    csv", usecols = ["Title", "Sales", "Publisher(s)"])
best_selling_games_df
```

When no index column is specified, pandas will add an additional auto-incremental column and use it as the index column, as shown in Fig. 11.5 by the most-left column. When a column index is specified, pandas will use that column as index column. An example is given below. The result is shown in Fig. 11.6.

```
best_selling_games_df = pd.read_csv("best-selling-video-games.
    csv", index_col = "Title", usecols = ["Title", "Sales", "
    Publisher(s)"])
best_selling_games_df
```

	Rank	Title	Sales	Series	Platform(s)	Initial release date	Developer(s)	Publisher(s)
0	1	Minecraft	238000000	Minecraft	Multi-platform	November 18, 2011	Mojang Studios	Xbox Game Studios
1	2	Grand Theft Auto V	175000000	Grand Theft Auto	Multi-platform	September 17, 2013	Rockstar North	Rockstar Games
2	3	Tetris (EA)	100000000	Tetris	Multi-platform	September 12, 2006	EA Mobile	Electronic Arts
3	4	Wii Sports	82900000	Wii	Wii	November 19, 2006	Nintendo EAD	Nintendo
4	5	PUBG: Battlegrounds	75000000	PUBG Universe	Multi-platform	December 20, 2017	PUBG Corporation	PUBG Corporation
5	6	Mario Kart 8 / Deluxe	60460000	Mario Kart	Wii U / Switch	May 29, 2014	Nintendo EAD	Nintendo
6	7	Super Mario Bros.	58000000	Super Mario	Multi-platform	September 13, 1985	Nintendo R&D4	Nintendo
7	8	Red Dead Redemption 2	50000000	Red Dead	Multi-platform	October 26, 2018	Rockstar Studios	Rockstar Games
8	9	Pokémon Red / Green / Blue / Yellow	47520000	Pokémon	Multi-platform	February 27, 1996	Game Freak	Nintendo
9	10	Terraria	44500000	None	Multi-platform	May 16, 2011	Re-Logic	Re-Logic / 505 Games
10	11	Wii Fit / Plus	43800000	Wii	Wii	December 1, 2007	Nintendo EAD	Nintendo

FIGURE 11.5

The simplest data frame importing using pandas.

Title	Sales	Publisher(s)
Minecraft	238000000	Xbox Game Studios
Grand Theft Auto V	175000000	Rockstar Games
Tetris (EA)	100000000	Electronic Arts
Wii Sports	82900000	Nintendo
PUBG: Battlegrounds	75000000	PUBG Corporation
Mario Kart 8 / Deluxe	60460000	Nintendo
Super Mario Bros.	58000000	Nintendo
Red Dead Redemption 2	50000000	Rockstar Games
Pokémon Red / Green / Blue / Yellow	47520000	Nintendo
Terraria	44500000	Re-Logic / 505 Games
Wii Fit / Plus	43800000	Nintendo

FIGURE 11.6

Specifying index column and reading only selected columns using pandas.

It is possible to read the full-size data frame into pandas first, then sub-sequently select only a few (or event only one) columns to form a new sub data frame. It is possible to take only one column from the data frame, and convert it into a series. Notice that a single-column data frame is different from a series from data type perspective. Examples are given below.

```
# form sub data frame
best_selling_games_df = pd.read_csv("best-selling-video-games.
    csv")
sub_games_df = best_selling_games_df[["Title", "Sales", "
    Publisher(s)"]]
sub_games_df
# form single-column sub data frame
sub_games_df = best_selling_games_df[["Title"]]
sub_games_df
# convert single-column data frame to series
best_selling_games_titles = sub_games_df.squeeze('columns')
best_selling_games_titles
# get series from data frame directly
best_selling_games_df = pd.read_csv("best-selling-video-games.
    csv")
```

```
best_selling_games_titles = best_selling_games_df["Title"]
best_selling_games_titles
```

It is worth mentioning that when generating series from data frames, the index column of the data frame is inherited by the series. This introduces an important feature of pandas series: unlike Python array where it is just single-stream sequence of data, pandas series has a separate measure of index for each element in the series, essentially making it multi-stream of data. More are introduced in later sections.

11.4.2   Series and Data Frame

12

Python-Powered Artificial Intelligence

CONTENTS

This chapter focuses on the introduction of commonly used Python-supported ANN engines used in data science. ANN relationships with AI, machine learning and deep learning as well as theories and mechanisms behind ANN can be found on other notebooks, hence is not given here. The introduction only contains the basic usage of these ANN engines from the implementation perspective, and it may not reflect the state-of-the-art technologies such as transformer, LLM, instruction-tuned LLM, etc. These state-of-the-art technologies are introduced on other notebooks.

There are many ANN engines for Python. Among all, TensorFlow and PyTorch are very popular and powerful generic-purpose ANN engines. They both cover a large range of supervised, reinforcement and unsupervised learning applications including classification, regression, pattern recognition, computer vision, natural language processing, clustering, abnormality detection,

and many more. They both offer variety of tools to quickly and flexibly design and deploy different types of AI models such as conventional dense networks, CNN models, RNN models, and many more. Both of them can be used to train, evaluate and run networks. Both of them provide server solutions, cloud solutions and edge computing solutions. TensorFlow and PyTorch are introduced in this chapter.

Installation of TensorFlow and PyTorch can be found in there websites. Although it is possible to run all the calculations on CPU, these ANN engines are more powerful when GPU/TPU are enabled. Depends on the OS and the GPU/TPU brands of the local system, different methods may apply to enable GPU/TPU. For example, if NVIDIA GPU is used, a software called CUDA can be used to configure and enable GPU for ANN training. The installation of TensorFlow, PyTorch, and the enabling of GPU/TPU modules are not covered here.

Alternative to running the code on a local system, consider using online platforms such as Google Colaboratory, which already have all necessary packages pre-installed and the CPU/GPU/TPU pre-configured.

12.1   Quick Review

This section briefly reviews the basic concepts used in this chapter. Details of the concepts can be found elsewhere in other notebooks.

12.1.1   AI Pipeline

AI pipeline is a set of (automated) steps used to build, train, evaluate and deploy AI models. An AI pipeline usually includes at least the following steps (for supervised learning):

1. Data collection

2. Data preparation

3. Model design

4. Model training

5. Model evaluation and analysis

6. Model deployment and testing

where notice that model design and training might need to be carried out iteratively. After the training, the performance of the model is validated using the validation set, according to which the model and its hyper parameters can be modified.

12.1.2 Data Preparation and Model Evaluation

Model training is where the magic happens. Nowadays, with the help of AI engines, it is done automatically via back propagation and other techniques. Given the same training data and model design (including training methods), the almost-the-same trained model can be reproduced by the machine. It is done in a standardized, systematic and consistent manner, hence does not distinguish the performance of the model.

It is rather the data preparation (pre-processing), model design, and model evaluation that require human guidance. Depending on the experience and skill level of the data scientist and engineer, this is where the model performance may differ. Model design and fine-tuning are closely related to model evaluation, as model evaluation results and analysis decides how the model should be tuned. Therefore, it can be concluded that good data preparation, model evaluation and analysis skills are the critical factors that affect model behavior.

To be more precise by breaking down the aforementioned critical factors:

- Data preprocessing:

  - Data cleaning. This is a critical step that greatly influences the model's performance. It includes cleaning the data, handling missing values, and dealing with outliers. Often, domain knowledge plays a significant role in these steps. In practice, some of the above procedures are done using AI engine built-in functions, while others are done using other packages such as pandas.

  - Feature engineering. This involves creating new features from the existing data that might help improve the model's performance. Feature selection is also crucial to reduce overfitting and improve the model's interpretability. Again, domain knowledge can be very beneficial here.

- Model design, evaluation and tuning:

  - Model selection. Choosing the right model or architecture for the problem at hand requires a solid understanding of the strengths and weaknesses of different models.

  - Hyperparameter tuning. While there are systematic approaches like grid search or random search, often, practical experience and intuition play a significant role in choosing the right hyperparameters.

  - Model evaluation and analysis. Evaluating a model goes beyond looking at a single metric. It involves understanding the model's errors, checking its performance on different subsets of data, and considering aspects like fairness and interpretability.

Given that many, if not all, of the above steps involve a lot of human interaction, data visualization tools often play a very important role to assist humans on their tasks.

12.1.3 Commonly Seen ANN Use Cases

"nobreak

12.1.4 Computer Vision

CV, as an important part of AI, has evolved in the past decades. In the early 2010s, ANN was not used in CV. Instead, conventional deterministic approaches were widely used. With the development in deep learning, the primary approach for CV has changed to CNN. There are a few milestones along the way that together make the change happen:

- Development of GPU

- CNN with deep neural network

- Introduction of rectified linear unit (ReLU) activation function

- Regularization techniques

Recently, with the development in transformer model and large language model, CV is able to be combined with LLM for image comprehension, reasoning, and even artwork generation.

The commonly seen objectives of CV include:

- Image classification

- Object detection

- Image generation

- Image search

- Image comprehension and generation

12.1.5 Natural Language Processing

"nobreak

12.2 TensorFlow

TensorFlow is an open-source software library for machine learning developed by Google in 2015. TensorFlow 2.X is released in 2019 and it is know the official latest major updated version. TensorFlow is backed up by a large community and it supports Python as well as a few other programming languages.

Don't confuse TensorFlow with Keras, later of which is a Python library built on top of deep learning libraries such as TensorFlow, and provides a simple and useful API. In TensorFlow 2.X, Keras is officially adopted as its API. Therefore, when TensorFlow 2.X is used, there is no need to install or import Keras separately.

12.2.1   TensorFlow Basics

Unless otherwise mentioned, the following packages are imported and the command executed in the beginning of all the relevant scripts.

```
import numpy as np
import pandas as pd
import tensorflow as tf

tf.test.is_gpu_available()
```

where `.is_gpu_available()` tests whether GPU is enabled on the machine.

The well-known `numpy` package defines "NumPy arrays" to store vectors, matrices and tensors. Package `tensorflow` also defines counterparts "TensorFlow tensors" for the same purpose. NumPy arrays and TensorFlow tensors can be converted from one to the other. A key difference of the two is that TensorFlow tensors related calculations are executed on GPU wherever possible, making it more efficient when comes to large-scale calculation that can be paralleled. An example is given below.

```
x = np.array([1, 2, 3, 4, 5])
y = tf.convert_to_tensor(x, dtype=tf.float64)
x = x*0.3 // cpu calculation
y = y*0.3 // gpu parallel calculation
```

12.2.2   Classification and Regression

Classification

Regression

The following data frame *kaggle.com/datasets/shree1992/housedata* is used in this example to as a demonstration to predict house pricing using regression model.

The following class `SimpleRegressionModel` serves as an example that can be used for the above task. Notice that this model design is only a demonstration, and it is not optimized.

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import MinMaxScaler
import tensorflow as tf
import seaborn as sns

class SimpleRegressionModel:
        """
        SimpleRegressionModel trains and tests a regression model using
            the given data frame.
        """
        def __init__(self):
                self.num_input = None
                self.num_output = None
                self.scalar = None
                self.num_training = None
                self.X_train = None
                self.Y_train = None
                self.num_validation = None
                self.X_val = None
                self.Y_val = None
                self.history = None
                self.num_test = None
                self.X_test = None
                self.Y_test = None
                self.model = None
        def import_dataset(self, df_total: pd.DataFrame, iplist: list,
            oplist: list, validation_size: float, test_size: float):
        total_X = df_total[df_total.columns.intersection(iplist)]
        object_columns = total_X.select_dtypes(include=['object'])
        if not object_columns.empty:
            dummy_columns = pd.get_dummies(object_columns)
            total_X = pd.concat([total_X.drop(object_columns, axis=1),
                dummy_columns], axis=1)
        self.num_input = len(total_X.columns)
        total_Y = df_total[df_total.columns.intersection(oplist)]
        self.num_output = len(total_Y.columns)
        train_val_X, self.X_test, train_val_Y, self.Y_test =
            train_test_split(total_X, total_Y, test_size=test_size,
            random_state=None)
        self.num_test = len(self.X_test.index)
                if validation_size == 0:
                        self.X_train = train_val_X
                        self.Y_train = train_val_Y
                        self.X_val = []
                        self.Y_val = []
                        self.num_training = len(self.X_train.index)
                        self.num_validation = 0
                else:
                        self.X_train, self.X_val, self.Y_train, self.
                            Y_val = train_test_split(train_val_X,
```

```python
                    train_val_Y, test_size=validation_size/(1-
                    test_size), random_state=None)
                self.num_training = len(self.X_train.index)
                self.num_validation = len(self.X_val.index)
        print("Dataset␣size:␣training:␣{},␣validation:␣{},␣test:
            ␣{}".format(self.num_training, self.num_validation,
            self.num_test))
        self.scalar = MinMaxScaler()
        self.X_train = self.scalar.fit_transform(self.X_train)
        if validation_size == 0:
                pass
        else:
                self.X_val = self.scalar.transform(self.X_val)
        self.X_test = self.scalar.transform(self.X_test)
    def design_model(self, hidden_layer_model: list, optimizer: str,
        learning_rate: float, loss: str):
        """
        The input model describes the design of the model. It is
            a list of layers. Each layer is given by a
            dictionary describing layer type, number of nodes,
            etc.
        """
        self.model = tf.keras.Sequential()
        for ind in range(len(hidden_layer_model)):
                if ind == 0:
                        if hidden_layer_model[ind]["type"] == "
                            dense":
                                layer = tf.keras.layers.Dense(
                                    hidden_layer_model[ind]["node"
                                    ], activation =
                                    hidden_layer_model[ind]["
                                    activation"], input_shape = (
                                    self.num_input, ))
                        else:
                                pass
                else:
                        if hidden_layer_model[ind]["type"] == "
                            dense":
                                layer = tf.keras.layers.Dense(
                                    hidden_layer_model[ind]["node"
                                    ], activation =
                                    hidden_layer_model[ind]["
                                    activation"])
                        else:
                                pass
                self.model.add(layer)
        self.model.add(
                tf.keras.layers.Dense(self.num_output, activation
                    ='relu')
```

```
            )
            if optimizer == 'adam':
                    optimizer = tf.keras.optimizers.Adam()
            self.model.compile(optimizer=optimizer, loss=loss,
                metrics=['mae'])
    def train_model(self, epochs: int, batch_size: int):
            self.history = self.model.fit(self.X_train, self.Y_train
                , epochs=epochs, batch_size=batch_size,
                validation_split=0.2, verbose=2)
    def evaluate_model(self):
            if self.num_validation == 0:
                    print("Validation␣set␣is␣empty.")
            else:
                    loss, mae = self.model.evaluate(self.X_val, self.
                        Y_val, verbose=2)
                    print("Validation␣set␣test␣result:␣loss={},␣mae
                        ={}".format(loss, mae))
    def test_model(self):
            loss, mae = self.model.evaluate(self.X_test, self.Y_test
                , verbose=2)
            print("Test␣set␣test␣result:␣loss={},␣mae={}".format(
                loss, mae))
```

The following code uses the above defined class to predict house price.

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
import tensorflow as tf
import seaborn as sns

df_house_pricing = pd.read_csv("house_pricing.csv").dropna()
srm = SimpleRegressionModel()
srm.import_dataset(
    df_total=df_house_pricing,
    iplist=["bedrooms", "bathrooms", "sqft_living", "sqft_lot", "floors
        ", "sqft_above", "sqft_basement", "yr_built", "yr_renovated", "
        condition", "city"],
    oplist=["price"],
    validation_size=0, test_size=0.2)
hidden_layer_model = [
        {
                "type": "dense",
                "node": 128,
                "activation": "relu"
        },
        {
                "type": "dense",
                "node": 64,
```

```
                "activation": "relu"
        },
        {
                "type": "dense",
                "node": 64,
                "activation": "relu"
        },
        {
                "type": "dense",
                "node": 32,
                "activation": "relu"
        },
        {
                "type": "dense",
                "node": 16,
                "activation": "relu"
        },
]
srm.design_model(hidden_layer_model=hidden_layer_model, optimizer='adam
    ', learning_rate=0.001, loss='mse')
srm.train_model(epochs=100, batch_size=50)
srm.evaluate_model()
srm.test_model()
```

The above example is self-explanatory. Some key components of the codes are:

- Use `train_test_split()` from `sklearn.model_selection` to split the data set into training set, validation set and testing set.

- Use `MinMaxScaler()` from `sklearn.preprocessing` to normalize the inputs to the AI model.

- Use `tf.keras.Sequential()` and `tf.keras.layers` to design the AI model structure.

- Use `.compile()` to configure optimization engine, loss function, validation matrix, etc., during the training.

- Use `.fit()` to train the model using the training set.

- Use `.evaluate()` to evaluate the AI model performance.

- Use `.predict()` to carry out prediction of input points.

Use `.save("name")` to save a model, and `load_model()` from `tensorflow.keras.models` to load a model. Some popular formats to store a model include H5 file.

12.2.3   Computer Vision

"nobreak

12.2.4 General Sequential Data Processing

"nobreak

12.2.5 Natural Language Processing

"nobreak

12.2.6 TensorFlow on Different Platforms

"nobreak

12.3 PyTorch

TensorFlow is another open-source software library for machine learning originally developed by Meta AI in 2016. It is now under the Linux foundation umbrella.

12.3.1 PyTorch Basics

"nobreak

12.3.2 Classification and Regression

"nobreak

12.3.3 Computer Vision

"nobreak

12.3.4 General Sequential Data Processing

"nobreak

12.3.5 Natural Language Processing

"nobreak

12.3.6 PyTorch on Different Platforms

Part IV

Semantic Web

13

Semantic Web Basics

CONTENTS

Ontology is the philosophical study of the nature of being, existence, or reality. In the context of the Semantic Web, ontology refers to a formal specification of concepts, relationships, and categories of a particular domain or knowledge base. Ontology allows the creation of a shared understanding of a particular domain, enabling machines to interpret and reason about data more effectively. Ontology is essentially about "what is everything". It is about how we as humans mark our existence, and how we preserve and add to our knowledge base for the generations to come.

Ontology inspires people in computer science how we can store and exchange information more efficiently by building an internet based knowledge base. The solution is called the semantic web. An internet powered by semantic web (together with other technologies) defines web 3.0.

13.1 Web of Data

The Internet, as we call it today, is a technology that enables information exchange between machines, as well as between machines and human beings. With the power of the Internet, humans can obtain the information they need from remote servers, databases, or knowledge bases.

Initially, this information exchange was difficult, and professional computer skills were required to operate the Internet and its human-machine interface. Nowadays, everyone can access the Internet by just using a graphical browser. Public knowledge bases like Wikipedia make obtaining information much easier.

Following this trend, AI knowledge bases combined with human-readable semantic web technologies such as RDF and OWL, will become more and more popular and provide additional help for knowledge transfer and sharing.

13.1.1 Web 1.0 and 2.0

In the age of Web 1.0, knowledge was stored on individual servers, often in a static manner. Users could open a browser and check the knowledge as it was presented in the form designed by the server owner. It was essentially a one-way data transmission. The experts provide information, and the ordinary users consume the information.

In the age of Web 2.0, thanks to more advanced communication technologies, users can interact with web service providers. They can search and filter data from the servers and even upload their own data to the servers and share it with others. The Internet became easier and more flexible to use for information exchange, and data transfer became bidirectional. Ordinary users are both providers and consumers of information on the Internet. Databases are used in the backend which allows the user to push and pull data from the server.

Web 2.0 is useful but still far from perfect, especially given the background of industry 4.0 and data-driven everything.

One of the biggest challenges that we encounter with Web 2.0 is how to quickly locate the information we want hidden among the vast amount of data on the Internet. Among all the information available, which is important, and which is redundant? Humans have contextual knowledge and experience to solve the problem, but it will take forever.

Keyword search engines are widely used in the Web 2.0 framework, and it is the main method one can rely on when pulling information from the Internet. Keyword search engines may struggle with polysemous, synonyms, unintended personalizing, and implicit information from pictures.

Today most of the information on the Internet is stored in HTML format. HTML tells the contents links to the information, but not the meaning behind

it. It is difficult for a machine to understand the meaning of the information from HTML files. As will be introduced later, this potentially makes it difficult for a machine to retrieve data efficiently.

13.1.2 Web 3.0

The ultimate goal of Web 3.0 is to allow more efficient information search and transfer using the Internet, given the vast amount of existing data hosted on distributed servers and sensors.

In order to search, exchange, and abstract useful information from distributed sources in an efficient manner, there are at least two things we can do:

- Let the machine do the job. The AI should be smart enough to precisely capture what the user wants and get back to him with useful and accurate results.

- Make the information stored on the Internet more readable for both humans and machines, i.e., make the information semantic (meaningful).

The above forms the fundamentals of Web 3.0. The first point leads to smart chatbot, and the second to semantic web.

Web 3.0 relies on advanced data storage, data searching, and data processing technologies. We need to develop powerful AI so that it can interpret the user demands, search for useful information in its knowledge base efficiently, and finally come back with a human-readable answer. For that, we need advanced natural language processing techniques, and efficient (and potentially online) AI training techniques to keep the sophisticated knowledge base up to date.

In addition, we need to formally express the semantics (meanings) of the information for everything put online. This helps both humans and machines to interpret information and decide whether it is useful. This introduces the concept of the semantic web, which is a significant aspect of Web 3.0. In fact, semantics are so important that they are often considered synonymous with Web 3.0, although the broader vision of Web 3.0 includes other features as well.

Distributed storage has both advantages and disadvantages. On one hand, it provides higher resilience against data loss. On the other hand, it adds difficulty to the searching and collecting of information. Technologies like IPFS, blockchain, and distributed databases are being used to address these challenges.

This notebook focuses on the technologies used in semantic web, some of which such as RDF and OWL have already been implemented. We will see how we can use semantic web to collect and organize information, and from where how we can create a knowledge base for both humans and AI.

13.1.3   Semantic Web Vision

How to make contents on the Internet meaningful to machines?

Recall the 2 bullet points in the previous section. We always have the choice to use AI with powerful natural language processing capability. This of course relies on technologies in machine learning. We need to design ANN models that can be easily massively trained and it needs to be able to capture long text that humans usually use. With the development of transformer, this is becoming promising. There is another challenge though, which is that the knowledge base formed by AI is a black box, and cannot be interpreted by other machines with a different framework. In other words, you cannot build a knowledge base with AI and use it everywhere without any tuning. Transfer learning and knowledge distillation may help, but they also have constraints.

In the semantic web vision, we mainly focus on another approach, which is **annotating explicit semantic metadata to the existing content**. The semantic metadata shall be encoded the meaning of the content in a machine-readable way.

Semantic web approach is complementary with AI approach. It has some advantages:

- The knowledge is more transparent, unlike AI where the knowledge is hidden in the ANN as a black box.

- The knowledge is reusable and scalable in a easier way than AI knowledge base which relies on knowledge distilling (reuse) and transfer learning (rescale).

- RDF and OWL, the tools we use to organize and store knowledge in semantic web, are more reader-friendly for both humans and machines.

Many research facilities and organizations have tried building semantic web for either global or specific domains. An example is DBPedia *dbpedia.org*. The goal of DBPedia is to create something like Wikipedia, but semantic.

13.1.4   Semantic Web Stack

A commonly seen semantic web stack looks like Fig. 13.1. It is a summary of components and tools used in building a semantic web.

In the very bottom is the platform that hosts everything. It is also the representation of all the existing knowledge. Uniform Resource Identifier (URI) is the "address" that can be used to identify a information source. It can be in the form of an URL to the web page that introduces an item. For example, "About: Semantic Web (dbpedia.org)" *dbpedia.org/page/Semantic Web* is the URL to the defination of semantic web in DBPedia. The web page represents semantic web as part of the existing knowledge (this web page is not stand alone though, as later we will see that it links to other sources using linked data).
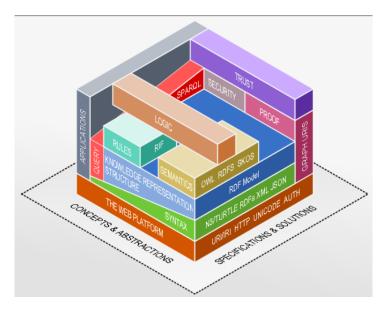
FIGURE 13.1
Semantic web stack introduced in [1].

N3, TURTLE, XML, JSON and Resource Description Framework (RDF) model are used to organize and store data. N3, TURTLE, XML and JSON is the syntax of the text, while RDF demonstrates the knowledge representation structure. The following is an example of using semantic web stored in TURTLE. Only the first few lines are given. Details to the syntax are introduced in later chapters.

```
@prefix dbo:   <http://dbpedia.org/ontology/> .
@prefix dbr:   <http://dbpedia.org/resource/> .
dbr:Semantic_web      dbo:wikiPageWikiLink dbr:Semantic_Web ;
      dbo:wikiPageRedirects dbr:Semantic_Web .
dbr:Knowledge_management      dbo:wikiPageWikiLink dbr:Semantic_Web .
dbr:Intelligent_agent dbo:wikiPageWikiLink dbr:Semantic_Web .
dbr:Wiki_software      dbo:wikiPageWikiLink dbr:Semantic_Web .
dbr:List_of_Brown_University_alumni dbo:wikiPageWikiLink  dbr:
    Semantic_Web .
dbr:NitrosBase dbo:wikiPageWikiLink dbr:Semantic_Web .
dbr:Arabic_Ontology   dbo:wikiPageWikiLink dbr:Semantic_Web .
dbr:Semantic_Web_Stack dbo:wikiPageWikiLink dbr:Semantic_Web .
dbr:Giant_Global_Graph dbo:wikiPageWikiLink dbr:Semantic_Web .
dbr:Computational_semantics  dbo:wikiPageWikiLink  dbr:Semantic_Web .
dbr:Deborah_McGuinness dbo:wikiPageWikiLink dbr:Semantic_Web ;
      dbo:academicDiscipline dbr:Semantic_Web .
```

Information is tagged in RDF. RDF, together with RDF Schema (RDFS)

provides limited semantics, and it created a foundation of semantic that allows Web Ontology Language (OWL) to build more complicated semantics on top of it, such as adding rules on restrictions and disjoint classes. In this sense, OWL is an extension of RDF.

SPARQL (SPARQL Protocol and RDF Query Language) is the standard query language for querying and manipulating RDF data on the Semantic Web. It allows users to search, retrieve, and modify information stored in RDF format.

The syntax of SPARQL is similar to SQL, as both are query languages designed for structured data. SPARQL was influenced by SQL, and many concepts in SPARQL have counterparts in SQL, such as SELECT, WHERE, GROUP BY, and ORDER BY. However, the backend technologies differ largely.

SQL is designed to query and manipulate data stored in relational databases, which are based on tables with rows and columns. On the other hand, SPARQL is designed for querying and manipulating RDF data, which is based on graphs with nodes and edges. In SPARQL, you work with RDF triples and graph patterns, and use operations such as graph pattern matching to find data that satisfies specific conditions. Their interfaces look similar, but the backend technologies differ largely.

An example of a knowledge base using RDF/RDFS/OWL is given below, together with a SPARQL query.

Example: Knowledge Base on Animals

The following examples demonstrates how to use RDF/RDFS and OWL to create a small knowledge base about animal. The example is provided by *ChatGPT-4*.

RDF/RDFS

```
@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
@prefix ex: <http://example.org/> .

ex:Animal rdf:type rdfs:Class .

ex:Mammal rdf:type rdfs:Class ;
    rdfs:subClassOf ex:Animal .

ex:Reptile rdf:type rdfs:Class ;
    rdfs:subClassOf ex:Animal .

ex:hasLegs rdf:type rdf:Property ;
    rdfs:domain ex:Animal ;
    rdfs:range rdfs:Literal .
```

```
ex:Dog rdf:type rdfs:Class ;
    rdfs:subClassOf ex:Mammal .

ex:Lizard rdf:type rdfs:Class ;
    rdfs:subClassOf ex:Reptile .

ex:Max rdf:type ex:Dog ;
    ex:hasLegs 4 .

ex:Lizzy rdf:type ex:Lizard ;
    ex:hasLegs 4 .
```

In this example, we used RDF and RDFS to:

- Define classes (Animal, Mammal, Reptile, Dog, and Lizard)

- Define a property (hasLegs)

- Set the domain and range of the property

- Create subclass relationships (Mammal and Reptile are subclasses of Animal; Dog is a subclass of Mammal; Lizard is a subclass of Reptile)

- Define individuals (Max and Lizzy) and their properties

OWL

```
@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
@prefix owl: <http://www.w3.org/2002/07/owl#> .
@prefix ex: <http://example.org/> .

ex:hasParent rdf:type owl:ObjectProperty ;
    rdfs:domain ex:Animal ;
    rdfs:range ex:Animal .

ex:hasChild rdf:type owl:ObjectProperty ;
    owl:inverseOf ex:hasParent .

ex:isWarmBlooded rdf:type owl:Class ;
    rdfs:subClassOf ex:Animal .

ex:Mammal rdfs:subClassOf ex:isWarmBlooded .

ex:Reptile owl:disjointWith ex:isWarmBlooded .

ex:Max ex:hasParent ex:Buddy .
```

```
ex:Buddy rdf:type ex:Dog ;
    ex:hasLegs 4 .
```

In this example, we used OWL to:

- Define object properties (hasParent and hasChild)

- Specify an inverse relationship between properties (hasParent and hasChild)

- Define a new class (isWarmBlooded) and set it as a superclass of Mammal

- Specify a disjoint relationship between Reptile and isWarmBlooded

- Define a new individual (Buddy) and his properties

- Specify a relationship between individuals (Max and Buddy)

In this demonstration example, RDF and RDFS provided the basic structure and hierarchy for the knowledge base, while OWL added more expressivity by defining complex relationships, additional semantics, and constraints.

Finally, the above knowledge base can be queried by SPARQL as follows.

```
PREFIX ex: <http://example.org/>

SELECT ?mammal
WHERE {
  ?mammal rdf:type/rdfs:subClassOf* ex:Mammal .
}
```

SPARQL query is not user-friendly. Commercialized semantic web applications such as WolframAlpha *wolframalpha.com* often provides a "search bar" with limited NLP capability to make the searching of the semantic web easier. It is worth mentioning that although this searching engine may look like chatbot at a first glance, their underlying information source and techniques differently largely. See ChatGPT-4's comments as follows.

I understand that chatbot like you works differently compared with these engines such as WolframAlpha. While their knowledge comes from semantic web, your knowledge comes from years of training using natural-language-written books and articles. Am I right?

Yes, that's correct. As an AI language model, I am based on the GPT-4 architecture developed by OpenAI. My primary source of knowledge

comes from the vast amount of text data that I have been trained on, which includes books, articles, and other documents written in natural language. My training data consists of diverse sources and covers a wide range of topics up to my knowledge cutoff date in September 2021.

Unlike search engines like WolframAlpha or tools that rely on the Semantic Web, my knowledge representation and understanding are not based on formalized structures like RDF or OWL. Instead, I analyze and generate text based on patterns and relationships learned from my training data.

This approach has some advantages and limitations. On the one hand, I can provide more human-like, natural language responses and can generate text on a wide range of topics, even if I have not been specifically trained on a particular domain. On the other hand, my understanding of the underlying structure and semantics of the data may be less precise than that of systems that use formal knowledge representation languages like RDF or OWL.

In summary, chatbots like me and engines like WolframAlpha operate differently in terms of their knowledge sources and representation, but both approaches can be useful for providing information and answering questions in their respective domains.

Chatbot and semantic web can work together. Chatbot is used to provide a human interface and basic domain knowledge, while the semantic web provides up-to-date precise informative facts.

13.1.5   Semantic Web Limitations and Challenges

It is worth mentioning that though semantic web is powerful, building semantic web can be challenging because it requires a lot of human labor. As of 2023, the majority of web sites and applications have not yet embraced the full potential of the semantic web, some of which only partially adopt semantic web concepts or technologies.

ChatGPT-4 comments on the current status of semantic web applications as follows.

The semantic web has made significant progress since its inception, but it hasn't been adopted as widely as initially envisioned. Many organizations, research institutions, and some large corporations have adopted semantic web technologies for specific use cases, such as knowledge management, data integration, and domain-specific applications. However, the majority of websites and web applications have not yet embraced the full potential of the semantic web.

That being said, some aspects of the semantic web have become popular and are widely used:

1. Structured data markup: Many websites use structured data markup, such as schema.org, to provide additional semantic information about their content. This helps search engines like Google to better understand the content and present it more effectively in search results. While not the full realization of the semantic web vision, this is a step towards making the web more machine-readable.

2. Linked Data: The concept of Linked Data, which is closely related to the semantic web, has gained some traction. Linked Data is a method of publishing structured data on the web so that it can be interlinked and easily accessed by machines. Examples include DBpedia, Wikidata, and various governmental and organizational datasets.

3. Research and academia: The semantic web has been a popular research topic in academia, with numerous conferences, workshops, and publications dedicated to its development and application. This has led to the development of various tools, algorithms, and methodologies that have been adopted in specific domains.

4. Domain-specific applications: Some industries and sectors, such as healthcare, life sciences, and digital libraries, have adopted semantic web technologies to manage their complex data and facilitate data integration and interoperability.

While the semantic web has not been adopted as widely as originally anticipated, its core ideas and technologies have had an impact on various areas of the web, and its principles continue to influence the development of new standards and technologies. It's possible that as AI and machine learning technologies continue to advance, there will be renewed interest in the semantic web and its potential to facilitate machine understanding of web content.

There are several challenges that have contributed to the semantic web not being embraced to its full potential:

1. Complexity: Semantic web technologies, such as RDF and OWL, can be complex and difficult to understand for many web developers, especially those without a background in knowledge representation, logic, or artificial intelligence. This complexity can act as a barrier to adoption.

2. Lack of incentives: Many website owners and content creators may not see immediate benefits or incentives for adopting semantic web technologies. Implementing semantic annotations and structured data can be time-consuming and resource-intensive, especially for smaller organizations or individuals without dedicated technical teams.

3. Incomplete coverage: The vision of the semantic web relies on a large amount of high-quality, machine-readable data being available. However, not all websites provide structured data or semantic annotations, which limits the effectiveness of semantic web applications and discourages further adoption.

4. Ontology development and maintenance: Creating and maintaining ontologies, which are key components of the semantic web, can be a complex and time-consuming task. Achieving consensus on common vocabularies and ontologies across different domains and stakeholders can also be challenging.

5. Scalability and performance: Semantic web technologies need to handle large-scale data processing and reasoning tasks, which can be computationally expensive. Some traditional semantic web tools and algorithms may not scale well for big data or real-time applications.

6. Privacy and security concerns: The semantic web promotes data sharing and integration, which can raise privacy and security concerns for both individuals and organizations. Ensuring that sensitive data is protected while still enabling the benefits of the semantic web can be a challenging balance to strike.

7. Competition from other technologies: Other technologies and approaches, such as keyword-based search engines, relational databases, and machine learning, have continued to evolve and provide effective solutions for many data management and retrieval tasks, which may have reduced the urgency for adopting semantic web technologies.

Overcoming these challenges will require continued research, development, and collaboration among researchers, developers, and stakeholders. As new tools, algorithms, and techniques are developed, and as the importance of machine understanding of data increases, there may be renewed interest in and adoption of the semantic web.

Despite the challenges in the widespread adoption of the semantic web, its core ideas and principles have had a significant impact on various fields, especially in the development of domain-specific knowledge bases and applications. The concepts of ontology, structured data, and linked data have inspired and facilitated numerous projects and use cases, such as:

1. Domain-specific ontologies: The development of domain-specific ontologies has allowed organizations and researchers to model and represent complex knowledge structures and relationships in their respective fields. Examples include the Gene Ontology (GO) in the life sciences, the CIDOC Conceptual Reference Model (CRM) in cultural

heritage, and the Financial Industry Business Ontology (FIBO) in finance.

2. Data integration and interoperability: Semantic web technologies have been instrumental in integrating and sharing data across different systems, platforms, and organizations. By utilizing shared ontologies and vocabularies, data can be mapped and linked more easily, enabling better interoperability and reducing the need for manual data reconciliation.

3. Knowledge management and discovery: Semantic web technologies have been used to build intelligent knowledge management systems that can infer new information and discover hidden relationships by reasoning over structured data. This can help organizations make better decisions, discover new insights, and optimize their processes.

4. Personalization and recommendation systems: By incorporating semantic web technologies and leveraging domain-specific knowledge bases, recommendation systems can provide more accurate and personalized suggestions based on users' preferences, behaviors, and the underlying relationships between items.

5. Natural language processing and AI: The ideas behind the semantic web have inspired research in natural language processing (NLP) and AI, as they both strive to enable machines to better understand and process human-generated content. Semantic web technologies can be used to enrich NLP tasks and enhance AI models with structured knowledge.

While the semantic web has not yet reached its full potential, its principles and concepts continue to influence and inspire the development of new technologies and applications, especially in the context of building domain-specific knowledge bases and fostering knowledge sharing and discovery.

With the population of LLM powered chatbot and IoT as part of Industrial 4.0, things might change. It has been extremely difficult to check the correctness and consistency of the existing data bases (both structured and unstructured) due to the large volume of data, let along converting them into semantic web. With the help from chatbot, the checking and conversion might be automated. In addition, as Industry 4.0 and IoT technologies continue to advance, the amount of structured data generated will increase significantly. This structured data can be a valuable resource for building semantic web applications more easily and effectively. A knowledge base that combines semantic web and chatbot might become a future standard. ChatGPT-4 comments as follows.

In the future, the combination of AI chatbots and semantic web technologies could become a powerful standard for handling various tasks. AI chatbots would serve as the natural language interface for users, providing basic knowledge and understanding of context, while the semantic web would provide accurate, up-to-date, and domain-specific data.

Integrating AI chatbots with SPARQL or other semantic web querying capabilities can enable them to search, analyze, and interpret structured data more effectively. This way, chatbots could provide users with more precise and relevant information based on the available structured data.

Such a combination of technologies would not only make it easier for users to access and interact with the vast amount of structured data, but it would also enable AI systems to better understand the context and semantics of information, leading to improved performance and more accurate results.

13.2 Ontology

Ontology has very rich meanings with both philosophy and semantic web perspective.

13.2.1 Philosophy Perspective

Knowledge is the intersect of ground truth and human beliefs, i.e., it is the truth that humans know of being truth. Ontology is a way of communicating knowledge.

In the context of philosophy, ontology discusses the meaning of an object "exists", how objects can be categorized, and how objects relate to each other. Ontology mainly discusses the following questions:

- What is existence?

- What are the fundamental categories of existence?

- What does it mean for something to exist or not exist?

- What types of entities exist, and what is their nature?

- What is the nature of abstract entities like numbers, properties, and relations?

- How do different entities relate to each other and interact?

- Can something exist independently of our perception or thought?

The discussion of these questions can be traced back to BC300 and earlier. Aristotle defined a system to structure and reasoning knowledge. The famous Aristotelian logic "major premise + minor premise -¿ conclusion" is in fact a standard and systematic way of reasoning unknown from knowledge. Aristotelian logic is an important tool of ontology, and it inspires how we build semantic web, even after these many years.

The first idea of creating a machine to categorize things came in 1200s. The first idea of creating a "universal precise language" to record knowledge was proposed in 1600s.

13.2.2   Semantic Web Perspective

In the context of the Semantic Web, an ontology is a formal, machine-readable representation of a specific domain's concepts, their properties, and the relationships between them. It serves as a shared vocabulary (unlike natural language, which is often ambiguous) for describing and reasoning about the knowledge within that domain. In practice, RDF/RDFS and OWL can be used to express the ontology.

The followings are defined in RDF/RDFS and OWL.

- Class. Classes are used to represent components and models. It represents a concept. It is an abstraction of objects sharing some similarities.

- Attribute. Attributes are used to describe classes.

- Relations. Relations are special attributes whose values are objects of (other) classes.

- Constraints. The attributes of classes might be restricted. For example, a class may have an multiple attributes of similar kind (consider a person has many houses). A class may not have two attributes at the same time (consider a person cannot be man and woman at the same time).

- Instance. An instance is an individual of an ontology. For example, "John is a person", where "person" is a class.

13.2.3   Knowledge and Logic

Different knowledge frameworks and logic types are briefly introduced in this section. They are covered in more details in later sections, including Section 13.3 for logic frameworks, logic equivalence, and logic reasoning, Section 13.4 for description logic (DL) and attributive language with complement (ALC), and Section 14.5 for OWL.

Knowledge should be stored in such a way that new knowledge can be derived from existing knowledge, and there should be a standardized way of

deriving new knowledge from existing knowledge. This introduces the important concept of logic.

Propositional logic is the fundamental of the logic that we know of today. In propositional logic, knowledge is represented by simple facts (simple proposition), and facts connected with AND, OR, NOT, IF THEN, IF AND ONLY IF (compound proposition).

First-order logic (FOL) is the most commonly used logic as of today. In FOL, quantifiers and logic connectives are introduced, including universal quantification $\forall$, existential quantification $\exists$, conjunction, disjunction and negation. A formal way of representing logic expressions are defined. For example, using the following to represent "all humans are mortal"

$$\forall x \, (\text{Human}(x) \rightarrow \text{Mortal}(x))$$

where $\text{CLASS}(x)$ is equivalent of a proposition "x belongs to CLASS", which can be either true or false. The above proposition says "for any item, if that it belongs to human is true, then that it belongs to mortal must be true". Given the current biology development, this proposition is still true.

Description logic (DL) is a subset of the first-order logic. DL is specifically designed to represent and reason about ontologies and their concepts, relations, and instances. It focuses on capturing the semantics of structured knowledge with a restricted set of constructs and reasoning capabilities. The idea of semantic web builds on top of DL. Some key features of DL are:

- Defining classes and subclasses

- Defining properties (roles) and their domain and range:

- Defining restrictions on classes

- Using existential and universal quantifiers

Semantic web is essentially an implementation of DL on computers. As can be seen later, many tools used in semantic web, such as RDF/RDFS and OWL, are essentially practicing DL in a machine-readable way.

Attributive language with complement (ALC) is one of the ways to describe a simple DL, and it forms an important subset of OWL which enables ontology in semantic web. Understanding ALC is helpful with learning OWL in later sections. A brief introduction of ALC is given below.

"Concept" (corresponding with class in RDF/RDFS) and "role" (corresponding with property in RDF/RDFS) are introduced. Top concept and bottom concepts are defined. Each property is associated with a range, which is a basically a set of values that the property can take. Concepts and roles are atomic types defined in ALC.

Constructors are used to describe ranges and their relationship with a property. Conjunction/Intersection, disjunction/union and negation/complement of a range are defined. Existential and universal quantifiers are defined.

For example, `\forall R.C` can be used to say that all attributes of property type "R" should stay within range "C". Similarly, `\exists R.C` means that there should be at least one attribute of property type "R" whose value stays within range "C".

Class relations are defined. Commonly used class relation constructors are inclusion (to describe sub class) and equality (to assign class), union, intersection, complement, etc.

Knowledge is expressed in two formats, either being terminological knowledge, or assertional knowledge. As a terminological knowledge example, we can define a teacher as

$$\text{Teacher} \equiv \text{Person} \wedge \exists \text{ HasStudent.Student} \wedge \exists \text{ Teaches.Class}$$

which translates to "a teacher is a person, and has relationship HasStudent with at least one Student element, and has relationship Teaches with at least one Class element, and "Person", "Student", "Class" are classes, "HasStudent", "Teaches" properties, defined in RDF/RDFS. There are people who teaches only tutorials but not classes. To include them into the teacher class, consider using

$$\text{Teacher} \equiv \text{Person} \wedge \exists \text{ HasStudent.Student} \wedge (\exists \text{ Teaches.Class} \vee \exists \\ \text{Teaches.Tutorial})$$

It can be seen that with terminological knowledge, a lot of restrictions and rules can be added when defining a class.

As an assertional knowledge, we can state an instant or a sub class to be of Teacher as follows.

```
Teacher(Peter)
Teaches(Peter, LinearAlgebra)
```

All the above can be translated into OWL, which largely enriches the capability of RDF model building relations and restrictions among classes. More details are given in later sections.

13.2.4　Ontology Types and Categories

There are different layers of ontology. The higher the layer, the more general it is. The lower the layer, the more specific it is to a particular research domain, a particular application, or even a particular task. An example is given in Fig. 13.2. In practice, there might not be a clear boundary between two adjacent ontology layers.

Lower-layer ontology can import concepts and knowledge from the upper-layer ontology. If needed, it can modify the imported concepts.

Other ways of leveling ontology include using the expressiveness. The "light-weight" ontology is informal, less semantic, and supports less comprehensive logic. The "heavy-weight" ontology, on the other hand, is formal, more semantic, and supports more comprehensive logic and reasoning up to first-order logic. The vocabulary also grows with the ontology complexity.
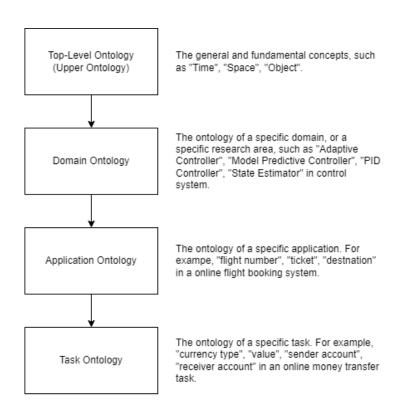
FIGURE 13.2
A demonstration of ontology level.

13.3 Logic

Logic and logical expression are by themselves a research area. Details on logical expression is not included in this section. Only a brief scratch is given.

Humans are good at deriving new knowledge from semantically defined existing knowledge. The tool that enables this is the logic inference. Logic is briefly introduced in this section. Logic helps to distinguish between truth and lie, so that we can accept truth and reject lie.

The idea "formal logic" systematically defines logic reasoning methods and procedures to automate logic inference. Logic inference engines or reasoners are used to automate this process of making logical conclusions based on the given knowledge.

13.3.1 Syntax and Semantics

We use "syntax" to represent "semantics". Syntax is just symbols without meaning that comes naturally. We humans give meaning to the syntax. The meaning of the syntax is its semantic.

Consider the following example. This is a piece of code written in Python. As anyone who has decent Python programming experience, the behavior of the code can be easily spotted, or even emulated in our brain. However, what is the semantics behind this syntax?

```
ax = []
for i in range(20):
      if i<=1:
              ax.append(i)
      else:
              ax.append(ax[i-1]+ax[i-2])
```

When considering the semantic understanding of the simple Python script above, it's interesting to note the differences between a Python interpreter, an AI model, and human cognition.

The Python interpreter is essentially syntax-driven. When it executes a script, it follows the provided instructions one by one without any broader understanding or anticipation of the results. In the context of our Fibonacci series, the interpreter doesn't recognize that it's generating a famous mathematical sequence or why that sequence might be important or interesting. It simply follows the rules of the script, calculating and outputting each number in the series as instructed.

On the other hand, an AI model like GPT-3 does possess a level of semantic understanding, but it's quite different from that of a human. Through training on a diverse range of internet text, it has learned to associate the Python script for generating a Fibonacci series with the mathematical concept of the Fibonacci sequence. This understanding is not innate but rather a result

of recognizing patterns in the data it was trained on. When asked about the script, it can provide a summary or explanation based on its learned associations, including the name of the series, and the purpose of each line of the code. It follows what it has been trained on.

Finally, there's the human level of semantics, which is by far the most advanced and intuitive. Not only can a human understand the Python script and the concept of the Fibonacci series, they can also infer its broader mathematical properties, such as its relation to the golden ratio, and its general behavior. Humans can understand that the series will converge to the golden ratio not only when starting with 0 and 1, but with any two initial positive integers. This level of understanding is a combination of learned knowledge, pattern recognition, and the ability to extrapolate or generalize from existing information.

One of the goals of semantic web is to enable machine to understand the semantic as much as possible.

The semantics can be categorized into different types.

- Intentional semantics refer to the information that the producer of the information wants others to interpret, usually in natural language.

- Formal semantics refer to expressing the semantic in a formal language.

- Procedural semantics refer to the meaning of a programming language expression.

- Model-theoretic semantics refer to the interpretation of natural languages by identifying meaning with an exact and formally defined model.

In the context of logics, model-theoretic semantics are the main interest of study.

13.3.2   Logical Expression

Any logic consists of a set of statements S and an entailment relation $\models$ (this can be an element or a set; in this example just take it as an element for simplicity). A logic

$$L := (S, \models)$$

shall claim $\Phi \subseteq S$ and $\phi \in S$, where ϕ is a logical consequence of Φ. If two logical assertions satisfy $\Phi \models \Psi$ and $\Psi \models \Phi$, then they are logically equivalent $\Phi \equiv \Psi$.

Interpretation (I) and formula (F) are two very important terms in logic expression. Interpretation is a formal construct that defines the meaning of a symbol in a formal language. For example, in the ontology of people, an interpretation can map an instance symbol, such as "Alice", to an actual person in the real world, and class symbol "Person", to the concept of a human being, etc.

Formula, on the other hand, is a statement formulated by string of symbols from a formal language. It can be a rule, or an assertion, trying to represent some fact. For example, a formula can be "Alice hasSibling Bob". Is it true or false? This depends on the interpretation. If "Alice" and "Bob" are indeed mapped to two siblings, and the relation "hasSibling" is literally what we understand it, then it is true.

Here is another example. Consider "10 is greater than 5". Intuitively, this is a universal truth. But in fact, from the interpretation and formula relationship perspective, this also depends on the interpretation. It really depends on what "10" and "5" map to. If "10" and "5" are interpreted as numerical quantities and "is greater than" as the standard numerical greater-than relation, then it is true. However, if "10" and "5" are interpreted as amounts of debt and "is greater than" as a relation indicating "is a better trader than" (with less debt being better), then it would be false under this interpretation.

Given a formula F alone, we cannot tell whether it is true or false. It is meaningful only if we discuss it with the interpretation I also given. For those interpretations that make the formula true, we call them the model of the formula, and we can denote $I(F)$ or $I \vDash F$, which read as "I models F" or "I satisfies F". Different interpretations of the same symbols can lead to different conclusions. With the above, it is possible to express FOL using formal language. For example, "all kids love icecream" can be represented by

$$\forall X : \text{Child}(X) \to \text{lovesIcecream}(X)$$

Multiple formulas can be grouped into a theory (T), which can be used interchangeably with F. A theory can be treated as a knowledge base.

13.3.3   Logical Equivalence

If two formula hold the same true and false value under the same interpretation, i.e., $F \vDash G$ and $G \vDash F$, then F and G are logically equivalent. The proof of logical equivalence can be done using the truth table.

There is a huge table of logical equivalence formulas. The proof of the formulas is not given here. Commonly seen equivalence is given in Table 13.1.

Canonical forms have been defined for logical expressions. There are at least 6 different canonical forms for a logical expression, namely conjunctive normal form, disjunctive normal form, Prenex normal form, Skolem normal Form, negation normal form and clausal normal form. Canonical forms are not necessarily the easiest and most intuitive forms of an expression (it is often the opposite, as a matter of fact), but somethings they are simple for further analysis and process, such as finding contradictions.

Notice that when deriving certain canonical forms from the original logical expression, information may get lost and they are not necessarily always equivalent, but they are usually.

TABLE 13.1

Numerical calculations.

Statement	Equivalent	Comment
$\neg\neg p$	p	Double negation law
$(p \wedge q)$	$(q \wedge p)$	Commutative laws
$(p \vee q)$	$(q \vee p)$	Commutative laws
$(p \wedge (q \wedge r))$	$((p \wedge q) \wedge r)$	Association laws
$(p \vee (q \vee r))$	$((p \vee q) \vee r)$	Association laws
$(p \wedge (q \vee r))$	$((p \wedge q) \vee (p \wedge r))$	Distributive laws
$(p \vee (q \wedge r))$	$((p \vee q) \wedge (p \vee r))$	Distributive laws
$(p \rightarrow q)$	$(\neg p \vee q)$	Conditional statements
$(p \rightarrow q)$	$(\neg q \rightarrow \neg p)$	Conditional Statements
$(p \leftrightarrow q)$	$((p \rightarrow q) \wedge (q \rightarrow p))$	Conditional Statement
$(\neg(p \wedge q))$	$(\neg p \vee \neg q)$	De Morgan's laws
$(\neg(p \vee q))$	$(\neg p \wedge \neg q)$	De Morgan's laws

13.3.4   Logical Reasoning

Logical reasoning is about deriving the true or false of a formula F, given a theory or a knowledge base T.

There can be many ways of doing the reasoning. It is worth mentioning in particular that there is a systematical way of proving the truth or false of a statement using the law of contradiction. For example, to prove that $\Phi \vDash F$, simply append $\neg F$ to Φ, and find a contradiction. Logic reasoning via law of contradiction is the fundamental to semantics. However, do notice that different logic framework such as propositional logic and FOL may have different practice details when using this method.

To look for contradictions, we need to simplify (resolve) the logical expressions. Here "resolution" refers to systematic procedures of simplifying logical expressions, and it plays an important role in looking for contradictions, hence in semantic reasoning. Resolution usually requires the logical expression to be given in the clausal form. Different logic framework may have different approaches for resolution.

Tableaux algorithm is another important algorithm for reasoning. Similar with resolution, Tableaux algorithm also checks the consistency of a logical expression, and looks for contradictions. Instead of working on clausal form, Tableaux works on disjunctive normal form.

13.4   DL and ALC

DL and ALC have been briefly introduced in Section 13.2.3. More details are given in this section.

13.4.1   Recap

Knowledge can be represented using formalisms with varying levels of logical expressiveness and complexity. FOL and DL are examples of highly expressive and complex formalisms. RDF/RDFS and rule-based expressions, while still grounded in logic, offer simpler, less expressive formalisms. Data-driven approaches such as ANN and statistical reasoning, represent a different category of knowledge representation that does not rely on the principles of formal logic. The more logically expressive and complex the formalism, the more formal the knowledge representation and the more capable it is of complex semantic reasoning.

We would surely want to introduce highly expressive and complex formalisms to the existing RDF/RDFS model. This motivates OWL, which enables DL in semantic web. More details of OWL is introduced later in Section 14.5. To build a solid foundation, DL and ALC are introduced here. Notice that DL instead of FOL is widely used in semantic web, mainly because FOL a bit over complicated. DL is a subset of FOL. ALC is one of the simplest way of expressing DL.

ALC defines atomic types such as concepts and roles. It defines constructors including negation, conjunction, disjunction, existential quantifier and universal quantifier to build up more sophisticated expressions. It defines class inclusion and class equivalence corresponding with interpretation and formula relationships. ALC uses terminological knowledge (TBox) and assertional knowledge (ABox) to form the knowledge base.

13.4.2   DL Inference and Reasoning

Before talking about knowledge in the knowledge base, we need to discuss what to return if knowledge is not in the knowledge base. In open world assumption (OWA), it is assumed that everything is possible when DL ontology is empty. As long as nothing can be found in the knowledge base that is against the formula, then it is probably true. On contrary, in closed world assumption (CWA), as long as nothing supports the formula, it is regarded as false.

In this perspective, when OWA is made, the knowledge base regards itself as still growing, and when knowledge is not recorded, it would response "unclear", "maybe". While when CWA is made, the knowledge base regards itself as completed knowledge base. If something is not recorded, then it does not exist.

Consider an example where it is asked "are Alice's children all males?", and in the database there are two children of Alice, both of which are male. Under OWA, the inference would be "maybe", as it does now know whether there are more children of Alice. While in CWA, the inference would return "yes", as it checks both children to be male, and it assumes they would be all the children Alice has. However, when Alice does have female child registered, both OWA and CWA would inference as "no", because the female child contradicts the assertion regardless of the completeness of the database.

The following is a summary of problems that can one can ask a knowledge base:

- Is the knowledge base consistent (is there conflicts)?

- Is class A empty?

- Is class A an inclusion of class B?

- Are class A and B equivalent?

- Are class A and B disjunctive?

- Is an individual contained in a class?

- Find all individuals in a class?

Intuitively, to answer the above inference problems, it is possible to use algorithms defined in FOL such as resolution and Tableaux algorithm. But a problem here is that FOL algorithms do not always terminate (which is another reason whey DL is preferable than FOL in semantic web, the former of which is less expressive, but more computationally manageable).

While FOL inference is not decidable, DL inference, on the other hand, is decidable. Consider using ALC tableaux algorithm to solve the aforementioned problems (there are other more complicated algorithms as well). The problems need to be first transformed into the associated unsatisfiable problems, and consequentially tableaux algorithm can be used to find if there is any contradiction.

The mechanism of implementing tableaux algorithm with ALC is neglected here.

13.4.3 DL in Semantic Web

It ultimately comes to how to implement DL in semantic web to enhance its capability of reasoning, and the answer is to introduce OWL. More details are given in Section 14.5.

14

Semantic Web Architecture

CONTENTS

Semantic web uses URI to identify an existing name or resource, RDF/RDFS and OWL to store new information, and SPARQL to query from the knowledge base. The semantic web stack is already given earlier in Section 13.1.4. More details are introduced in this chapter.

14.1   Uniform Resource Identifier (URI)

Humans uses symbol and concept to help interpret things and link to objects in real life. This is demonstrated by semiotic triangle as shown in Fig. 14.1. The interpretation of "apple" is used as an example. The concept is shared among humans in the form of knowledge.
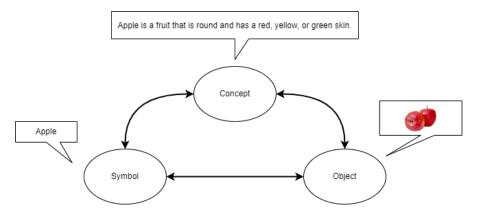
FIGURE 14.1
Semiotic triangle. The symbol is an "representation" of a category of items. The concept is the abstracted general properties of an item.

Uniform resource identifier (URI) is the way to identify a resource of information or knowledge. URI is unique from resource to resource.

For example, DBpedia gives concept of the apple here: *dbpedia.org/page/Apple*. This is an URL to the web page that contains information of the apple. This URL is an example of URI. In this example, the information can be accessed by HTTP call to the URL.

In semantic web, the URI of a resource identifies an item or a class. Using the URI, the representation (meta data) of the item/class can be accessed. The representation is like the symbol in Fig. 14.1, and although it is not a specific instance of the item, it is usually sufficient for us to use the symbol to discuss the item. This is shown in Fig. 14.2.

URI shall contain at least two pieces of information: the address (locator), and the identity (name). URI is often ASCII encoded, but it is possible to extend to unicode. The generic syntax is given below, and the URL is a subset of the syntax.

`scheme:[//authority]path[?query][#fragment]`

where

- `scheme` specifies the protocol used to access the resource. Common schemes
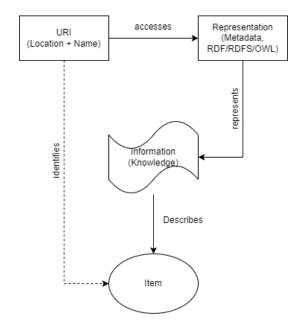
FIGURE 14.2
Identification flowchart.

include `http`, `https`, `ftp`, `mailto`, `file`, and `data`. The scheme is followed by a colon `:`.

- `authority` component is optional and typically includes the user information, host, and port, separated by an at sign `@` and a colon `:`, respectively. The authority is preceded by a double forward slash `//`.

- `path` identifies the specific resource within the context of the scheme and authority. It is a sequence of segments separated by forward slashes `/`.

- `query` component is optional and provides additional information that the resource can use for processing. It is a series of key-value pairs separated by an ampersand `&`. The query starts with a question mark `?`.

- `fragment` component is optional and allows for the identification of a specific part or section within the resource. It is typically used with HTML documents to indicate a specific anchor or location within the page. The fragment starts with a hash sign `#`.

More details can be found on

`https://www.rfc-editor.org/rfc/rfc3986.html#section-3.1`

which happens to be a good example of an URI in https scheme.

14.2 Resource Description Framework (RDF)

The representations of the knowledge shown in 14.2 need to be presented in a consistent way. In practice, RDF is introduced to represent knowledge.

Notice that RDF itself is not a programming language. It is the structure using which information is stored. Therefore, it needs a markup language to host the RDF framework. Commonly used markup language the following. There are tools to automatically transfer them from one to another.

- RDF/XML: has good compatibility with old machines.

- Terse RDF Triple Language (Turtle): easy to use, human-readable.

- JSON for Linked Data (JSON-LD): popular in web applications and APIs where JSON-based format is required.

- N-Triples: simple, machine-readable format for data exchange between tools.

Turtle is very self-explanatory. Unless otherwise mentioned, Turtle is used in this notebook.

14.2.1 Triple Representation

RDF uses "node-edge-node" triples to represent knowledge. The triple refers to the combination of:

- Subject

- Predicate

- Object

For example, consider representation of knowledge "Einstein was born in Ulm". In RDF, "Einstein" is the subject of the knowledge, and "Ulm" the object. The predicate is "has birthPlace", where "birthPlace" is a property assigned to "Einstein". The graph representation is given by Fig. 14.3. The

FIGURE 14.3
Graph representation of triple for knowledge "Einstein was born in Ulm".

RDF is given as follows.

```
@prefix dbo: <http://dbpedia.org/ontology/> .
@prefix dbr: <http://dbpedia.org/resource/> .

dbr:Albert_Einstein dbo:birthPlace dbr:Ulm .
```

where `Albert_Einstein` and `Ulm` are defined in `dbpedia.org/resource/` as name and place, and `birthPlace` in `dbpedia.org/ontology/` as an attribute, respectively. To give more details:

- `@prefix` keyword is used to define prefixes for namespaces, making it easier to write URIs.

- `dbr:Albert_Einstein` is the subject, representing Albert Einstein as a resource in the DBpedia namespace.

- `dbo:birthPlace` is the predicate, representing the "birthPlace" property from the DBpedia ontology.

- `dbr:Ulm`is the object, representing the city of Ulm as a resource in the DBpedia namespace.

- `.` indicates the end of a statement.

  We use

```
<object> <predicate> <subject> .
```

to claim a statement, and

```
<object> <predicate1> <subject1>; <predicate2> <subject2>; <predict3> <
    subject3> .
```

to assign multiple predicates to a subject to avoid repeating the same object in statements.

14.2.2 Multi-valued Relation and Blank Node

It is possible to use multi-valued relations and blank nodes to enforce combining of information. Consider an example given in Fig. 14.4, where the graph is used to demonstrate a lecture taking place at a specific room at given time slot. What if the lecture takes place twice a week, each time at a different location? With the help of multi-valued relations and blank nodes, this can be demonstrated clearly as shown in Fig. 14.5.

Turtle and other markup languages provides syntax for creating blank nodes that looks like the following.

```
<subject1> <predicate1> [
      <predicate2> <object2>;
      <predicate3> <object3>
] .
```
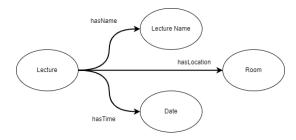
FIGURE 14.4
An example that demonstrate when and where a lecture takes place.
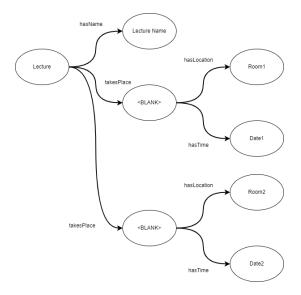
FIGURE 14.5
An example that demonstrate a lecture taking place at multiple locations and time slots using multi-valued relations and blank nodes.

To refer to a blank node, it is also possible to give the blank node a name. The syntax looks like the following.

```
<subject1> <predicate1> _:<blank-node-name> .
_:<blank-node-name> <predicate2> <object2> .
_:<blank-node-name> <predicate3> <object3> .
```

where `_:<blank-node-name>` is used to declare a blank node and assign it a name.

14.2.3   Lists

Lists help to make the code clean and readable. There are two types of lists, the container (open list, extendable) and the collections (closed list, fixed). Container is helpful to handle the situation given in Fig. 14.6. Notice that
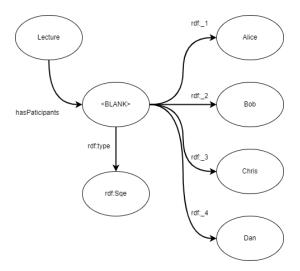
FIGURE 14.6
An example of a container.

the blank node has a type `rdf:Seq`. This tells that the items stored in the container follows an ordered set. There are other container types, such as `rdf:Bag` (unordered set) and `rdf:Alt` (alternatives of elements; only one element is relevant for the application).

The collection, on the other hand, defines a closed list as shown by Fig. 14.7. In turtle, this can be done by using nested `[]` iteratively. A short cut is to use `()`, with the items in the collection listed down in the bracket as follows.

```
<subject> <predicate> (<object1> <object2> <object3>) .
```
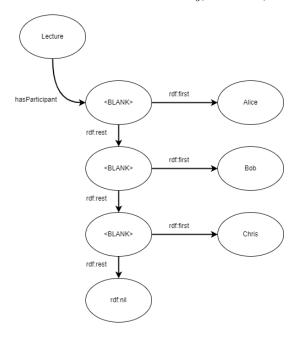
FIGURE 14.7
An example of a collection.
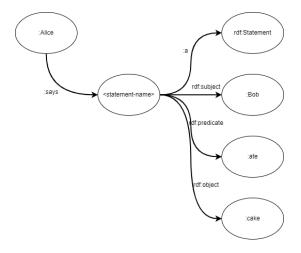
14.2.4 Reification

RDF permits interleaving of statements, i.e., to make a statement about another statement. For example, consider "Alice says that Bob ate the cake". This example contains nested statement, where "Bob ate the cake" is a statement, and "Alice says . . . " is a statement on that statement. RDF reification follows Fig. 14.8.

14.2.5 Converting RDB to RDF

For small-scale relational database, it is possible to convert it to semantic web RDF manually. For example, consider an RDB for all the books in a study. There are three tables in the database, for books, authors and publishers, respectively. Each table contains a few dozens of entries. The RDB can be converted to RDF as shown in Fig. 14.9. It can be realized very easily, simply by defining everything as nodes and stack triples for all relationships.

However, when comes to a large and complicated database, converting it to RDF manually would consume too much labor. Several ways to systematically do the conversion have been proposed. Details are not covered here. See [4] for more details.

FIGURE 14.8
An example of RDF reification.
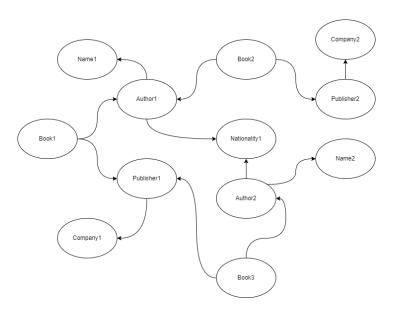
FIGURE 14.9
Semantic web of a few books, and their authors and publishers.

14.2.6 Conclusion

In summary, RDF, different from a plain XML or JSON which can also store independent classes, defines not only the objects (nodes) themselves but also the relationships among objects. This is essentially how RDF differs from a conventional key-value based NoSQL. RDF can be easily scaled up as nodes with the same name naturally merge together.

The underlying mechanism behind RDF, such as how machine stores graphical database including nodes and relationships, and how it enables query using SPARQL, is out of the scope of this notebook. In short, there are several ways to do that, for example by transforming the graph links into multiple small tables, etc. Different approaches may have their pros and cons. Details are neglected here.

14.3 RDF Schema (RDFS)

RDFS enforces schema to the RDF model. By using RDF/RDFS, a more consistent and semantic RDF model can be achieved compared with using RDF along.

14.3.1 RDFS Motivation

RDF is flexible. It is so flexible that sometimes it becomes difficult to maintain consistency of the RDF model, let alone performing sophisticated reasoning from it. RDFS, also known as RDF vocabulary description language, enforce schema to the RDF model by adding more "meta information" which builds more connections among the nodes.

RDFS expands the vocabulary of RDF. It introduces the concepts of "class" and "subclass" to RDF. It provides built-in predefined classes such as `rdfs:Literal`, `rdfs:Resource`, `rdfs:Datatype`, etc., and enforce the nodes to be linked to these classes. RDF already defines `rdf:Property`. RDFS further expands the properties and relations. All above makes the RDF modeling more consistent and semantic.

In the deeper insights, RDFS helps to add "ontology" to the RDF model by introducing the schema. It essentially integrate common understanding and domain knowledge to tine information. For example, by creating a property ":hasSpouse" whose domain and range person, it demonstrates the common knowledge that a person can be married to another person.

14.3.2   RDF Versus RDF/RDFS via an Example

The following example applies RDF versus RDF/RDFS on the same statement. Consider statement "banana is yellow, apple is red, and orange is orange color". In a RDF implementation, this would be simply be Fig. 14.10 green-colored elements. It is a disconnected graph, where the three statements are completely independent.
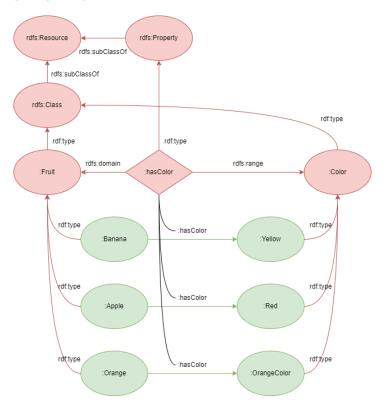
FIGURE 14.10
Semantic web of a fruits, with the green color RDF and green + red RDF/RDFS implementations.

In a RDF/RDFS implementation, class/subclass and property are enforced in the RDF model, with property domain and range specified. All classes eventually root to `rdfs:Resource`. This is shown by the green + red color in Fig. 14.10. The corresponding Turtle is

```
@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
@prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
@prefix example: <http://example.org/> .

# Define classes (optional; they are implicit)
```

```
rdfs:Class rdfs:subClassOf rdfs:Resource .
rdf:Property rdfs:subClassOf rdfs:Resource .

example:Fruit rdf:type rdfs:Class ;
rdfs:subClassOf rdfs:Resource .

example:Color rdf:type rdfs:Class ;
rdfs:subClassOf rdfs:Resource .

# Define properties
example:hasColor rdf:type rdf:Property ;
rdfs:domain example:Fruit ;
rdfs:range example:Color .

# Define fruits and colors
example:Banana rdf:type example:Fruit .
example:Apple rdf:type example:Fruit .
example:Orange rdf:type example:Fruit .

example:Yellow rdf:type example:Color .
example:Red rdf:type example:Color .
example:Green rdf:type example:Color .

# Define relationships between fruits and colors
example:Banana example:hasColor example:Yellow .
example:Apple example:hasColor example:Green .
example:Orange example:hasColor example:Orange .
```

where only the last 3 lines would be used in RDF implementation. It can be seen from this example that RDF/RDFS uses more "complicated" structure to enforce schema of the model. In this example, "a fruit has color of a color" is enforced. Notice that predicate `rdf:type` can be used interchangeably with turtle keyword `a`. They both declares a instance of a class.

14.3.3   RDFS Expanded Class and Properties

Special classes and properties are defined, some of which already used in the above example. Some commonly seen ones include:

- `rdfs:Resource`

- `rdfs:Class`

- `rdf:Property`

- `rdfs:subClassOf`

- `rdfs:subPropertyOf`

- `rdfs:domain`

- `rdfs:range`

Commonly seen further properties of a node include the following. They help to make the RDF model more human-readable.

- `rdfs:seeAlso` points to a resource where a detailed explanation of the node can be found.

- `isDefinedBy` defines the relation of a resource to its definition.

- `rdfs:comment` points to text comment.

- `rdfs:label` assign a more human-readable name to the node.

More details of the latest version of RDF/RDFS defined classes and properties are given by W3C and can be found at *w3.org/TR/rdf-schema/*.

14.3.4   Semantics inside RDF/RDFS

The semantics in RDF/RDFS, such as Fig. 14.10, is stored in the format of connections linked via properties, with domain and range specified. The name of a class or a property makes no sense to a machine when it is doing reasoning and query. The connection matters.

The following gives some intuitive examples for semantic reasoning.

- Via class inheritance. In Fig. 14.10, if apple is a subclass of fruit, and fruit a subclass of plant, then apple must also be a subclass of a plant. There is an "hidden arrow" pointing from the apple to the plant.

- Via property domain and range. In Fig. 14.10, if an unknown object "has-Color", then its object must be a fruit, and the subject a color. There are associated "hidden arrows", even if these arrows are not specified in the programming.

- Via property inheritance. Consider two properties, "isMotherOf" and "isParentOf". Both properties have the same domain and range of "person", and "isMotherOf" is a sub property of "isParentOf". In this case "A is the mother of B" can derive "A is the parent of B".

Though not given by the input directly, the hidden information exists in the semantic web and should be accessible from API

14.4   SPARQL Protocol and RDF Query Language (SPARQL)

SPARQL is the query language for semantic web. It is SQL-like in the API syntax perspective, but notice that the underlying technologies of SQL and

SPARQL differ largely, as SQL is for operations on structured tabular, while SPARQL for those of graphical database. The details of the internal mechanism of SPARQL, which involves a lot of graph theory, pattern searching and data structuring, are not covered in this chapter.

In the recent update by W3C standard for SPARQL, namely SPARQL 1.1, more operations such as advanced query and interfering, updating triples, etc., have been added to the existing SPARQL 1.0 standard, making it more powerful and capable. SPARQL 1.1 is already supported by many triplestores.

Notice that SPARQL is not only a language, but also a protocol layer. The input and return have specific formats which are also illustrated by the SPARQL standard. However, introducing of SPARQL from a protocol perspective is not the focus of this chapter, hence it is not covered in details.

More details of SPARQL 1.1 can be found at *w3.org/TR/sparql11-query/*.

14.4.1   SPARQL for Basic Query

Many query types are defined in SPARQL, each with a different purpose and associated with a different return type. For example,

- SELECT returns a tabular just like SQL.

- CONSTRUCT returns a new RDF graph based on the query result.

- ASK returns true and false of whether a query has a solution.

- DESCRIBE returns the schematic of a resource; this is useful when the structure of RDF data in the data source is unclear.

The basic syntax of SPARQL is inspired by SQL as follows.

```
SELECT [?s ?p ?o]
[FROM rdfSource]
WHERE {
       [?s/s] [?p/p] [?o/o]
}
[ORDER BY (?s/?p/?o)]
[LIMIT limitNum]
[PFFSET offsetNum]
```

where `?variableName` denotes a variable, and without `?`, a constant. Recall the fruit example used in Section 14.3.2. To perform a query to look for fruits with yellow color, use the following

```
PREFIX ex: <http://www.example.com/>
PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>

ASK WHERE {
       ?fruit rdf:type ex:Fruit .
       ?fruit ex:hasColor ex:Yellow .
} # return yes
```

```
SELECT ?fruit
WHERE {
        ?fruit ex:hasColor ex:Yellow .
} # return a table with one element, ex:Banana

CONSTRUCT {?fruit ex:hasColor ?color}
WHERE {
        ?fruit rdf:type ex:Fruit .
        ?color rdf:type ex:Color .
        ?fruit ex:hasColor ?color .
        FILTER (?color = ex:Yellow || ?color = ex:Red)
} # returns 2 triples, banana has color yellow and apple has color red

DESCRIBE ?fruit
WHERE {
        ?fruit ex:hasColor ex:Yellow .
} # return triples related to ex:Banana
```

Always put in mind that SPARQL is essentially pattern matching. The triples given in the WHERE clause are the patterns to be matched. Only the elements or triples that matches all the patterns in the WHERE clause can be returned, on top of which other filtering functions (such as FILTER which is introduced later) can be carried out.

Notice that some triplestore may require each query to be passed to the engine one at a time, instead of putting everything into a gigantic file and "run all".

When there are string and numerical attributes, "`FILTER`" keyword can be used to filter the result and returns only the entries matching the filter condition. An example is given below. It is worth mentioning that only triples in the WHERE clause need to end with a period ".". The filter condition lead by `FILTER ()` does not require the period.

```
PREFIX ex: <http://example.org/>

SELECT ?person
WHERE {
        ?person ex:age ?age .
        FILTER (?age >= 30 && ?age <= 40)
}
```

Commonly used keywords and operators in FILTER clause include `&&` (and), `||` (or), `!` (not), as well as string functions `STR`, `LANG`, `CONTAINS`, `STRSTARTS`, `STRENDS`, `STRLEN`, `SUBSTR`, `REGEX` (regular expression matching), etc., and numeric functions `+`, `-`, `*`, `/`, `>`, `>=`, `<`, `<=`, `ABS`, `ROUND`, `CEIL`, `FLOOR`, etc., and comparison operators `=`, `!=`.

Similar with FILTER clause, there are also other clauses such as OPTIONAL, UNION, etc. The OPTIONAL clause indicates that if the entry does not have the attribute to be evaluated in the clause, it still passes the

query. An example of is given below, where RDF and SPARQL are given separately. RDF:

```
@prefix foaf: <http://xmlns.com/foaf/0.1/> .

:alice
a foaf:Person ;
foaf:name "Alice" ;
foaf:mbox <mailto:alice@example.com> .

:bob
a foaf:Person ;
foaf:name "Bob" .
```

SPARQL:

```
PREFIX foaf: <http://xmlns.com/foaf/0.1/>

SELECT ?name ?email
WHERE {
        ?person a foaf:Person ;
        foaf:name ?name .
        OPTIONAL { ?person foaf:mbox ?email . }
}
```

The above returns both names and Alice's email, and it won't fail although Bob does not have an attribute of email.

The UNION clause allows the combination of two queries as one return, given that the structure of the queries matches. An example is given below.

```
PREFIX foaf: <http://xmlns.com/foaf/0.1/>
SELECT ?name
WHERE {
        {
                ?person a foaf:Person .
                ?person foaf:name ?name .
        }
        UNION
        {
                ?organization a foaf:Organization .
                ?organization foaf:name ?name .
        }
}
```

which would return both people and company names in one go.

14.4.2   SPARQL for Advanced Operations

SPARQL 1.0 provides basic query functions based on graph pattern matching, as already explained in the previous section. SPARQL 1.1 provides advanced query functions as follows.

Advanced Query

In the SELECT clause, it allows simple calculations on top of the returned result, and name it as a new column. For example,

```
SELECT ?x (?y*1.1 AS ?z)
WHERE {
       ?x ex:hasProperty ?y .
}
```

returns `?y*1.1` instead of `?y`, and furthermore rename it as `?z`.

SPARQL 1.1 enables aggregate functions. For example,

```
PREFIX ex: <http://www.example.com/>
PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>

SELECT (Count(?fruit) AS ?numOfFruit)
WHERE {
       ?fruit rdf:type ex:Fruit .
}

SELECT (Count(DISTINCT ?fruit) AS ?numOfFruit)
WHERE {
       ?fruit rdf:type ex:Fruit .
}
```

counts the total number of fruits in the database. Notice that when using aggregate functions, new variable name (in this example, `?numOfFruit`) must be assigned. Otherwise, there will be an syntax error.

Data Editing

SPARQL provides operations to insert, edit, and delete elements in a semantic web as follows.

- INSERT inserts triples into a graph.

- DELETE deletes triples from a graph.

Recall the fruit example used in Section 14.3.2. To create that RDF/RDFS database from scratch using SPARQL, use the following (notice that "orange is orange color is changed to pear is green")

```
PREFIX ex: <http://www.example.com/>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>

INSERT DATA {
       # Define resources
       ex:Banana rdf:type ex:Fruit .
       ex:Apple rdf:type ex:Fruit .
       ex:Pear rdf:type ex:Fruit .
```

```
       ex:Yellow rdf:type ex:Color .
       ex:Red rdf:type ex:Color .
       ex:Green rdf:type ex:Color .

       # Define properties
       ex:hasColor rdf:type rdf:Property ;
       rdfs:domain ex:Fruit ;
       rdfs:range ex:Color .

       # Link resources with properties
       ex:Banana ex:hasColor ex:Yellow .
       ex:Apple ex:hasColor ex:Red .
       ex:Pear ex:hasColor ex:Green .
}
```

14.4.3   Default Graph and Named Graph

It is worth mentioning that it is possible to create and query "disjoint" graphs using SPARQL. If the graph name is not specified during the inserting of the data, the data is inserted into the default graph. A query can have multiple FROM clause, each specifying a graph. If the graph name is not specified during the query, all graphs will be searched. An example of inserting and querying named graph is given below.

```
INSERT DATA {
       GRAPH <http://example.org/mygraph> {
               <http://example.org/book1> <http://purl.org/dc/elements
                   /1.1/title> "A new book" .
       }
       GRAPH <http://example.org/myothergraph> {
               <http://example.org/book2> <http://purl.org/dc/elements
                   /1.1/title> "Another book" .
       }
}

SELECT ?book ?title WHERE {
       GRAPH <http://example.org/mygraph> {
               ?book <http://purl.org/dc/elements/1.1/title> ?title .
       }
}
```

The relationship of the graphs is somehow like the different tables in an RDS. They usually represent logically separated things, but they can have some overlaps. When doing analysis information can be extracted from multiple graphs at the same time.

14.4.4 SPARQL Returns

As introduced earlier, SPARQL is an HTML based protocol. The returned result of a SPARQL query can be specified in the HTML header. There are several return types defined in the standard. When SELECT or ASK are used, the return types can be:

- XML

- JSON

- TSV (table-separated values, similar to CSV)

When CONSTRUCT or DESCRIBE are used, the return is an RDF that can be in

- RDF/XML

- Turtle

- N-Triples

- JSON-LD

and a few more. The decoding of the return can be done my relative tools and packages, and are not given in details here.

14.4.5 A Bit More Insights on RDF/RDFS

RDF can be represented as a set of triples, which makes it readable for both humans and machines. When using SPARQL to query or edit RDF database, RDF is treated as directional graphs with some sort of hierarchy. This is the interface of RDF database, following W3C standard.

How about the underlying data structure that ultimately supports the interface, making sure that everything runs smoothly and efficiently? Different triplestores may have different preference when choosing the underlying data structure. Some of them are introduced below.

Structured Tables

An intuitive way of storing triples is to use a structured table that contains 3 columns: subject, predicate (property), and object. To save some space, all subjects, predicates and objects can be encoded into integers, where there are additional translation tables that can be used to decode the data.

In this case, a SPARQL query needs to be convert into SQL (introducing self-join) then performed on the table. One of the major problems of this architecture is that it is too computationally expensive when the query is complicated.

RDF should support multiple graphs. To enable multiple graphs, an additional column indicating graph ID needs to be added to all the tables mentioned above.

RDB

To reduce self-join in the query, table size needs to be made smaller. The big structured tables need to be split into multiple small tables via aggregation. Due to the nature of triples, there will be many "NULL" in the tables. Implementing multi-value property, etc., can also be complicated. Overall, it is just too difficult to design the optimal RDB schematics, let along creating a database of that schematics.

To systematically and automatically split big tables into small tables without warring too much about the schematics, someone has proposed destruct the table to the very ground level. Each property becomes a small table that contains its subjects and objects. The problem of this approach, however, is that it becomes time consuming to loop over all the tables when the RDF model is large. The performance of the architecture will be bad, and things such as inserting will become expensive.

NoSQL

Structured table and RDB approaches are rarely used in commercialized triplestores due to the expensive computational cost. The most widely used approaches nowadays are NoSQL based.

Two commonly seen NoSQL data structures are vertically partitioned tables and hexastores. These structures help to either reduce or speed up joins, making the computation more efficient than the RDB based methods.

14.5   Web Ontology Language (OWL)

By adding schema to RDF, RDFS already boosts the capability and adds ontology to the RDF model. However, there are still a number of limitations using RDF/RDFS alone that can cause problems and issues. OWL is proposed to solve these problems and further expand the capability of the RDF model.

Notice that it is OWL, not WOL, for historical contexts reasons. OWL, the English word for the bird, also embodies qualities like wisdom and perceptiveness, which align well with the goals of the language.

14.5.1   RDF/RDFS Limitations

Several examples are given to demonstrate the limitations of RDF/RDFS.

RDF/RDFS brings schema to the RDF model by introducing hierarchy to

classes and properties. For example, in the schematic part of the model we can create "animal eats food", where "eat" is a property with domain "animal" and range "food". We can then add instances to animal and food classes. This is demonstrated in Fig. 14.11. The SPARQL to create such an RDF model is given below.

```
PREFIX ex: <http://www.example.com/>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
PREFIX rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#>

INSERT DATA {
        # Define resources
        ex:Human rdf:type ex:Animal .
        ex:Cow rdf:type ex:Animal .
        ex:Vegetable rdfs:subClassOf ex:Food .
        ex:Meat rdfs:subClassOf ex:Food .
        ex:Cabbage rdf:type ex:Vegetable .
        ex:Ribeye rdf:type ex:Meat .

        # Define properties
        ex:eat rdf:type rdf:Property ;
        rdfs:domain ex:Animal ;
        rdfs:range ex:Food .
}
```
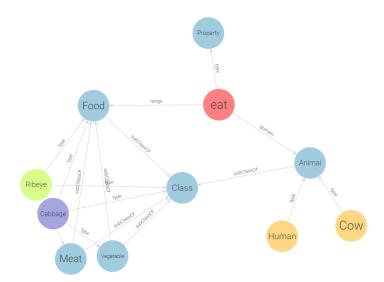

FIGURE 14.11
An RDF model that demonstrates animal eats food.

From Fig. 14.11, it can be seen that in this RDF model, we can derive all the following statements:

- Human eats rib-eye.

- Human eats cabbage.

- Cow eats rib-eye (this is not what we expect).

- Cow eats cabbage.

Since both vegetable and meat are defined as sub classes of food, and animal eats food in general, "cow eats rib-eye" can be deduced. RDF/RDFS cannot exclude cow from eating meat, or to say it in general, RDF/RDFS cannot add restrictions the property for a specific object.

A walk around of this is to make "vegetable" and "meat" completely separated classes without the shared parent "food", and define 3 properties, namely "human eats vegetable", "human eats meat" and "cow eats vegetable". However, this destroys the schema of the model and make null the effort of RDFS.

Similarly, consider an example where "hasVisualAcuity" is used to describe human's clarity of vision. A person usually have two visual acuity values for the two eyes. In RDF, that means a person should have two and only two of this property. This cannot be enforced by RDF/RDFS.

Consider another example where "man is a sub class of human" and "woman is a sub class of human". With only RDF/RDFS, however, there is no way to add restrictions that says "an instant cannot be man and human at the same time", i.e., to disjoint man class and woman class.

RDF/RDFS does not support class combinations. For example, for existing classes "Car", "Motorcycle", "Bicycle", "Ship", "Plane", RDF/RDFS cannot create a new class "AllTranportationTool" to automatically include every instant from the above classes. In general, with RDF/RDFS alone, it is not possible to define a class which is a union/intersect/complement of other classes.

In semantic web stack Fig. 13.1, one layer above RDF/RDFS, OWL is proposed to solve the above problems.

14.5.2   OWL Vision

As explained earlier in Section 13.4, OWL essentially introduces DL inference to semantic web, making it more expressive and capable of more complicated reasoning. With that in mind, OWL provides additional features to enforce and enhance the schema of an RDF model, and solve the problems mentioned in Section 14.5.1. It allows adding relations and attributes constraints (rules) to the REF model. The following summarizes the main features introduced by OWL.

- Allow the disjoint of sub classes (an instant cannot be man and woman at the same time).

- Allow enforcing the number of attributes of each property type (one can have only one social security number, but can have many mail addresses)

- Allow enforcing the range of the values an attribute can take.

- Offer more expressive class definitions including union, intersection, complement, etc.

- Enlarge the vocabulary sets on top of RDF/RDFS.

which align well with what is capable in DL. This is not surprising as OWL is DL practice in semantic web technology.

For the machine to interpret the additional semantics, there are at least two things to consider: the interface, which concerns how a user tells the rules to the machine; the underlying structure and mechanism, which concerns how the machine enforce the rules. In this section, we mainly focus on the interface.

The latest version of OWL is currently OWL 2, as has been recommended by W3C since 2009. It makes the following assumptions on top of DL:

- OWA is made. In this context, the knowledge base is considered open, and absence of information must not be valued as negative information.

- An instance can have non unique names, i.e., multiple syntax can be mapped to the same instance via interpretation.

In practice OWL comes in different flavors, including OWL Lite, OWL DL and OWL Full, just to name a few. OWL 2 supports different choices of syntax, including RDF-syntax, XML-syntax, functional-syntax, Manchester-syntax, and turtle.

14.5.3   OWL Basic Syntax

Unless otherwise mentioned, turtle is used when introducing OWL syntax. Classes, individuals, and properties are already introduced in RDF/RDFS. OWL also defines these concepts, each with enriched features. It is worth mentioning that OWL defines two special classes, `owl:Thing` and `owl:Nothing`, corresponding with the top class (universal set) and bottom class (empty set) in DL, respectively.

Class and Subclass

Defining a class from the root `owl:Class` using OWL is given in the following example in turtle.

```
:Fruit a owl:Class ;
```

where `a` can be replaced by `rdf:type`. To define an individual via class membership, simply use

```
:Apple a :Fruit .
```

which in DL this is equivalent of saying `Fruit(Apple)`. It is also possible to define an individual without a named class as follows, which indicates that the named class will be added in a later stage. In this case, `owl:NamedIndividual` is used.

```
:ANewThing a owl:NamedIndividual .
```

Subclass can be defined similarly as follows.

```
:Fruit a owl:Class ;
       rdfs:subClassOf :Food .
:Meat a owl:Class ;
       rdfs:subClassOf :Food .
:Food a owl:Class .
```

Notice that it was impossible to enforce class disjoint using RDF/RDFS alone. This is made possible in OWL as follows.

```
:Fruit owl:disjointWith :Meat .
```

Or alternatively,

```
[] a owl:AllDisjointClasses ;
       owl:members
       (       :Fruit
               :Meat
               :SoftDrink
               :Beer ) .
```

which is a shortcut when disjoint is claimed on multiple classes. Similarly, class equivalence can be defined using `owl:equivalentWith` keyword.

Property

There are two types of properties defined in OWL, namely the object property which links to another resource URL (another object), and datatype property which links to a literal. Examples are given below.

```
:hasColor a owl:ObjectProperty ;
       rdfs:domain :Fruit ;
       rdfs:range :Color .
:hasShelfLife a owl:DatatypeProperty ;
       rdfs:domain :Fruit ;
       rdfs:range xsd:integer .
```

where notice that XML schema definition (XSD) defines many data types and sub data types, including `xsd:integer`, `xsd:string`, `xsd:decimal`, `xsd:boolean`, `xsd:date` (a calendar date), `xsd:time` (time in a day), `xsd:dateTime`, `xsd:double` (64-bit floating number), `xsd:float` (32-bit floating number), `xsd:anyURI`, `xsd:language`, etc.

14.5.4 OWL Advanced Syntax

Equivalent and Different Individuals

In RDF/RDFS, different individuals are distinguished by their names, and each individual can have one and only one name (it is technically possible for an individual to have no name, but it is not often a good practice). In OWL, it is possible to map multiple names to the same individual, and to emphasize that two names are pointing to different individuals (this is sometimes necessary due to the OWA). Examples are given below.

```
:Computer a owl:Class .
:HarryPc a :Computer .
:PotterPc a :Computer ;
      owl:sameAs :HarryPc .
:DumbledorePc a :Computer ;
      owl:differentFrom :HarryPc .
```

To state that individuals are different, alternatively use the following

```
[] a owl:AllDifferent ;
      owl:distinctMembers
      ( :HarryPc,
        :DumbledorePc,
        :HermionePc,
        :HagridPc   ) .
```

Closed Class

It is possible to define an enumerate class, whose instances are taken from individuals. An example is given below.

```
:Weekdays a owl:Class ;
      owl:oneOf
      (       :Monday
              :Tuesday
              :Wednesday
              :Thursday
              :Friday ) .
```

which indicates that any individuals to be defined under class "Weekdays" must be one of the 5 individuals "Monday", "Tuesday", etc. Do not confuse `owl:oneOf` with `owl:unionOf`, as the former creates class from individuals, and the later creates class from classes.

Class Constructors

Intersection, union and complement are the commonly seen class (set) constructors. They can be defined as follows.

```
:MyFavourateBook a owl:Class .
```

```
:HisFavourateBook a owl:Class .
:OurFavourateBook a Class ;
      owl:intersectionOf (:MyFavourateBook :HisFavourateBook) .
:Wine a owl:Class ;
      owl:equivalentClass [
            owl:unionOf (
                  :DryWine
                  :MediumWine
                  :SweetWine
            )
      ] .
:Plant a owl:Class ;
      rdfs:subClassOf [
            owl:complementOf :Animal
      ] .
```

Property Restrictions

It has been introduced earlier that the type of the property can be specified by using either `owl:ObjectProperty` (specify range as a class) or `owl:DatatypeProperty` (specify range as a datatype such as string, decimal, integer, etc). In the context of OWL, more restrictions can be enforced to properties, including the number of each property, and the values each property can take. Details are given below.

The syntax of adding different property restrictions may differ slightly. A general syntax is given by

```
:<ClassOfProperty> rdfs:subClassOf [
    rdf:type owl:Restriction ;
    owl:onProperty :<PropertyName> ;
    owl:<RestrictionType> <RestrictionDetails>
] .
```

where notice that blank node is used for property restrictions.

For example, to indicate that the value of a property must be taken from a class, use

```
:<Class> owl:restriction [
      owl:onProperty :<PropertyName> ;
            owl:allValuesFrom :<RestrictionClass>
] .
```

And to indicate that among all the property values of a property, at least one of them must take values from a class, replace `owl:allValuesFrom` with `owl:someValuesFrom`.

To indicate that a property must take specific value, use `owl:hasValue` at the restriction type, and put the value as `<RestrictionDetails>`, whether an individual or a datatype value.

To restrict the number of a property, use `owl:maxCardinality`, `owl:minCardinality` or `owl:cardinality`, and put a number as `<RestrictionDetails>`.

It is possible to define a class from property restrictions, essentially saying "the class is defined as a collection of individuals that satisfy the following property restrictions". To do that, use `rdfs:subClassOf` with property restrictions as given in the example below.

```
:MyBook a owl:Class ;
      rdfs:subClassOf [
              a owl:Restriction ;
              owl:onProperty :hasOwner ;
                      owl:hasValue :Me
      ] .
```

where `:Me` is a predefined individual that maps to myself. To put it in words, a book is considered an instance of `:MyBook` if it has a `:hasOwner` property with the value `:Me`.

Property Hierarchy

Properties can have hierarchy like classes. Properties hierarchy can be realized using `rdfs:subPropertyOf` in RDF/RDFS framework. OWL further enriches that idea, by introducing new concepts such as `owl:inverseOf`. Examples below are used to demonstrate property hierarchy.

To define property hierarchy, use

```
:isOwnerOf a owl:ObjectProerty ;
      rdfs:subPropertyOf :isMasterOf ;
      rdfs:domain :Person ;
      rdfs:range :Book .
:isOwnedBy a owl:ObjectProperty ;
      owl:inverseOf :isOwnerOf ;
      rdfs:domain :Book ;
      rdfs:range :Person .
```

Notice that when defining a inverse property, it is not necessary, but still good practice, to point out the domain and range. This is for better readability and easier error checking.

OWL further defines the following features of a property.

- `owl:TransitiveProperty`: if `Property(a,b)` and `Property(b,c)`, then `Property(a,c)`. An example is "isAncestorOf".

- `owl:SymmetricProperty`: if `Property(a,b)`, then `Property(b,a)`. An example is "isColleagueWith"; there is also `owl:AsymmetricProperty`, which indicates that if `Property(a,b)`, it is not possible that `Property(b,a)`.

- `owl:FunctionalProperty`: if `Property(a,b)` and `Property(a,c)`, then a=c. An example is `hasMother`.

- `owl:InverseFunctionalProperty`: if `Property(a,c)` and `Property(b,c)`,

then `a=b`. Examples are `isMotherOf`, `hasId` (assuming ID is unique for each individual).

Similar with class declaration where `owl:disjointWith`, `owl:AllDisjointClasses` can be used to claim disjoint of classes, it is possible to declare disjunctive properties using `owl:propertyDisjointWith` and `owl:AllDisjointProperties` as follows. Disjunctive properties cannot take the same value. For example,

```
:hasParent a owl:ObjectProperty ;
      rdfs:domain :Person ;
      rdfs:range :Person .
[] a owl:AllDisjointProperties ;
      owl:members ( :hasParent :hasChild :hasSibling ) .
```

In OWL, top and bottom classes are defined, corresponding to the universal set and empty set, respectively. Similar ideas apply to properties. Object and datatype properties each has its top and bottom properties.

It is possible to enforce NOT related via a property. An example is given below.

```
[] a owl:negativePropertyAssertion ;
      owl:sourceIndividual :SherlockHolmes ;
      owl:assertionProperty :isMurderOf ;
      owl:targetIndividual :Watson .
```

Notice that negative assertion is supported only at individual-to-individual level. It is currently not possible to claim that Sherlock Holmes is not a murder of any victims, i.e, `:Watson` cannot be replaced by a class such as `:Person`. In FOL, this would have been possible. But currently OWL does not adopt FOL for a good reason (FOL can be undecidable, and too computationally expensive).

General Role Inclusion

If we define "B is the father of A" and "C is the brother of B", then naturally, "C is the uncle of A" can be derived. This role chain can be realized via OWL general role inclusion. However, this adds uncertainty to the reasoning, and may cause the model to be undecidable. This is one of the reasons why different tiers of OWL and OWL 2 have been proposed. More features usually mean more computations and higher risks of undecidable results, and the user needs to decide which tier to use.

And do notice that due to the Gödel's incompleteness theorems, there is no sophisticated enough system or knowledge base that can use finite input to derive all the knowledge in a complete (every statement can be proved true or false) and consistent (no contradiction) way.

Using role chain, we can define `isUncleOf` as follows.

```
:isUncleOf a owl:ObjectProperty ;
      owl:PropertyChainAxiom ( :hasFather :hasBrother ) .
```

It is not recommended ot use role chain in a semantic web.

14.5.5   Beyond OWL: Rules

Rules (rule-based systems, also known as rule-based engines) are used to express logic beyond DL, i.e., beyond what RDF/RDFS and OWL can describe.

The basic syntax of a rule is simply

```
IF A THEN B
```

but there are many variations. There are different types of rules, such as logical rules (FOL rules, similar with $A \rightarrow B$), procedural rules, and logic programming rules. There are semantic web triplestores that support rules.

15

Semantic Web Practice

CONTENTS

This chapter studies the design and development of ontology and semantic web model in practice. Both methodologies (ontological engineering) and tools (triplestores) are introduced. Examples are given.

15.1   Brief Introduction to Ontological Engineering

Ontological engineering focuses on the methodology to design the ontology or the semantic web for a specific application or task (recall Fig. 13.2).

In practice, to model a physical system, we need to define classes, properties, and instances. For class, we need to define class hierarchy and classes relationships (intersection, union, etc.). For properties, we need to define property type, hierarchy, domain and range, property features (transitive, symmetric, etc.) and restrictions. There might be many different ways to design and describe the same physical system. Ontological engineering studies the systematical way of designing, comparing, and merging of ontologies.

Ontology engineering has the following sub areas:

- Ontology design: it studies the methods to systematically design, develop, and devploy ontology models.

- Ontology mapping: it studies the methods to efficiently compare different ontology models.

- Ontology merging: it studies the methods to efficiently combine ontology models.

- Ontology learning: it studies the automatic learning of new knowledge by an existing ontology model, when new data sets are provided.

15.2   Ontology Design

Ontology design methodology describes all activities necessary for the construction of an ontology model. It helps with consistent, efficient, and distributed (collaborative) ontology model design.

15.2.1   Ontology Management Activity

Design, develop and deploy ontology model for a sophisticated system can be time and manpower consuming. It is important to manage each step during the entire procedure to ensure healthy development of the system. This include

1. Scheduling: identify tasks and problems to be solved by the semantic web; plan and arrange resources, time, manpower and money ahead.

2. Control: guarantee correct execution of tasks and problems to be solved.

3. Quality assurance: guarantee all steps are done correctly, including using the correct software, and everything is documented in details.

In pre-development stage, environment study needs to be carried out. This is mainly to identify what software to use to host the semantic web, and what interface/API shall the semantic have, and what applications would call the API to talk to the semantic web. A feasibility study is also necessary. For example, we need to consider whether it makes sense to develop a semantic web for the application, and whether the semantic web design is realizable.

In development stage, domain expert needs to come in to structure domain knowledge in a conceptual model. Knowledge engineer or data scientist then formalize the conceptual model into a computable (formal) model, then into ontology representation language.

Finally, in the post-development phase, pipeline needs to be designed to maintain, update and scale up and down the ontology model. In the case

where the ontology model needs to be migrated into different platforms, used by unplanned applications, or merged with other models, necessary changes need to be made to the model. In the case when the knowledge in a model needs to be exported, knowledge recycle needs to be supported.

There are many ontology support activities. These activities need to be carried out during the different stages of ontology development. Below is a list of these activities summarized in bullet points.

- Knowledge acquisition. Interview experts in the field and learn domain knowledge from them. This is referred as ontology learning. This is often done manually by the knowledge engineer or data scientist in the beginning stage. It is also possible to develop tools to automatically gather information and transfer it into ontology models, and even merge it with existing models.

- Technical evaluation. A domain expert checks occasionally to make sure the ontology model is correct.

- Integration and merging. This refers to the case where a big (scaled-up) ontology model can be built from a small existing ontology model.

- Alignment. Where there are multiple ontology models describing the same physical thing, alignment needs to be made to make sure knowledge consistency.

- Documentation and configuration management.

15.2.2 Ontology Design Basics

To design a complicated ontology model, many methods and techniques need to be used comprehensively. It is too complicated to be covered here.

This section introduces the basics of ontology design. It covers a subset of the aforementioned methods and techniques in a very brief manner.

15.3 Linked Data Engineering

"nobreak

15.4 Semantic Search

"nobreak

15.5 Triplestore

Triplestore is the DBMS of RDF/RDFS and OWL.

When it comes to relational database, there are many choices of DBMS such as Microsoft DBMS, Oracle DBMS, MySQL, MariaDB, and many more. Similarly, there are many choices of triplestore for semantic web. To name a few, a list of widely used triplestores is given below.

- GraphDB: a commercialized enterprise-tier semantic graph database management system compliant with W3C standards. It is famous for its performance and inference capabilities. It also provides free-tier for learning and for small projects, with limited capability.

- Apache Jena: an open-source Java framework for building semantic web and linked data applications. It has RDF APIs that can read and process RDF and SPARQL written in XML, Turble, JSON-LD and N-Triples.

- Virtuoso: a multi-model DBMS for both RDB and NoSQL databases such as RDF. It is famous for its scalability and standards compliance.

- Many more, such as Stardog, Blazegraph, AllegroGraph, etc.

More triplestores can be found at W3C, where a list of triple stores is maintained. At this moment of 2023, there are 50 triplestores registered at W3C.

For the demonstrations given in this notebook, GraphDB is used, unless otherwise mentioned. A full instruction on using GraphDB, including applying for access and installation of the software, can be found at *graphdb.ontotext.com*.

Notice that RDF/RDFS/OWL/SPARQL APIs can be enabled using packages or libraries of many programming languages. For example, in Python, there are several packages available for semantic web operations, including `rdflib`, `Owlready2`, `SPARQLWrapper`, etc. Some of these packages have "lite" triplestore engine built-in which provides some SPARQL features for in-memory operations. They do not have all the features of a triplestore. However, they can connect to a triplestore API, in which case they serve as Python-to-triplestore interfaces.

15.6 Example: Semantic Web for Home Assets

As an example, we are creating a semantic web for a the assets in a household. Here "assets" refer to the electrical products, furniture, LEGOs, and other non-consumable, relatively static elements.

We will start with defining classes to divide everything into large groups, including electrical product, furniture, toy, etc. Under each class, sub-classes are defined, such as TV, computer, game console under electrical product, bed, chair, sofa, lamp under furniture, and LEGO under toy. Lastly, we will define instances under each sub-class. For example, bedroom TV and living room TV under TV, Nintendo SWITCH under game console, living room sofa under sofa, study computer, living room computer, TV attached computer under computer, etc.

We will then use RDFS to enforce schema to the RDF model as follows. Consider electrical product class for example. All elements in this class shall have a property called "hasBrand", which maps them to a pre-defined "electricalBrand" class, inside which are commonly seen electrical brands such as Boche, Siemens, Nintendo, Sony, Google, etc. The electrical product shall also have a "hasPrice" property, "warrantyExpiresAt" property, etc. The similar concept applies to all other produces including furniture, etc.

Finally, use OWL to setup some limits of the properties. For example, for electrical products, the price is usually between 0 to 5000 dollars, etc.

The result should be something like Fig. 15.1, after visualization.

15.6.1 Define Classes Hierarchy

"nobreak

15.6.2 Define Properties Hierarchy

"nobreak

15.6.3 Add OWL

"nobreak

15.6.4 Data Retrieval Examples

"nobreak

15.7 Reference: Commonly Used Namespace

Commonly used built-in name spaces are summarized in Table 15.1.

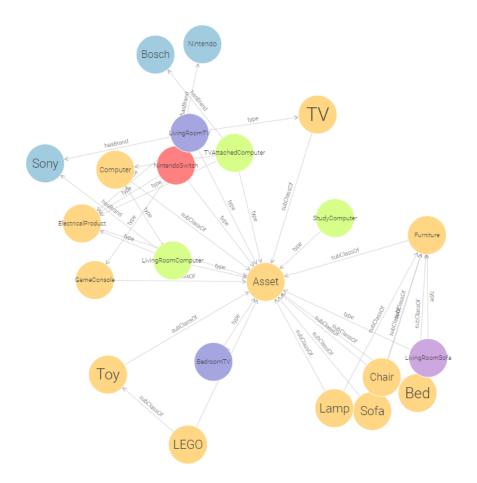To search URI for a name space, use *prefix.cc*.

FIGURE 15.1
An example of an RDF model in GraphDB that describes house assets. This is only a demonstration graph and the information inside is artificial and not true.

TABLE 15.1
Commonly used name spaces in RDF models. URI is neglected since they can be easily found online.

Namespace	Description
rdf	RDF syntax.
rdfs	RDFS syntax.
xsd	XML syntax.
foaf	Friend-of-a-friend. It describes people, their activities and relations to other people and object.

Bibliography

[1] The common layered semantic web technology stack.

[2] The R project for statistical computing.

[3] Rstudio.

[4] Franck Michel, Johan Montagnat, and Catherine Faron Zucker. *A survey of RDB to RDF translation approaches and tools.* PhD thesis, I3S, 2014.